

From Scratch:
The Design and Implementation of a SAT Solver
(First Draft)

Oliver Lie

August 21, 2026

Contents

Chapter 1	Introduction	5
1.1	What Do SAT Solvers Even Solve?	6
1.2	Our Environment	7
1.3	CNF Encoding	7
1.4	Reading a CNF Equation Programmatically	11
Chapter 2	Brute Force	15
2.1	Optimizing Brute Force	17
2.2	Brute Force - Results	20
Chapter 3	DPLL	21
3.1	Recursive Backtracking	22
3.1.1	Recursive Backtracking – Results	26
3.2	Unit Propagation	26
3.2.1	Changes To Our Implementation	27
3.2.2	“Textbook” Unit Propagation	28
3.3	Better Unit Propagation	31
3.3.1	Better Unit Propagation - Results	36
3.4	Pure Literal Elimination	37
3.4.1	Implementing Pure Literal Elimination	38
3.4.2	Pure Literal Elimination - Results	39
3.5	Watched Literals	40
3.5.1	A Visual Example of Watched Literals	42
3.5.2	Algorithmic Changes	44
3.5.3	Why Two Watches?	51
3.5.4	Watched Literals - Results	52
3.6	Iterative DPLL	53
3.6.1	Iterative DPLL - Results	59
Chapter 4	CDCL	61
4.1	Clause Learning	62
4.1.1	Clause Resolution	69
4.1.2	Clause Learning - Results	72
4.2	Non-Chronological Backjumping	73
4.2.1	Non-Chronological Backjumping - Results	76
4.3	First UIPs	76
4.3.1	Implementing First UIPs	77
4.3.2	First UIPs - Results	80

4.4	A Second Round of Improvements	80
4.4.1	Removing Hashing	81
4.4.2	Fixing Clause Resolution	82
4.4.3	Improvements - Results	88
4.5	VSIDS	88
4.5.1	cVSIDS	89
4.5.2	mVSIDS	91
4.5.3	EMA/Normalized VSIDS	93
4.5.4	AMA/adaptVSIDS	96
4.5.5	VSIDS - Results	97
4.6	Clause Database Management	99
4.6.1	N-Order Learning	99
4.6.2	M-Size Relevance-Based Learning	104
4.6.3	Combining M-Size and K-Bounded Learning	110
4.6.4	“Clause VSIDS”	111
4.6.5	Clause Database Management - Results	117
4.7	Restarts	121
4.7.1	Assignment Polarity	123
Chapter 5	Conclusion	127
5.1	Improvements	127
5.1.1	This Paper is Poorly Written	128
5.1.2	The Code Needs to be Cleaned	128
5.1.3	We Don’t Explore Other Real Solvers	128
5.1.4	There Are No Actual Citations	129

Preface

Hello, thank you for picking up this book. I have just finished writing it, except for this portion which I'm writing right now! It has been one epicly long journey of 9 weeks in my last quarter of college before I graduate. I wasn't even sure if I would be able to complete this by the end, but it seems just like my assignments, I finished this at the last minute.

To be completely honest, I am not an expert when it comes to computer science. After all, my degree is in environmental sciences (just kidding, I am a computer science major). However, even though I don't know everything about every data structure, algorithm, design pattern, or whatever, I think that this book would be great for someone getting started in the world of SAT solving. These things that I am not an expert in are my passion, and by doing this technical research, I hope to become a little less of "not an expert" in those fields.

We are not going to revolutionize SAT solving, in fact, we're barely going to scratch the surface in this book (it's a big world out there). What I aim to do in writing this is to help you understand how SAT solvers work and how they are implemented. I hope that this remains accessible and I'm not too confusing, however, I am just a ~~girl~~ undergraduate. Because we're covering how SAT solvers work and are implemented, if you are a student tasked with building one, don't use this to cheat. That would hurt my feelings. Instead use this to help your own understanding instead of copy-pasting.

Throughout this book we will build our own modern SAT solver from scratch, complete with all the bells and whistles. We will begin with the most basic algorithm imaginable, brute force, and gradually work our way through DPLL, CDCL, and beyond.

Throughout this journey we will encounter thought experiments, exercises, and a handful of experiments intended to solidify our understanding. Any external resources will be cited (hopefully correctly), and in keeping with the goal of accessibility:

- Questions that are explicitly asked will eventually be answered.
- Proofs will be explained in informal language before becoming formal.
- Code will be provided and linked through GitHub (or another hosting platform).
- Revisions to this book will be made whenever I discover mistakes or better explanations.

There is one more thing that deserves acknowledgement: some pieces of code were AI-assisted. To be completely transparent, "AI-assisted" means that instead of spending

days repeatedly banging my head against a wall trying to understand a bug or some subtle algorithmic detail, I occasionally asked an AI to explain concepts or help debug an implementation. Every word of this book, however, is my own. AI did not generate any of the explanatory text you will read here, although it certainly helped produce some of the code that accompanies it. I know that AI is very unpopular, and I do get that, however, even with literature and code I could find, it wasn't enough **for me** to figure out how everything is implemented from the ground up. That being said, I would still encourage you to read the other literature out there and look at code on your own, and perhaps you'll find me an idiot for needing AI assistance in the first place.

There are practical reasons for this. SAT solvers are difficult to implement correctly, and introducing subtle bugs is remarkably easy. In fact, before deciding to write this book, I had already attempted to build a SAT solver once before. I never managed to get a working CDCL implementation.

That failure, however, taught me a tremendous amount. So in a sense, we are learning SAT solving together (isn't that comforting?). My previous mistakes have shown me many of the pitfalls, and my hope is that by documenting this journey I can help you avoid making the same ones. So without further ado, let's begin. By the time we reach the end, I hope I will have taught you more than I myself learned.

Chapter 1

Introduction

SAT solvers, in my opinion, are one of the most underrated pieces of algorithmic technology of the modern day. They are used in both software and hardware verification (circuitry testing), product configuration, package management, computational biology, particle physics, solving graph problems, and much much more.¹

In a more computer science-focused world, they have applications in NP-Completeness. Since (nearly)² every problem has a reduction to a Boolean SAT problem, creating one really really good SAT solver lets us solve all those problems in a “fast” manner. So they are important everywhere despite the vast majority of us never really thinking about their existence. One note on what “fast” means: even if solving a SAT problem takes a small amount of time, the algorithmic complexity isn’t necessarily “fast” as there is no known polynomial algorithm for any NP-Complete problem (I am a firm believer that $P = NP$, and yes I know that makes me crazy).

Before we start writing out code, proofs, and the like, we need to make sure we’re all on the same page on things such as input and ideas. Reading through this introduction is important so that you understand the basic notation used throughout this book.

One last thing to note: we will write out pseudocode rather than the code for one specific programming language. However, although the idea behind pseudocode is that it’s programming language agnostic (it doesn’t carry artifacts of any real programming language (usually)), this idea is poorly executed when it comes to formatting. My algorithms professor will probably scream at me for this, but I’m making the executive decision that we will follow a general C-based syntax for readability purposes. That is, we will use braces “{ }” for block scope, parentheses “()” for grouping, “=” for assignment, “==” for equality, and “<=”, “>=”, etc. for their respective uses. Everything that can be inferred visually will be used. Another thing to note is that we will pretend our arrays can dynamically resize, so we won’t pre-allocate space,³ and we will also assume 1-based indexing for an easier map from variables to indices, as you will see

¹<https://www.cs.utexas.edu/~isil/cs389L/lecture4-6up.pdf>

²Before anyone comes after me, the Halting Problem doesn’t have a reduction to a boolean satisfiability problem, so take that.

³If you are implementing this in real code, this tiny detail of no pre-allocation of memory will come back to bite you later, so if you are encountering segfaults, this *should* be one of the first things you should check!

later.

1.1 What Do SAT Solvers Even Solve?

“What do SAT solvers even solve,” I hear you say? SAT solvers solve *boolean satisfiability* problems. Another way to say that is they solve *propositional logic* problems. There are many forms of propositional logic, and you may be familiar with it. Here are some basic examples:

- $P \rightarrow Q$, an implication statement
- $(A \wedge B)$, an and statement
- $(A \vee B)$, an or statement
- $\neg A$, a negation/not statement

Of course these are very basic examples, and we haven’t included every operator such as xor, nand, etc. As you know, we can rewrite implications to contain only and, or, and not operators, which is what SAT solvers operate on. In other words, SAT solvers only operate on Boolean algebra equations.⁴

However, SAT solvers don’t operate on just any Boolean algebra equation. No no, they must be in a proper form, called CNF (more on that in the next section). CNF (Conjunctive Normal Form) equations are a “conjunction of disjunctions,” i.e., clauses of ors, combined with ands.

Some examples of CNF are as follows:

- $(A \vee B)$
- $(A \vee B) \wedge (\neg C)$
- $(A \vee B) \wedge (\neg C \vee D)$

Some examples that are **not** CNF:

- $\neg(A \vee B)$, where there is a NOT in front of a clause
- $\neg(A \vee B) \wedge (C)$, where there is a NOT in front of a clause
- $(A \vee B \wedge C)$, where there is an AND within a clause
- $(A \wedge B) \vee (C)$, where it is in DNF form.⁵

⁴A Boolean algebra equation is one containing only and ($\wedge$), or ($\vee$), and not ($\neg$) operations (at a basic level).

⁵DNF (Disjunctive Normal Form) is a “disjunction of conjunctions,” i.e., clauses of ands combined by ors. Solving these is trivially easy, as we just need to find one clause where no element evaluates to false.

1.2 Our Environment

Before we begin, we should talk about the environment that this book uses.

- Hardware: an Apple M1 MacBook Air, running MacOS Sequoia 15.7.8 with 8GB RAM and 512GB of storage space.
- Programming Language: C++20 compiled by Apple clang version 17.0.0 (clang-1700.6.4.2) targetted for arm64-apple-darwin24.6.0
- IDE: Visual Studio Code version 1.132.1
- Testing and Benchmarking: Testing and benchmarking is done by Test++, a C++ unit testing program I wrote (and is open source in case you want to try it out).⁶ At the time of writing, it does not optimize the compiled code. It is also not technically a benchmarking program. The way tests are set up is that each category of test/testing-set is given its own Test++ test which is run to completion. The resulting time to complete that Test++ test is used as the benchmark time. Differences in benchmarking frameworks and the way tests are set up may affect the result, but our testing workflow is as follows:

1. Create a Test++ testing function
2. Read in the testing file(s) from the given file path/directory
3. Create an instance of our SAT solver
4. Solve the equation
5. Check that the result is correct
6. Repeat steps 3-5 until no more testing files need to be processed

1.3 CNF Encoding

Now that we know what our inputs look like, the next question (I hope you're asking questions) is "how do we encode that?" Well good thing you're reading this, because that's what we're here to talk about! In this section, we'll expand our base of boolean algebra, including some definitions, transformations, and finally good encoding.

Definitions We'll Use

Clause A clause is a collection of variables, in which each variable is separated by an or ($\vee$), and each clause is separated by an and ($\wedge$).

Variable A variable is an element of a clause, taking on a value of True, False, or Unassigned. Each variable can be negated ($\neg$) or not, affecting its evaluated value.

⁶Not a shameless plug because it's awesome.

Literal A literal is an evaluated variable, evaluating to one of True, False, or Unassigned. Since a literal is an evaluated variable, it cannot be negated. A literal whose corresponding variable is unassigned evaluates to Unassigned, regardless of whether or not the variable is negated.

A quick example of the difference of a literal value, using all three of our definitions:

$$(\neg A)$$

is a clause of one variable “A”, which is negated. If the variable $A = \text{False}$, then its literal value is *True*, because

$$\neg \text{False} \equiv \text{True}.$$

In other words, A was assigned the value False, and then it evaluated to True. This will be very important in the future, so please take the time to understand it well enough.

As discussed earlier, a CNF equation only has three operators: and ($\wedge$), or ($\vee$), and not ($\neg$), but in propositional logic, there are many more operators, and their clauses are not guaranteed to be in the format we want. There are many algorithms to transform any propositional logic equation into CNF, but frankly, that’s beyond the scope of this book (it’s also beyond my scope). Instead, we will just look at how we transform each operator into an equivalent format using just our three operators.

Implication ($\rightarrow$)

p	q	$\neg p$	$p \rightarrow q$	$\neg p \vee q$
T	T	F	T	T
T	F	F	F	F
F	T	T	T	T
F	F	T	T	T

Xor ($\oplus$)

p	q	$\neg p$	$\neg q$	$p \oplus q$	$p \vee q$	$\neg p \vee \neg q$	$(p \vee q) \wedge (\neg p \vee \neg q)$
T	T	F	F	F	T	F	F
T	F	F	T	T	T	T	T
F	T	T	F	T	T	T	T
F	F	T	T	F	F	T	F

Biconditional ($\leftrightarrow$)

p	q	$p \leftrightarrow q$	$p \rightarrow q$	$q \rightarrow p$	$(\neg p \vee q) \wedge (\neg q \vee p)$
T	T	T	T	T	T
T	F	F	F	T	F
F	T	F	T	F	F
F	F	T	T	T	T

We won't look at negations of operators such as nand, nor, xnor, or any others. I'm not a boolean algebra expert when it comes to what operators exist. But if you were given an equation with these extra operators, you could easily apply the correct equivalency, then apply an algorithm that converts any boolean algebra equation into one of CNF form.⁷

DIMACS CNF Format

Now that we know what CNF looks like, we can now discuss how it's formatted for a SAT solver's consumption. Typically, CNF is formatted in *DIMACS CNF* form.

Some rules for DIMACS CNF form are:

1. The file may begin with comment lines. The first character of each comment line must be a lower case letter "c". Comment lines typically occur in one section at the beginning of the file, but are allowed to appear throughout the file.
2. The comment lines are followed by the "problem" line. This begins with a lower case "p" followed by a space, followed by the problem type, which for CNF files is "cnf", followed by the number of variables followed by the number of clauses.
3. The remainder of the file contains lines defining the clauses, one by one.
4. A clause is defined by listing the index of each positive literal, and the negative index of each negative literal. Indices are 1-based, and for obvious reasons the index 0 is not allowed.
5. The definition of a clause may extend beyond a single line of text.
6. The definition of a clause is terminated by a final value of "0".
7. The file terminates after the last clause is defined.

There are some things about this syntax to be wary of:

1. The definition of the next clause normally begins on a new line, but may follow, on the same line, the "0" that marks the end of the previous clause.
2. In some examples of CNF files, the definition of the last clause is not terminated by a final "0".
3. In some examples of CNF files, the rule that the variables are numbered from 1 to N is not followed. The file might declare that there are 10 variables, for instance, but allow them to be numbered 2 through 11.
4. Comments can appear anywhere in the file, but the comment must be on its own line and it must start with "c".⁸

For a quick example, here is a simple way to format a DIMACS CNF equation:

⁷I honestly may be talking out of my ass here. I will do more research on this subject at a later date.

⁸<https://jix.github.io/varisat/manual/0.2.0/formats/dimacs.html>

```
p cnf 4 2
1 2 -3 0
-1 4 0
```

And one with comments:

```
c this is a comment
p cnf 4 2
c this is another comment
1 2 -3 0
-1 4 0
```

In order to increase the compatibility of our SAT solver, we will be bending some rules.

We will firstly allow missing header lines. In the case where no header is passed in, we will instead infer the number of clauses and variables. We will also be very lenient in where comments can appear, but we will enforce that variables must be between 1 and N .

1. Comments can appear anywhere in the file, but they **MUST** be on their own line and start with “c”.
2. A header is not necessary, and if one is not provided, then the number of variables and clauses shall be inferred. However, if one is provided, then the number of variables and clauses **MUST** match the header.
3. All clauses, with the exception of the last clause, must end with “0”, and can extend multiple lines.
4. All variables are numbered 1 through N , where N is the number of variables. That is, if you say there are 10 variables, they must be labeled 1, 2, ..., 10 and not 2, 3, ..., 11 or any other convention.
5. (For compatibility) Upon encountering a “%” character outside of a comment, we will stop parsing immediately.

Thus, valid input can take on many forms:

```
p cnf 4 2
1 2 -3 0
-1 4 0
```

```
4 2
1 2 -3 0
-1 4 0
```

```
p cnf 4 2
c this comment doesn't have to be at the top
1 2 -3 0
-1 4 0
```

```
p cnf 4 2
c this comment doesn't have to be at the top
1 2 -3 0
-1 4
c the first two numbers are still the number
c of variables and the number of clauses
```

```
p cnf 4 2 1 2 -3 0
-1 4 0
```

More information can be found here.⁹

1.4 Reading a CNF Equation Programmatically

Now that we know what our input looks like, we can finally discuss reading it programmatically. Firstly we must discuss how our SAT solver will be represented in code. At the moment, we just need to keep track of

- the number of variables
- the number of clauses
- each clause in our equation
- whether or not we've tried to solve the equation yet
- what our solution is

We also want some methods to parse the equation, construct a SAT solver, solve the equation, print out the solution if it is satisfiable, and whether or not there's a solution. Therefore, our object will look like so:

```
1 public interface:
2     SATSolver(string input)
3     bool solveCNF()
4
5 private interface:
6     void parse(string equation)
7     int numVars
8     int numClauses
9     int[] [] clauses
10    optional<bool>[] vars
```

Our pseudocode is as follows:

⁹<https://people.sc.fsu.edu/~jburkardt/data/cnf/cnf.html>

```
1 SATSolver(string input)
2 {
3     vars = no-value
4
5     if (input is a file)
6     {
7         string content = readFile(input)
8         parse(content)
9     }
10    else
11    {
12        parse(input)
13    }
14 }
```

```
1 parse(equation)
2 {
3     // The method of parsing if there's a header
4     if (contains(equation, "p cnf"))
5     {
6         int[1..N] Arr
7
8         for each (line in input)
9         {
10            if (startsWith(line, 'c'))
11            {
12                skip to next line
13            }
14            else
15            {
16                for each (token in line)
17                {
18                    append(Arr, token)
19                }
20            }
21        }
22
23        numVars = Arr[1]
24        numClauses = Arr[2]
25
26        idx = 2
27
28        for (i = 1 up to numClauses)
29        {
30            int[1..N] clause
31
32            while (Arr[idx] != 0)
33            {
```

```

34         append(clause, Arr[idx])
35         idx = idx + 1
36     }
37
38     append(clauses, clause)
39 }
40
41
42 // Method of parsing when there's no header
43 else
44 {
45     numVars = 0
46     numClauses = 0
47
48     int[1..N] Arr
49
50     for each (line in input)
51     {
52         if (startsWith(line, 'c'))
53         {
54             skip to next line
55         }
56         else
57         {
58             for each (token in line)
59             {
60                 numVars = max(numVars, token)
61
62                 if (token == 0)
63                 {
64                     numClauses = numClauses + 1
65                 }
66
67                 append(Arr, token)
68             }
69
70             idx = 0
71
72             for (i = 1 up to numClauses)
73             {
74                 int[1..N] clause
75
76                 while (Arr[idx] != 0)
77                 {
78                     append(clause, token)
79                     idx = idx + 1
80                 }
81
82                 append(clauses, clause)
83             }

```

```
84         }  
85     }  
86 }  
87 }
```

This code will be implemented slightly differently in the example code, as I am using C++. There is also a big bug, which is that the code assumes that because “p cnf” appears in the file, there is a header. This is not always true because it does not consider whether or not the header is on the same line as a comment, so this will break it:

c	p	cnf
1	2	-3 0
-1	4	0

Because it assumes that there is a header when in fact there is not. I will leave fixing this bug as an exercise to both the reader and the writer, however, I will not be implementing this fix because if you’re purposely trying to break the constructor, you’re not trying to solve a CNF equation anyways. In case you’re wondering, the fix is to first find “p cnf” then scan backwards until you reach either the start of the file, or the end of the previous line. If you don’t encounter a “c” then you know that this header is actually a header.

Chapter 2

Brute Force

Since the dawn of man, whenever encountered with a difficult problem, we have decided “eh, let’s do it in the hardest way possible because I can’t think of a better solution.” The history of brute forcing things goes all the way back to our ancestors discovering fire,¹ all the way to our current solution to climate change.²

The idea of brute forcing something is really as simple as it gets. We have n number of literals, and therefore 2^n possible assignments, and we will check them all until we have either found a solution, or until we’ve run out of assignments to check. Because of how simple the idea is, it would just be better to show code and explain that instead of explaining then showing code.

```
1  solveCNF()
2  {
3      bool[1..numVars] Assignment
4
5      for (i = 0 up to 1 << numVars)
6      {
7          for (j = 0 up to numVars - 1)
8          {
9              Assignment[j] = i & (1 << j)
10         }
11
12         bool solved = true
13
14         for each (clause in clauses)
15         {
16             bool satisfied_clause = false
17
18             for (j = 1 up to clause.size)
19             {
20                 int variable = clause[j]
```

¹This is probably not true.

²Which is doing nothing and hoping everything sorts itself out, which is really disappointing.


```

21         int v_idx = abs(variable)
22
23         if (variable > 0)
24         {
25             satisfied_clause = Assignment[v_idx]
26             || satisfied_clause
27         }
28         else
29         {
30             satisfied_clause = !Assignment[v_idx]
31             || satisfied_clause
32         }
33     }
34
35     solved &= satisfied_clause
36 }
37
38 if (solved == true)
39 {
40     vars = Assignment
41     return true
42 }
43 }
44
45 return false
46 }

```

Unfortunately, the formatting across the page kind of sucks, but let's try to explain it anyways.

First, we create an array to hold each boolean value, and then next we start our loop from 0 all the way to “ $1 \ll \text{numVars}$.” If you are unfamiliar with what the left bit shift operator ($\ll$) is, essentially, it is a fast way of computing some number times 2 times the number of places to shift. That is all the explanation I am going to give because that's not the point. But by computing “ $1 \ll \text{numVars}$ ”, we are computing $1 \rightarrow 2^n$, which happens to be the total number of possible variable combinations (from all False to all True).

Then in the first inner for loop, we put the current variable combination into the ‘Assignment’ array. The index variable ‘i’ holds the current variable combination, and we isolate the j th variable using “ $1 \ll j$ ” and put it into the j th spot in the “Assignment” array.

Next, we evaluate the CNF equation against the current assignment. We first assume that the equation is solved (innocent until proven guilty), and then we evaluate each clause. We initially assume the clause is false because logically, if one literal is True, then OR-ing it with False makes the “clause satisfied” condition true (we could exit out of the evaluation loop, but it's easier to read if we don't). Then, we “and” the “clause satisfied” evaluation with the “equation solved” condition.

At the end of the evaluation loop, we check if we found a satisfiable assignment. If we have, we then make our *vars* field, which holds the solution if there is one, the

assignment, and then we return True to indicate that we have found a satisfiable assignment. If each possible assignment is not satisfiable, then we return False, and *vars* will hold nothing.

Since it's trivially easy to see why the code works, we will move onto something more interesting: optimizing brute force.

2.1 Optimizing Brute Force

I'll be completely honest with you, if you want to skip this section, go ahead. It's more or less just something interesting enough we can do to make the (arguably) worst algorithm³ better.

Since we're dealing with pseudocode, we're not going to consider the hardware side of computers such as branch prediction, cache locality, or anything that the hardware uses to execute our code (other than the CPU). We will consider it from a pure Big-O perspective (excluding constant factors unless we're choosing between two algorithms with the same computational complexity).

The first thing we will add is short circuiting to our clause evaluation. After line 26 where we finish OR-ing the "clause satisfied" boolean, we will add in a new block:

```
1  ...
2  bool satisfied_clause = false;
3
4  for (j = 1 up to clause.size)
5  {
6      int variable = clause[j]
7      int v_idx = abs(variable)
8
9      if (variable > 0)
10     {
11         satisfied_clause =
12             Assignment[v_idx] || satisfied_clause
13     }
14     else
15     {
16         satisfied_clause =
17             !Assignment[v_idx] || satisfied_clause
18     }
19
20     if (satisfied_clause == True)
21     {
22         continue
23     }
24 }
25
26 ...
```

³I do not know of any algorithm worse than brute forcing.

This just says “once this clause is satisfied, go evaluate the next one immediately.”

We can then add another conditional after this loop, which will move onto the next variable assignment the instant the equation is no longer satisfied:

```

1  ...
2  bool solved = true;
3
4  for each (clause in clauses)
5  {
6      bool satisfied_clause = false;
7
8      for (j = 1 up to clause.size)
9      {
10         ...
11     }
12
13     if (satisfied_clause == True)
14     {
15         continue
16     }
17
18     solved &= satisfied_clause;
19
20     if (solved == False)
21     {
22         break
23     }
24 }
25
26 ...

```

The next algorithm optimization we can make is removing the loop where we move the current assignment into the “Assignment” array. Rather than doing it upon every iteration, why don’t we do it once we find a satisfiable assignment?

```

1  solveCNF()
2  {
3      bool[1..numVars] Assignment
4
5      for (i = 0 up to 1 << numVars)
6      {
7          bool solved = true
8
9          for each (clause in clauses)
10         {
11             bool satisfied_clause = false
12
13             for (j = 1 up to clause.size)
14             {

```

```
15         int variable = clause[j]
16         int v_idx = abs(variable)
17
18         if (variable > 0)
19         {
20             satisfied_clause = (i & (1 << j))
21             || satisfied_clause
22         }
23         else
24         {
25             satisfied_clause = !(i & (1 << j))
26             || satisfied_clause
27         }
28
29         if (satisfied_clause == True)
30         {
31             continue
32         }
33     }
34
35     solved &= satisfied_clause
36
37     if (solved == False)
38     {
39         break
40     }
41 }
42
43 if (solved == True)
44 {
45     for (j = 0 up to numVars - 1)
46     {
47         Assignment[j] = i & (1 << j)
48     }
49
50     vars = Assignment
51     return true
52 }
53 }
54
55 return false
56 }
```

With that, that's all the optimization I can think of. In case you're wondering, no this does not change the asymptotic runtime of the method, however, it's easy to see why these changes would presumably make the algorithm run faster as we're adding code that stops evaluation the moment we've figured just enough out to know if it's worth continuing evaluation and deferring extra calculation (recording the current assignment) until once we know it's satisfiable.

As an exercise to the reader, I will ask you, how can you parallelize this algorithm? And once you've figured that out, can you program it in a way where once one thread finds a satisfactory assignment, it terminates all other threads immediately instead of waiting for the others to finish?

2.2 Brute Force - Results

To see how well our current algorithm does, running some tests would be a good idea. Why run tests when we have already proven that our algorithm works "perfectly?" To see how quickly it runs of course.

To test its speed (and as we add more features, to ensure correctness is kept), we use a database of problems from here:

<https://www.cs.ubc.ca/~hoos/SATLIB/benchm.html>

We gave our current implementation 1000 SAT instances (uf20-91), each with 20 variables and 91 clauses each, all satisfiable. It took this baseline brute force algorithm **2,119,301 milliseconds**, which is roughly 35.32 minutes. This result is actually from the first time I wrote a brute force algorithm. The brute force result you can find on GitHub, it actually took **1,770,851 milliseconds**, which is about 29.51 minutes. I have no idea what I did in the rewrite that shaved off 4 minutes. I won't even lie to you reader, I am a busy undergraduate student, so I didn't even do the "best practice" for testing, which would be to quit all other processes and then benchmark it. However, while I was testing, I was doing a Canvas assignment, so I was doing lots of background work.

Normally in science you have to run an experiment at least three times for data accuracy. But you must understand that I don't want to give up an hour and a half of my life to an algorithm that we're going to be improving upon anyways.

Here are the results for all brute force algorithms, including the short-circuited and parallel versions:

On all 1000 equations in uf20-91:	
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).

For the parallel version, the implementation isn't perfect. I believe that we don't prevent false sharing amongst cache lines (ie we get more cache swaps between threads leading to more cache misses), which could explain why on my M1 MacBook Air, we don't get close to a 8 times speedup. Since this was for fun, I'm not going to perfect the parallelized version, that can an exercise to the reader :)

Chapter 3

DPLL

The Davis-Putnam-Logemann-Loveland (DPLL) algorithm is a complete backtracking based search algorithm for finding satisfiable assignments for SAT problems. It is, dare I say, the foundational algorithm for modern SAT solvers, as the state-of-the-art algorithm used by basically every SAT solver builds upon the foundations that DPLL provides, as we will see later on.

It consists of three components:

- Recursive backtracking
- Unit propagation
- Pure literal elimination

How does it work?

Imagine a binary tree of variables, each with two branches representing their assignment, one for true and one for false. Suppose we don't know anything about our CNF equation. That is, from looking at it, any variable could have any value, and we don't know how each assignment helps us in finding out whether or not the equation is satisfiable. How can we figure out if we can satisfy the equation without just brute forcing it?

The answer is ~~making an ass out of you and me~~ to make assumptions.




Figure 3.1: Example DPLL decision tree (the outcomes may not correspond to an actual CNF instance, sue me).

Going back to our hypothetical, we don't know anything about anything. We could have a giant decision tree to explore with an exponential number of leaves to evaluate. How does making assumptions help us? Do we just run a random number generator and if it returns 1, we go "okay, it's satisfiable," and false otherwise?

No. We make assumptions about our literal assignments. Unlike brute forcing where we try every possible combination of variable assignments at once, in DPLL we start with a literal, we'll call it A . Suppose there is nothing in the equation that indicates whether or not A should be true or false; either could be valid.

So to tackle our problem, initially we will assume A is true and see where that leads us. We propagate that value throughout our clauses, and suppose there is a clause that can't be satisfied. That's a contradiction, because if the equation was satisfiable then $A = \text{True}$ couldn't have led us there.

So we've figured out that at the very least $A = \text{True}$ leads to a contradiction. We've now eliminated half of the search space immediately and only need to consider the cases where $A = \text{False}$.

This may not make total sense immediately, so instead of giving one explanation of the algorithm and then building it, we will build up the algorithm and explain how each part works.

1
2

3.1 Recursive Backtracking

First, we need to rewrite our brute force algorithm to be recursive, as that's the foundation for the entire algorithm (duh). To do this, we will introduce a new method `solve()` which will be our recursive solving method. To also make things simpler, we will extract the equation evaluation into its own separate method as well.

This means our SAT solver becomes like so:

```
1 public interface:
```

¹Note for future me: write a better introduction.

²https://en.wikipedia.org/wiki/DPLL_algorithm

```

2     SATSolver(string input)
3     bool solveCNF()
4
5     private interface:
6         void parse(string equation)
7         bool solve(bool[] assignment)
8         bool evaluate(bool[] assignment)
9         int numVars
10        int numClauses
11        int[] [] clauses
12        optional bool[] lits

```

The flow essentially becomes like this: *solveCNF()* calls *solve()*, which calls itself over and over and returns whatever solution it finds back up to *solveCNF()*, who will handle any sort of cleanup.

Why can't we just make *solveCNF()* call itself recursively? Good question. That's because *solveCNF()* will handle any setup that *solve()* will need in order to function properly, and it will handle any cleanup that *solve()* may create.

Let's now discuss how *solve()* will work. It will be passed in a *bool[]* for the current assignment of literals. When the size of our assignment is equal to the number of variables, we will then evaluate the satisfiability of the assignment, and if it's satisfiable we return true. Otherwise, we return false, backtrack, retry, until we either find a satisfiable assignment or we've exhausted them all.

The pseudocode is as so, starting with *evaluate()*:

```

1     bool evaluate(bool[] assignment)
2     {
3         bool solved = true;
4
5         for each (clause in clauses)
6         {
7             bool satisfied_clause = false;
8
9             for (j = 1 up to clause.size)
10            {
11                int variable = clause[j];
12                int v_idx = abs(variable);
13
14                if (variable > 0)
15                {
16                    satisfied_clause = assignment[v_idx]
17                        || satisfied_clause;
18                }
19                else
20                {
21                    satisfied_clause = !assignment[v_idx]
22                        || satisfied_clause;
23                }
24            }

```



```

25         if (satisfied_clause == True)
26         {
27             continue;
28         }
29     }
30
31     solved &= satisfied_clause;
32
33     if (solved == False)
34     {
35         break;
36     }
37 }
38
39 if (solved == True)
40 {
41     lits = assignment;
42 }
43 else
44 {
45     lits = no-value;
46 }
47
48 return solved;
49 }

```

As you can see, all we've done is move our equation evaluation from brute force into its own function. One thing to note is that this method also handles updating our *lits* field. Thus, if *evaluate()* returns true, then *lits* will hold the satisfying assignment, and if it returns false, then *lits* will hold nothing.

Now our recursive *solve()* method:

```

1  bool solve(bool[] assignment)
2  {
3      if (assignment.size == numVars)
4      {
5          return evaluate(assignment);
6      }
7
8      append(assignment, True);
9
10     if (solve(assignment))
11     {
12         return true;
13     }
14
15     assignment.back() = False;
16
17     if (solve(assignment))

```

```

18     {
19         return true;
20     }
21
22     truncate(assignment, 1);
23     return false;
24 }

```

This might be slightly confusing at first, so let's go over it.

We first check our base case. If our assignment has all of its variables, we evaluate it and return the result whether it be true or false. And you'll notice that whenever we return true from a call of *solve()*, we immediately return true afterwards, so therefore, if we find a satisfiable assignment, we will go up our call stack back to *solveCNF()*, so no need to worry that it will terminate if we find a satisfiable assignment.

However, if our assignment still needs some variables, we append true to our assignment, then we do a recursive call on *solve()*. If that recursive call leads to a satisfiable assignment, we return true, otherwise, we swap out our last value for false and try *solve()* again. And if that ends up working out, we return true, but if it doesn't then we remove the last value and return false.

So how do we know that it will guarantee termination for an unsatisfiable problem? Notice how when assigning the current variable true or false and it doesn't work, we just remove the last one and go up the call stack by one? Well suppose we're at the top of the call stack (we've assigned the first variable) and the first variable is assigned false, and that didn't lead to a satisfiable assignment. Then we remove that first variable from the assignment, and then we return false. So our assignment array is empty, and we've returned to *solveCNF()*. Thus, *solve()* will terminate eventually.

And how do we know that we'll try every single assignment? Well for the very last variable, it's easy to see, if it being true doesn't satisfy the equation, we assign it false, and if that doesn't work, we remove it and go up the call stack. Then for the previous variable, if it was true, we assign it false and go back to the last variable where it tries true and/or false again. Then that doesn't work, we go up the call stack twice, and that continues until we've found a satisfiable assignment or we've exhausted them all.³

Lastly, our *solveCNF()* method:

```

1  bool solveCNF()
2  {
3      bool[1..numVars] assignment;
4
5      return solve(assignment);
6  }

```

Because we're dealing with pseudocode, we don't have to do things such as reserving or resizing the assignment array. And since *evaluate()* takes care of making sure our *lits* field holds a valid assignment (if found) and holds nothing if not, we don't need to do anything else in *solveCNF()*.

³This is a very informal proof, but what's most important is that we build up our understanding of what's going on instead of dumping notation. Perhaps I will do a formal proof in the future.

You may wonder though, why not pass in *lits* to *solve()* or have no setup and make *solveCNF()* purely recursive? Those are very valid questions to ask.

This has to do with the differences between pseudocode and actual implementation. In the actual implementation, some methods depend on the state of *lits*, such as whether or not it holds anything, and by passing it into *solve()*, we force it to hold something, even if there is no satisfiable assignment. And it's true, that in *solveCNF()* if we had no setup, we could make it recursive on *lits*, however, I chose to implement it differently.

When you are reading this, please ask yourself these kinds of questions, because oftentimes, the way you read my implementation is not the best way, and even if it is, there is no reason as to why you can't think of different implementations, try them out yourself, and compare. Perhaps, and I expect this will happen often, you will find a better implementation than the one written down here.

3.1.1 Recursive Backtracking – Results

Running this code on the same 1000 instances (uf20-91), it took **1,846,644ms**, which is roughly 30.78 minutes. This is actually a pretty big improvement from our 35.32 minutes via brute forcing. Although this seems like a big gain in speed, it should be noted that at this point the algorithm is still just a fancier brute forcing algorithm. The next time I rewrote this algorithm, it took **1,540,613ms**, which is about 25.68 minutes. So once again, that's about a 5 minute speedup just from rewriting it. Perhaps we should rewrite every piece of software in the world to make it run 5 minutes faster.

Here's a question for you, why might assigning our variables true instead of false first lead to an increase in speed? **Hint: Look at the number of negations present in the testing set.**

On all 1000 equations in uf20-91:

Recursive Backtracking:	1,540,613ms (~25.68 minutes).
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).

3.2 Unit Propagation

Unit propagation is the next step in implementing DPLL, and it introduces the idea of implication. Imagine a clause with two variables: *A* and *Z*, and we know that *A = False*. What can we say about *Z*? Well, if our equation is satisfiable, then the clause must be satisfiable, therefore *Z = True* because if it didn't, that would be a contradiction.

So instead of assigning variables *B*, *C*, ..., *Y*, then assigning *Z*, why don't we just do it now, and use that knowledge in other clauses containing *Z* or $\neg Z$?

That is essentially the idea behind unit propagation: the decisions we make imply other assignments, and so we should force them now instead of waiting until later.

However, this introduces a light hiccup into our implementation. You see, our initial invariant about assigning variables is that they are assigned in numerical order instead of (possibly) out of order.

“Not a problem! We’ll just make the assignment array have space for every variable so we can index it anywhere,” I hear you think. Unfortunately that interferes with our base case because it relies on the size of the assignment array. If it is initially at a fixed size of *numVars*, then our base case will always fire because our assignment array always has a size of *numVars*. So we would either need to change our base case or change how we store learned variables.

What about a learned set? One that holds whether or not we’ve learned a variable (positive or negative), and then as we traverse down the call stack, we just check if a variable is already learned, and if so, we give it its learned value?

That would 100% work. In fact, the first time I implemented DPLL, that is exactly what I used. However, it has the disadvantage of being annoying to maintain. Instead, why don’t we create a new data type called something like *var* that represents the state of each variable: *True*, *False*, and *Unassigned*.

Then we can just maintain a simple counter of how many are not *Unassigned*, and then that keeps our base case nice and neat. Yes dear reader, we are going to do that, however, if you’d like to use the learned set, reading about how we do it without a learned set will help you and offer clues on how to implement it your way.

3.2.1 Changes To Our Implementation

```

1      public interface:
2          SATSolver(string input)
3          bool solveCNF()
4
5      private interface:
6          enum var
7          {
8              False
9              True
10             Unassigned
11         }
12
13         void parse(string equation)
14         bool solve(var[] assignment, int level)
15         bool evaluate(var[] assignment)
16         void propagate()
17
18         int numVars
19         int numClauses
20         int numAssigned
21
22         int[] [] clauses
23
24         optional var[] vars

```

Here we just add in a new enum for the states that a variable can be, and adjust the methods using *bools* to use the new enum type instead.

3.2.2 “Textbook” Unit Propagation

Before we start changing our algorithm, let’s implement the *propagate()* method.

As stated before, the idea is that given a clause with either only one variable or only one *Unassigned* variable (and the rest evaluate to *False*), we can immediately say that *Unassigned* variable will be whatever value it needs to be to evaluate to *True* for that clause.

In the “textbook” way of implementation,⁴ once a clause is satisfied because of a literal, you remove every other clause that is satisfied by that literal. And for every other clause containing the negation of the literal, you remove that negation entirely. This is done to shrink the search space.

We will implement this algorithm with as much malicious compliance as you can carry at an office job without getting fired:

```

1      bool solveCNF() {
2          var[1..numVars] assignment
3
4          for each (var in assignment) {
5              var = Unassigned
6          }
7
8          propagate(assignment)
9          solve(assignment, 0)
10     }
```

We want to propagate before we start solving in case there are clauses with only one variable (it’s already unit).

```

1      bool solve(bool[1..numVars] Assignment, int level) {
2          if (level == numVars - 1) {
3              return evaluate(Assignment)
4          }
5
6          bool[1..N] backupAssignment = assignment
7          int[][] backupClauses = clauses
8
9          if (Assignment[level] == Unassigned) {
10             Assignment[level] = True
11             propagate(Assignment)
12         }
13     }
```

⁴https://en.wikipedia.org/wiki/Unit_propagation

```

14         if (solve(Assignment, level + 1)) {
15             return True
16         }
17
18         clauses = backupClauses
19         Assignment = backupAssignment
20
21         if (Assignment[level] == Unassigned) {
22             Assignment[level] = False
23             propagate(Assignment)
24         }
25
26         if (solve(Assignment, level + 1)) {
27             return True
28         }
29
30         return False
31     }

```

Before we get onto the code for unit propagation, let's understand this code first. Because unit propagation no longer guarantees that variables are assigned in order, we need to pass in a "level" as a parameter so we know which variable we need to assign. We also now need to check that propagation didn't already assign this variable. If it's currently Unassigned, we need to assign it then propagate, but if it's been assigned before we assign it, then we know that it's the result of propagation, so we can continue onto the next level. Other than that, it's practically the same as our recursive backtracking code, except now we need to have a "backup" of our clauses and our assignment.

Why? Because *propagate()* **mutates** the clauses and the assignment in a way that breaks the assumption that our literals are assigned in "numeric" order. Since it returns nothing, we don't have a way to undo that easily, so instead, for readability, we just have a backup copy and whenever propagation leads to an unsatisfying assignment, we return false and undo whatever changes we made by copying back into our solver's fields. And the reason why we don't need to undo the propagation before returning False is because when we move back up the call stack by one level, the method continues by undoing the effects for us. Now let's look at how we perform textbook unit propagation:

```

1     void propagate(int[1..N] Assignment) {
2         while (true) {
3             Set<int> unitClauses;
4             Set<int> unitVars;
5
6             for (i = 0 up to numClauses) {
7                 clause = clauses[i]
8
9                 bool isUnit = True
10                int unassignedIdx = -1
11                int numUnassigned = 0
12

```

```

13         for (j = 0 up to clause.size) {
14             int var = clause[j]
15             int idx = abs(var)
16
17             if (Assignment[idx] == Unassigned) {
18                 numUnassigned = numUnassigned + 1
19                 unassignedIdx = j
20             } else if (var > 0 && clause[j] == True
21                 || var < 0 && clause[j] == False)
22             {
23                 isUnit = false
24             }
25         }
26
27         isUnit = isUnit && numUnassigned == 1
28
29         if (isUnit) {
30             add(unitClauses, i)
31             add(unitVars, clause[unassignedIdx])
32
33             int var = clause[unassignedIdx]
34             int idx = abs(var)
35
36             if (var > 0) {
37                 assignment[idx] = True
38             } else {
39                 assignment[idx] = False
40             }
41         }
42     }
43
44     int clauseIdx = 0;
45     for each (clause in clauses) {
46         for each (var in clause) {
47             if (contains(unitVar, -var)) {
48                 remove(clause, -var)
49             }
50
51             if (contains(unitVar, var)
52                 && contains(unitClause, clauseIdx))
53             {
54                 remove(clauses, clause)
55             }
56         }
57
58         clauseIdx = clauseIdx + 1
59     }
60
61     if (unitClauses.size == 0 && unitVars.size == 0) {
62         break

```

```

63         }
64     }
65 }

```

Okay, now let's understand what's happening here. We immediately go into an infinite loop (nice), and we start looking at each clause for two things:

1. The clause has only one Unassigned variable
2. Every other literal in the clause is False (remember, literals are evaluated variables)

If both things are true, then the Unassigned variable will be assigned whatever value is required to make it evaluate to True within that clause, and we know we have found a unit variable and a unit clause. After we have found all of our unit variables and unit clauses, we then rescan every other clause and delete the each unit variable's negative, (if we learn that -5 is unit, then we remove $+5$), and if the clause contains a unit variable, we just delete that clause entirely from the database. Once we stop finding new unit variables and new unit clauses, we stop entirely. You may wonder (I reread my code and thought I found a bug), "how do we know that we always terminate?" That's a great question. You see, this confusion may come from the fact that we rescan the database every iteration of the *while()* loop searching for unit clauses. What is stopping us from re-finding the same clause as unit upon a new iteration? The answer is much simpler than you'd expect: since we assign our unit variables, the clauses upon the next iteration no longer follow our criteria for being unit propagate-able because there is now a literal in that clause that is True. Therefore, the number of unit variables and clauses monotonically decrease between iterations (a fancy way of saying the number can only go down).

Now, you may be wondering, "what's the performance of our solver now?" Let me answer that as plainly as possible: worse than brute force. How? Easy! We are doing SO MANY things that hurt the computer. We are allocating TONS of memory in order to backtrack, we are allocating even more memory to propagate (hash sets aren't cheap), and hashing, while $O(1)$, is not a cheap operation to do. So yeah, I can believe that it's worse than brute force because at least the brute force algorithm was easy on the CPU. Let's learn how this is *usually*⁵ implemented!

3.3 Better Unit Propagation

So the textbook style unit propagation sucked. Let's change that. As discussed before, the biggest bottleneck/inefficiency in the actual implementation is the number of copies we need to make of our clauses and assignments on each decision level. Assigning one variable creates a new decision level, so 20 variables means 20 levels. For an unsatisfiable CNF equation, I estimate we need to make about 3 copies of our clauses

⁵I wanted to say "actually implemented", however, that's a big generalization, and again, since I'm not the expert here, I don't think I have the authority to say that. The next section is how I implemented it, and it roughly follows the idea on how SAT solvers implement this.

and assignments: one for the initial storage, another for the first failure, and a third for the second. Even if we ignore the inefficiencies in how we propagate (memory allocations, traversals, and adjustments our arrays would need to make when we delete things), this is a huge amount of memory to store. If you think that 20 variables and 91 clauses is a lot, modern SAT solvers can handle equations with **millions** of variables, so creating an algorithm that makes brute force look attractive is actually an achievement.⁶

So what's the strategy? Firstly, it would be nice if we didn't actually mutate our clauses at all. Ie, no clause deletion, and no variable deletion. Next, it would be nice if we didn't need hash sets in order to store our unit variables. Instead, it would be nice if we had a data structure where we put in our decided variables, and then it slowly shrinks as we iterate through. Ie, we propagate one on variable at a time and never again. This sounds like we need a *queue*.⁷

Lastly, instead of copying our assignment, and constantly reiterating through the *entire* clause database, what if we just marked which variables we assigned, and which clauses we've satisfied? That way when we need to undo, we just mark those variables as Unassigned, and those clauses as no longer satisfied. That way we save a lot of memory (and besides, copying is NOT $O(1)$ as you might be led to believe by that deceptively simply "=" operator).

Let's change our interface:

```

1      public interface:
2          SATSolver(string input)
3          bool solveCNF()
4
5      private interface:
6          enum var
7          {
8              False
9              True
10             Unassigned
11          }
12
13          void parse(string equation)
14          bool solve(var[] assignment, int level)
15          bool evaluate(var[] assignment)
16          (int[] vars, int[] clauses) propagate()
17
18          int numVars
19          int numClauses
20          int numAssigned
21
22          int[][] clauses
23          bool[] satisfied
24
```

⁶They might revoke my degree for that actually

⁷That wasn't shoehorned in at all!

```
25 optional var[] vars
```

We need to make some minimal changes to *solveCNF()*:

```
1 bool solveCNF() {
2     var[1..numVars] assignment
3
4     for each (var in assignment) {
5         var = Unassigned
6     }
7
8     propagate(assignment, 0)
9     solve(assignment, 0)
10 }
```

Since *propagate()* now takes a decision level/variable index, we need to pass one to the method. Although no variable has an index of 0 (because it's reserved to mark the end of a clause), passing in 0 just tells it “Hey, there isn't a literal to propagate on, so just find all clauses with only one variables and mark them as satisfied and ‘learn’ those variables.” And since this is at the root level, we don't need to undo anything if it turns out that the equation is unsatisfiable, since we're gathering things that HAVE to be a certain value if this equation is to be satisfiable.

Since we don't need to change *solveCNF()*, we can go straight to *solve()*:

```
1 bool solve(int[1..N] Assignment, int level) {
2     if (level == numVars - 1) {
3         return evaluate(Assignment);
4     }
5
6     if (Assignment[level] != Unassigned) {
7         return solve(Assignment, level + 1)
8     }
9
10    Assignment[level] = True
11    (int[] unitVars, int[] satisfiedClauses)
12        = propagate(Assignment, level)
13
14    if (solve(Assignment, level + 1) == True) {
15        return True
16    }
17
18    for each (var in unitVars) {
19        int idx = abs(var)
20        Assignment[idx] = Unassigned
21    }
22
23    for each (clauseIdx in satisfiedClauses) {
24        satisfied[clauseIdx] = False
25    }
```

```

26
27     Assignment[level] = False
28     (int new_unitVars, int new_satisfiedClauses)
29       = propagate(Assignment, level)
30
31     if (solve(Assignment, level + 1) == True) {
32         return True
33     }
34
35     for each (var in new_unitVars) {
36         int idx = abs(var)
37         Assignment[idx] = Unassigned
38     }
39
40     for each (clauseIdx in new_satisfiedClauses) {
41         satisfied[clauseIdx] = False
42     }
43
44     Assignment[level] = Unassigned
45     return False
46 }

```

This method is practically the same, except now, we just check if we've already assigned the variable at the current decision level. If we have assigned a variable before we reach its decision level, then it must have been inferred/learned from the *propagate()* method, so we don't need to call it again. Then, it's business as usual: assume the current variable is True, propagate, if it leads to a satisfiable assignment, return True, otherwise, we undo what we've assigned and marked as learned, then repeat with False.

What's more interesting is our shiny new *propagate()* method:

```

1     (int[] vars, int[] clauses) propagate(int[] assignment, int level
2     ) {
3         queue<int> unit
4
5         int[] unitVars
6         int[] satisfiedClauses
7
8         unit.enqueue(level)
9
10        while (unit.size > 0) {
11            unit.dequeue()
12
13            for (i = 0 up to numClauses) {
14                if (satisfied[i]) {
15                    continue
16                }
17
18                clause = clauses[i]

```

```

19         bool isUnit = True
20         int unassignedIdx = -1
21         int numUnassigned = 0
22
23         for (j = 0 up to clause.size) {
24             int var = clause[j]
25             int idx = abs(var)
26
27             if (Assignment[idx] == Unassigned) {
28                 numUnassigned = numUnassigned + 1
29                 unassignedIdx = j
30             } else if (var > 0 && clause[j] == True
31                       || var < 0 && clause[j] == False)
32             {
33                 isUnit = false
34
35                 satisfied[i] = true
36                 append(satisfiedClauses, i)
37             }
38         }
39
40         isUnit = isUnit && numUnassigned == 1
41
42         if (isUnit) {
43             add(unitClauses, i)
44             add(unitVars, clause[unassignedIdx])
45
46             int var = clause[unassignedIdx]
47             int idx = abs(var)
48
49             if (var > 0) {
50                 assignment[idx] = True
51             } else {
52                 assignment[idx] = False
53             }
54
55             unit.enqueue(var)
56
57             append(unitVars, var)
58             satisfied[i] = true
59             append(satisfiedClauses, i)
60         }
61     }
62 }
63
64     return (unitVars, satisfiedClauses)
65 }

```

So, let's discuss what's going on. *propagate()* now takes in the decision level, which

is also the same thing as an index for a variable (conceptually speaking). It then adds that to a queue, then while the queue is not empty, it scans the clauses as before. First it makes sure we're not operating on a clause that's already been satisfied (because we can't tell anything from it), and when it finds a clause that has a true literal, it marks the clause as satisfied. Then when it finds a unit clause, it marks the clause as satisfied, keeps track of the new unit variable its found, and then adds that unit variable to the queue. Then once it's done, it returns the unit variables and the indices of the satisfied clauses.

You may have some questions about this approach. Firstly, how do we make sure we don't ever add the same variable twice? The answer is the same as before: because we assign a variable once it's found to be in a unit clause, it can't be found as Unassigned again, and therefore can't be marked as a unit variable again. Then, since we skip over already satisfied clauses, the amount of clauses we scan each iteration decreases monotonically.

And going back to our *solve()* code, even if we mark the same clause as satisfied or the same variable as unit twice, it doesn't matter because when we backtrack, we simply restore it back to its state of unsatisfied or Unassigned, respectively. What's important is that we don't mark the same clause or variable twice *across decision levels* because then, and only then, would our backtracking code become a problem because we wouldn't be correctly restoring the state of the equation back to what it was before the decision level was entered. But since we can't mark the same variable or clause twice **until** we backtrack, there's nothing to worry about.

Before we end this section, you may also be wondering: can we use the *satisfied* array in our *evaluate()* method to skip clauses we know we've already satisfied? And the answer is yes, and I would encourage you to do so in your implementations.

3.3.1 Better Unit Propagation - Results

So after implementing better unit propagation, this took **868ms** to solve all 1000 satisfiable equations. The first time I implemented this version of unit propagation, it actually took **2363ms**, so this is somehow over twice as fast. And because it's so fast, just like the first time I implemented this, we will be now testing our solver on unsatisfiable equations as well.

Our second testing set comes from the same database as before.⁸ And our next testing set is called "UUF50.218.1000." In other words, it's full of equations that have 50 variables, 218 clauses, and there are 1000 such equations in the set. In order to solve all 1000 equations, it took **140,411ms** (2.34 minutes). That's also much faster than the first time I implemented this algorithm. The first time I did, it took 646,228ms (10.77 minutes). Honestly, with this much speedup, I don't know if I was just bad at implementing these algorithms the first time,⁹ or if downgrading my operating system (and subsequently wiping my harddrive clean)¹⁰ just has that much of an effect on execution speed.¹¹

⁸<https://www.cs.ubc.ca/~hoos/SATLIB/benchm.html>, in case you needed a reminder.

⁹Highly likely

¹⁰I HATED MacOS Tahoe. Sequoia is where it's at.

¹¹Also likely, depending on how much RAM and swap is being used to begin with.

Why didn't we test our previous versions out on unsatisfiable equations? That's a good question. For brute forcers, since there is no satisfiable assignment, it needs to check every single one to conclude that it's unsatisfiable, which would take a lot of time. Likewise with our recursive backtracking solver. And since our (my) implementation of textbook unit propagation was *worse* than brute force,¹² obviously we wouldn't want to test it on something that will take more time than a satisfiable equation. Also because the asymptotic runtime is *still* exponential, having a much larger equation would take a much, much larger amount of time, even if they were satisfiable.

Results so far:

On all 1000 equations in uf20-91:

Unit Propagation:	868ms (~0.87 seconds).
Recursive Backtracking:	1,540,613ms (~25.68 minutes).
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).

On all 1000 equations in UUF50.218.1000:

Unit Propagation: 140,041ms (~2.34 minutes).

3.4 Pure Literal Elimination

Suppose I gave you the following CNF equation:

1	2	-3	0
1	-4	5	0
1	-2	5	0

Do you notice anything odd about it? Perhaps "odd" is not the correct word here, do you notice an easy way to satisfy this equation? You probably noticed that the variable 1 shows up in every clause and is always in the positive. So, if you just assign $1 = T$, then it doesn't matter what the rest of the variables are assigned, because $1 = T$ satisfies every clause! This is idea behind **Pure Literal Elimination**. Essentially, every now and then, we scan our equation for "pure literals", which are variables that only show up in the positive or negative (but never both), and then give them whatever assignment is necessary to have their literals evaluate to True.

Now, suppose I give you this equation:

1	0		
-1	-2	-3	0
1	2	0	
-2	3	0	

¹²A difficult feat.

I have a question for you dear reader: is -2 a pure literal? You might be inclined to say “no” because -2 and 2 are both present in the equation. However, despite that, the answer is “yes.” Why? Well notice in the first clause that 1 is a unit variable, so $1 = T$ is forced by pure literal elimination. Then the third clause, 120 , because it’s already satisfied, the variable 2 doesn’t matter anymore. The other two clauses, which are unsatisfied, are the ones that matter. They both contain -2 , and because they’re unsatisfied, we can force $2 = F$ to satisfy them, thus making the assignment of 3 irrelevant.

Okay, that makes sense, sort of. How come we’re allowed to do that? Like, how do we know that forcing $2 = F$ wouldn’t screw up the equation elsewhere (if it were larger)? And the answer is: because it doesn’t hurt the other clauses. In the third clause, once it’s satisfied, the only thing that can make it unsatisfied is the backtracking, but since 1 is unit, there shouldn’t be any backtracking that would suddenly make its satisfiability rest on the assignment of 2 . Likewise, for the other two clauses, no matter what value 3 takes, $2 = F$ must be true to satisfy the other clause (try it out for yourself!), so we can go straight ahead and assign $2 = F$ and skip deciding.

Okay, sure, but there are 8 total assignments that we can get. How do we know that forcing $2 = F$ would lead to a satisfiable assignment? As in, how do we know we’re not prematurely boxing ourselves in? That’s another great question! The reason why we’re not boxing ourselves in prematurely is because, again, it only helps the satisfiability of the whole equation, and because we’re not doing anything detrimental, we can make that decision.¹³

3.4.1 Implementing Pure Literal Elimination

Now that we (sort of) understand why it works, how do we implement it? We need some form of efficient tracking for things we’ve already encountered, perhaps something like a map or a set. Then, we just need to make sure we’ve only encountered one version of a literal, and if we have, then we know it’s pure. Then we just need to gather our pure literals and the satisfied clauses, and then return them.

```

1  (int[] var, int[] clauses) pureElim(int[] assignment) {
2      map<int, int[]> var_to_clause
3      int[] pure
4      int[] satisfied
5
6      for (i = 0 up to numClauses) {
7          if (satisfied[i]) {
8              continue
9          }
10
11         clause = clause[i]
12         for each (var in clause) {
13             append(var_to_clause[var], i)
14         }

```

¹³Essentially, think of it as making the solution “more correct” by forcing the pure literal.

```

15     }
16
17     for each ([var, clauses] in var_to_clause) {
18         idx = abs(var)
19
20         if (contains(var_to_clause, -var) or assignment[idx] !=
21             Unassigned) {
22             continue
23         }
24
25         if (var > 0) {
26             assignment[idx] = True
27         } else {
28             assignment[idx] = False
29         }
30
31         for (i = 0 up to clauses.size) {
32             satisfied[categories[i]] = true
33             append(satisfied, clauses[i])
34         }
35
36         append(pure, var)
37     }
38
39     return (pure, satisfied)

```

So this is a single pass version of pure literal elimination. It is possible that after finding some pure literals, and marking some clauses satisfied, a second pass could find new pure literals. So why are we implementing a single pass version? Besides it not being a technique used in modern SAT solvers, in my experience, which isn't much since I'm being honest, it doesn't make a big (positive) difference in solve time efficiency. And besides, if we wanted a multi-pass version, we could wrap it in *while(true)* loop and break once we find no more new pure literals.

Now let's get into how the algorithm works. We first build a map from variables to their clauses (in unsatisfied clauses). Next, afterwards, we then check for each element in the map, that the map does not contain the negative variable of the current entry, and if it doesn't, we then need to make sure that the current variable is Unassigned. If the current entry passes both those conditions (pure and Unassigned), we assign it so it evaluates to True. Then, we mark each clause the variable is found in as satisfied, and then we add the indices of the satisfied clauses into the return values.

3.4.2 Pure Literal Elimination - Results

So. I have some terrible news. After implementing pure literal elimination, our solving time *increased*. I know, I know. It's not supposed to happen, but it can happen. There are many possible explanations for this. For instance, pure literal elimination takes a decent amount of work to perform. We have to scan all unsatisfied clauses, and since

there is no way to determine if a literal is impure until after everything is scanned (we could have used a hash set, but then we'd need to hash for every literal), we're accruing memory and runtime costs. Second, even if pure literal elimination finds something, it usually isn't much (at least for the testing sets currently in use), so we're doing so much work for so little in return. Third, it increases our backtracking cost because now we need to undo unit propagation and the pure literal elimination. Fourth, my implementation isn't perfect. It is possible that forcing a pure literal may force another pure literal in a different clause, so you would need multiple passes to do. However, our implementation is just one pass (wrap it in a *while* loop or something).¹⁴

However, no one cares *why* something performed badly. People only care about *how* badly something performed. So here are the results: For uf20.91: 1548ms, almost a 2x increase from just unit propagation. For UUF50.218.1000: 231,107ms (3.85 minutes), about a 1.5 minute increase from just unit propagation.

On all 1000 equations in uf20-91:	
Pure Literal Elimination:	1548ms (~1.55 seconds).
Unit Propagation:	868ms (~0.87 seconds).
Recursive Backtracking:	1,540,613ms (~25.68 minutes).
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).
On all 1000 equations in UUF50.218.1000:	
Pure Literal Elimination:	231,107ms (~3.85 minutes).
Unit Propagation:	140,041ms (~2.34 minutes).

Pure Literal Elimination had about a **2x** increase (which is bad) in solving time when compared to just Unit Propagation on uf20.91, which is the test set of satisfiable problems. It also had about a **1.6x** increase in solving time on UUF50.218.1000, which is the set of unsatisfiable problems.

3.5 Watched Literals

Before we get into the Watched Literal algorithm, I want to ease us into it. Why? Because it's so different from our current version of unit propagation. For a moment (and for the rest of the book) let's forget about pure literal elimination; it wasn't that great anyways. Unit propagation gave us a massive speedup opposed to recursive backtracking, taking runtimes down from minutes to seconds and from half-hours to mere minutes. Despite all of that, it's really inefficient in the sense that we're going so much work for so little in return.

¹⁴Modern SAT solvers don't tend to use pure literal elimination for reasons outside the "it wasn't very helpful" aspect, and next chapter we'll see why.

Let's define what a **contradiction** is, because we've been using that word here and there, but beyond "unsatisfiable," we haven't formally defined what it means. A **contradiction** is the result of a literal forcing a clause to evaluate to False, leading to the whole assignment to become unsatisfiable. You might remember how in the brute force solver, we added an *iff()* statement at the end of evaluating a singular clause that checked if that clause evaluated to False. If it did, we skipped that assignment and moved onto the next one, but if it didn't, we continue our evaluation.

This led to a pretty significant speedup in the brute forcer because instead of evaluating everything then deciding, we did it incrementally, allowing us to quickly leave when something was bad. At the moment, our original idea for unit propagation follows that first idea of delaying evaluation to the end. However, if you go back and look at the pseudocode, you'll notice that we do something *almost* the same as checking if a clause evaluates to False, at the moment, we check if everything except for one element evaluates to False, and that one element is Unassigned (since we made it so that unit propagation only operates on unsatisfied clauses). However, notice that when a clause evaluates to False, rather than stopping our propagation and assignment, we continue until we reach the evaluation phase?

Just like in brute force, the idea is the same: once a clause evaluates to False, no matter where we are in our assignment, there's no point in continuing on because we've found out that nothing will unfalsify that clause. As an exercise, I want you to try adding that into your unit propagation code and seeing how fast it takes to run on our test sets. I would do it, but I have a massive headache from implementing this algorithm last night.¹⁵ Once you're done, we'll discuss watched literals some more.

So that early contradiction detection is one of the ideas behind watched literals. It's not one the key ideas, but it is part of it. There are some other ideas behind watched literals. For instance: zero cost backtracking. Not $O(1)$ backtracking, 0 backtracking. Speaking of cost, do you notice something else wasteful about the way we propagate? If not for our *satisfied* marker array, we'd have to rescan the whole equation every time we propagated, which is about twice per level. That's pretty wasteful. After all, why do have to scan every clause in their entirety every single time? Especially because, usually, clauses don't become unsatisfied very often. And because they're not at risk of becoming unsatisfied very often, why do we need to scan it? And besides, it's not like every single variable appears in every single clause, so why not scan only the clauses we have to? I hope these questions are leading you to think of other questions regarding inefficiencies in our current implementation.

These questions were the driving force into the invention of the *Watched Literals* algorithm.

History can be found here¹⁶

Let's learn how the Watched Literals algorithm works. Earlier, I made a comment about how usually, each clause doesn't contain every single variable in the equation. So when we force a variable's assignment, we don't need to check every single clause, only at most the ones that contain the variable and its negation. (Ie, if we force a variable, say

¹⁵This headache was also induced by my lack of ability to fall asleep at all last night AND a really intense session at the gym. Computer science gives you headaches for sure, but not like this one.

¹⁶<https://school.a4cp.org/summer2011/slides/Gent/SATCP3.pdf>

$2 = T$, we only need to check the clauses containing 2 and the clauses containing -2 . However, do we *need* to check the clauses containing 2? Think about it. Those clauses are already satisfied by $2 = T$, so there's nothing to propagate on, for those clauses. The only clauses that have a *chance* to produce a propagation are the ones containing -2 because $-2 = F$, so if they have one Unassigned variable and everything else is now False, we can propagate!

3.5.1 A Visual Example of Watched Literals

Let's view a visual example. Suppose we had a clause with 5 literals (doesn't matter what they are). The first two literals are watched

■ ■ □ □ □

UUUUU

Suppose the first literal now evaluates to False. We now scan the clause for another literal that doesn't evaluate to False. We will find it at the third literal.

□ ■ ■ □ □

FUUUU

Now suppose the fifth literal becomes False. We do nothing because we're not watching the fifth literal.

□ ■ ■ □ □

FUUUF

Now the third literal becomes False. We scan again for another not-False literal.

□ ■ □ ■ □


FUFUF

Now the second literal becomes False. Other than the second watched literal (the fourth), there are no more not-False literals. We don't relocate the watch.

□ ■ □ ■ □


FFUFU

However, now the second watched literal knows it's the last Unassigned literal in the clause. It will now evaluate to True!



FFFTF

Our clause is now satisfied. Suppose instead, when we were here:




FFUFU

That the fourth literal also became False.



FFFFU

We've now encountered a *contradiction*, meaning that whatever assignment we have is impossible to satisfy the equation. Suppose we backtrack, and that restores us to a state where two literals in this same clause are now Unassigned. What happens to the watched literals?



UUFFF

The answer is that they stay where they are.



UUFFF

This is because one watched literal is Unassigned, so everything is okay. Although this last diagram isn't accurate to what a real restored state would look like, we can suspend our disbelief for the sake of simplicity. The reason why watched literals work is because of the **Watched Literal Invariant**, which is that at any given point in time, at least one of the watched literals is either True or Unassigned. This invariant is what allowed us to propagate the clause when there was only one Unassigned literal left, and we will see further how it helps.

3.5.2 Algorithmic Changes

```

1  public interface:
2      SATSolver(string input)
3      bool solveCNF()
4
5  private interface:
6      enum var
7      {
8          False
9          True
10         Unassigned
11     }
12
13     void parse(string equation)
14     bool solve(var[] assignment, int level)
15     optional int[] delta propagate(var[] assignment)
16     bool createWatched()
17
18     int numVars
19     int numClauses
20
21     int[][] clauses
22     bool[] satisfied
23     queue<int> unitProp
24
25     optional var[] vars
26     map<int, int[]> var_to_clause
27     map<int, (int, int)> clause_to_vars

```

So firstly, there are some major differences. We removed the *evaluate()* method. This is because now that propagation checks for contradictions, having all variables assigned is now **equivalent** to finding a satisfactory assignment of variables. So we don't need to evaluate our assignment anymore! That's significantly less work than what we had before, which is what we want. Next, we added in a queue, *unitProp*, kind of like what we had for unit propagation before. This time, it's now a member of the solver object instead of a local object within the method. We will get back to that in a second. Lastly, we added in two maps, one from variables to clauses, and the other from clauses to a pair of variables.

These data structures are used for fast lookup between variables to **indices** of clauses (for *var_to_clause*), and one from **indices** of clauses to **indices** of watched literals. It's important that we remember we're using indices instead of actual clauses because forgetting how we're using our mapping is going to create a lot of bugs and we're gonna have a bad time.

We also added in a new method, called *createWatched()* and what it does is it creates our initial maps for *var_to_clause* and *clause_to_vars*. Here is the pseudocode:

```

1      bool createWatched(int[] assignment) {
2          for (i = 0 up to numClauses) {
3              clause = clauses[i]
4
5              if (clause.size >= 2) {
6                  append(var_to_clause[clause[0]], i)
7                  append(var_to_clause[clause[1]], i)
8
9                  clause_to_var[i] = {0, 1}
10             } else if (clause.size == 1) {
11                 int var = clause[0]
12                 int idx = abs(var)
13                 append(var_to_clause[var], i)
14                 clause_to_var[i] = {0, 0}
15
16                 enqueue(unitProp, var)
17
18                 if (var > 0) {
19                     assignment[var] = True
20                 } else {
21                     assignment[var] = False
22                 }
23             } else {
24                 return False
25             }
26         }
27
28         return True
29     }

```

So here's how to read it: for each clause, it will fall into one of three cases:

1. It's size is greater than or equal to 2
2. It only has 1 variable
3. It's empty

In the first case, the two literals that will be watching that clause will be the first two literals. So the index i gets appended to their array of clause indices. Then, we record that clause i is being watched by its first and second literal.

In the second case, since there's only one literal, it gets watched by only that literal (duh). We also record that clause i is being watched by only its first literal. And lastly, we enqueue that literal to our unit propagation queue so we can do a root level unit propagation before we start solving. We also want to immediately force the assignment of the variable so that way we can use its value immediately in solving.

In the last case, since there are no literals, this clause is automatically unsatisfiable, so therefore this whole equation is unsatisfiable, so we return False so that when we call this method, we know that we can't solve it.

Now let's look at our modified solving methods:

```

1      bool solveCNF() {
2          var[1..numVars] assignment
3
4          for each (var in assignment) {
5              var = Unassigned
6          }
7
8          if (createWatched(assignment) == False) {
9              return False
10         }
11
12         if (propagate(assignment, 0) returns nothing) {
13             return False
14         }
15
16         solve(assignment, 0)
17     }

```

There are some slight differences: we create our watched literals and make sure that there isn't an unsatisfiable clause, and we make sure that *propagate()* returns something. If both those things are successful, then we can start solving away.

```

1      bool solve(int[] assignment, int level) {
2          if (level == numVars) {
3              return True
4          }
5
6          if (assignment[level] != Unassigned) {
7              return solve(assignment, level + 1)
8          }
9
10         int cur_var = level + 1
11
12         assignment[level] = True
13
14         optional int[] prop_result = propagate(assignment, cur_var)
15
16         if (prop_result is optional) {
17             assignment[level] = False
18
19             optional int[] new_prop_result
20                 = propagate(assignment, -cur_var)
21
22             if (new_prop_result is optional) {
23                 return False
24             }
25
26             if (solve(assignment, level + 1)) {
27                 return True

```

```

28         }
29
30         int[] new_delta = new_prop_result.value
31
32         for each (int var_idx in new_delta) {
33             assignment[var_idx] = Unassigned
34         }
35
36         return False
37     }
38
39     if (solve(assignment, level + 1)) {
40         return True
41     }
42
43     int[] delta = prop_result.value
44
45     for each (int var_idx in delta) {
46         assignment[var_idx] = Unassigned
47     }
48
49     assignment[level] = False
50
51     optional int[] new_prop_result
52         = propagate(assignment, -cur_var)
53
54     if (new_prop_result is optional) {
55         return False
56     }
57
58     if (solve(assignment, level + 1)) {
59         return True
60     }
61
62     int[] new_delta = new_prop_result.value
63
64     for each (int var_idx in new_delta) {
65         assignment[var_idx] = Unassigned
66     }
67
68     return False
69 }

```

This is almost the same as our old unit propagation code, except for that we don't keep track of which clauses are satisfied across levels, and the propagation method returns an array of variable indices instead of the actual variables. The code is also split up into multiple branches depending on whether or not propagation was successful. Some parts of the code seem really similar to each other because of this. I am sure that there may be a way to simplify the similar parts, but I like the way I wrote it because I

Now for the interesting part:

```

1 optional int[] propagate(int[] assignment, int decision) {
2     int[] delta
3
4     if (decision != 0) {
5         enqueue(unitProp, decision)
6
7         int idx = abs(decision)
8         append(delta, idx)
9     }
10
11     while (unitProp.size != 0) {
12         int prop = unitProp.pop()
13
14         if (contains(var_to_clause, -prop) == False) {
15             continue
16         }
17
18         int[] watched_clauses = var_to_clause[-prop]
19
20         for (i = 0 up to watched_clauses.size) {
21             clause_idx = watched_clauses[i]
22
23             int[] clause = clauses[clause_idx]
24             (int, int) watches = clause_to_var[clause_idx]
25
26             if (clause[watches.first] != -prop) {
27                 swap(watches.first, watches.second)
28             }
29
30             bool relocated = False
31             for (j = 0 up to clause.size) {
32                 if (j == watches.first or j == watches.second) {
33                     continue
34                 }
35
36                 int unwatched = clause[j]
37                 int unwatched_idx = abs(unwatched)
38
39                 bool evals_false =
40                     (unwatched > 0 and assignment[unwatched_idx]
41                      == False)
42                     or
43                     (unwatched < 0 and assignment[unwatched_idx]
44                      == True)
45                 if (!evals_false) {
46                     append(var_to_clause[unwatched], clause_idx)
47                     erase_at(watched_clauses, i)

```

```

48         watches.first = j
49         relocated = True
50         break
51     }
52 }
53
54 if (relocated == False) {
55     i = i + 1
56
57     int other_watch = clause[watches.second];
58     int idx = abs(other_watch)
59
60     if (assignment[idx] == Unassigned) {
61         if (other_watch > 0) {
62             assignment[idx] = True
63         } else {
64             assignment[idx] = False
65         }
66
67         append(delta, idx)
68         enqueue(unitProp, other_watch)
69     } else {
70         bool evals_false =
71             (other_watch > 0 and assignment[idx]
72              == False)
73         or
74             (other_watch < 0 and assignment[idx]
75              == True)
76
77         if (evals_false == False) {
78             continue
79         }
80
81         for (k = 0 up to delta.size) {
82             assignment[delta[k]] = Unassigned
83         }
84
85         clear(unitProp)
86
87         return null
88     }
89 }
90 }
91 }
92 return delta
93 }

```

This is a huge method, so let's walkthrough it so we understand every aspect of it. If we look back at our *solve()* code, we pass in *cur_var* and *-cur_var* as the *decision*

parameter into the propagation method. And in the beginning of `solveCNF()`, we pass in 0, to signify propagation at the root level. When that happens, we don't add anything into the queue, instead we just propagate on the literals that were forced during the creation of the watched literals.

Like I said before, we only need to check the clauses containing the negative version of the literal being propagated on because those are the clauses that the actual unit propagation could be forced upon. However, we only need to actually check the clauses being watched by the negative version of the literal being propagated on.

Next, we grab the indices of those clauses being watched and then for each one, we try to relocate the index of the first watched literal to a new literal that doesn't evaluate to False. Why only the first? Because swapping the two indices and operating on only the first index is far easier than creating a second branch dedicated to the second index. If a literal doesn't evaluate to False, we first make that unwatched literal now watch the current clause, we then make the current watched literal stop watching the current clause, and then we update the first index of the watched literals to be the index of the newly watching literal. That means we have successfully relocated the watched literal and no unit propagation needs to occur. Because we're modifying the array of clause indices, we don't want to increment the index variable in our implementation.

Aside: When you delete something in a dynamic array, in order to fill the blank space, every element to the right of what you deleted has to shift over to the left by one unit. This takes $O(n)$ time and on large arrays can be very costly. Because the next element has shifted over to the left, its index is now the current index, so you don't want to increment the index variable. In the implementation, the relocation logic is implemented as a "swap and pop" in which we swap the current element with the last one, then we delete the last element, which was the current element. This leads to no shifting of elements, and is logically equivalent (for our use case), even if the order of elements, relative to the deleted element, is no longer the same.

If we couldn't relocate the watched literal's index, we want to increment the index variable so that on the next iteration it will point to a new element. In the case that we couldn't relocate the watched literal, it means that every element in the clause, except for possibly the element pointed to by the second watched literal, evaluates to False. If the element pointed to by the second watched literal also evaluates to False, then we've encountered a contradiction in our assignment. However, if the second watched literal evaluates to True, then the clause is already satisfied and we need to do nothing except for move onto the next variable to propagate on. And if the second watched literal evaluates to Unassigned, then we can assign it the value it needs to evaluate to True and we then propagate on that newly assigned unit literal.

If we encounter a contradiction, we must undo all the propagation we've done, including the initial decision before returning a null value to signify that the propagation failed. However, if it has succeeded, then we return the array of the variable indices that we've assigned.

So one of the big aspects of watched literals is how we get away with doing so little amount of work. And the answer is something called the "watched literal invariant." This invariant is that at all times, at least one of the watched literals doesn't evaluate to False. I don't expect this to make everything click immediately, so let's explain more. A clause is only a contradiction when everything evaluates to False. However,

we don't want to check every single clause when we assign only one variable. And even checking the negative of the variable we assigned is too much. So instead, we change which variables are watching which clauses every now and then. We pick at most two variables to watch a singular clause with, and each one watches their respective clauses until they are assigned a value that makes them evaluate to False. In that case, there is the possibility that the watched literal invariant may be violated (ie, there is a possibility that both watched literals may evaluate to False), but just to be sure, we try to change the variable watching that same clause.

We scan the clause for any unwatched literal that is Unassigned or True, and if we find one, we've restored the watched literal invariant, so we don't need to propagate or do anything else for that matter. However, if we can't find one, then we're in serious danger of the watched literal invariant being violated. So our only hope is for the second watched literal to either already be True (we need to do nothing), or Unassigned (we assign it and propagate).

Thanks to that invariant, we also don't need to do any undoing when we backtrack other than undoing what we propagated. So long as the watched literal invariant holds true when we undo our assignments, the watched literals don't need to move, which means besides the undoing of what we propagated, we get zero cost backtracking.

3.5.3 Why Two Watches?

This now begs the question, why two watched literals? What is special about the number 2? Why can't we have just one watched literal? In order to answer this question, we need to ask ourselves: "what can we do with two watched literals that we can't do with one?" When we have two watched literals, we can efficiently check for whether or not we can propagate. How so?

After we try propagating on a literal, we scan the whole clause except for the indices of the first and second watched literals. Next, if everything in the clause evaluates to False, our only hope to propagate is on the second watched literal. So if that evaluates to False as well, we have found a contradiction, and if it doesn't we have found a literal we can propagate on.

But why is that more efficient than just one watched literal? Let's think of the worst case scenario for if we just had one watched literal. What does that look like? Suppose we were watching the first literal in the clause and every other literal in the clause evaluated to False. Ie, we need to propagate on this watched literal. How would it know it needs to propagate when it was never alerted by anything becoming False? It would have no idea it needs to do anything until that watched literal was falsified, and by that point, everything in the clause would evaluate to False, so the clause would be unsatisfied, only then would we backtrack, but when we backtrack, we'd still have no way of knowing if the clause is unit.

So with only one watched literal, in the worst case scenario, which is that the literal we're watching is the last literal in the clause to be falsified, we'd be doing the same amount of work as normal unit propagation, perhaps even more.

Visually speaking, this is what I'm talking about:



UUU

The last literal becomes False, but we're not watching it so nothing moves.

■□□

UUF

Next the second literal becomes False, but we're not watching it so nothing happens.

■□□

UFF

Although we should perform unit propagation, our watched literal is completely blind to the fact that it is the last Unassigned literal in the clause, so it remains as is. If luck has it, it could become True and all would be well, but this is reality, so more likely than not, it will also become False.

■□□

FFF

Now we scan the clause, and we see that we can't relocate the watched literal, so we backtrack on our assignment.

■□□

UFF

Since the watched literal was never aware if both of the other literals became False in one pass of the algorithm or if it was across multiple decisions, it still doesn't know that it's the last Unassigned literal. It has to scan the clause each time to know whether or not it's the last Unassigned literal, which is no better than regular unit propagation.

3.5.4 Watched Literals - Results

After implementing watched literals, our runtime decreases dramatically. For uf20.91, solving all 1000 instances takes **349ms**, as opposed to the **1548ms** from pure literal elimination and the **868ms** from normal unit propagation. And for UUF50.218.1000, it takes only **7130ms** as opposed to the **231,107ms** from pure literal elimination and the **140,411ms** from normal unit propagation. Watched literals took us down from minutes to seconds for the unsatisfiable problems, which is amazing for us.

```

On all 1000 equations in uf20-91:

Watched Literals:           349ms (~0.35 seconds).
Pure Literal Elimination:   1548ms (~1.55 seconds).
Unit Propagation:          868ms (~0.87 seconds).
Recursive Backtracking:    1,540,613ms (~25.68 minutes).
Brute Force:               1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):   49,874ms (~50 seconds).
Parallel:                  10,589ms (~10.5 seconds).

On all 1000 equations in UUF50.218.1000:

Watched Literals:           7,130ms (~7.13 seconds).
Pure Literal Elimination:  231,107ms (~3.85 minutes).
Unit Propagation:          140,041ms (~2.34 minutes).

```

3.6 Iterative DPLL

Now that we’ve implemented DPLL to the best of our ability, it’s time to take it one step further. We will now implement it iteratively rather than recursively. Why? Remember how I said in the beginning that DPLL is the foundation for modern SAT solvers?¹⁷ Well iterative DPLL is the foundation for the next chapter’s algorithm: CDCL. Implementing iterative DPLL is imperative in order to create a modern SAT solver worth its salt.¹⁸

As said in the beginning of the chapter, DPLL forces us to traverse the decision tree as well, a tree. And that is still true, even in iterative DPLL. What changes now is that instead of an implicit path, we now have an explicit one. Suppose we are on the first decision level, so we decide variable $1 = T$, and that forces $3 = F$, $5 = F$, and $6 = T$. In recursive DPLL, we assign these variables into the *assignment* or *vars* field, but we don’t really use them until we get to the decision level that matches their index or the evaluation phase.

With iterative DPLL, there are still decision levels, but the representation is now very, very different because instead of a decision *tree*, we now use a decision *trail*. The trail is reminiscent of a queue in the sense that we only operate at the end of the data structure. It contains some information that we need in order to continue mimicking the behavior of recursive DPLL. Each trail entry contains the literal that was forced, the decision level it’s a part of, and the clause that forced that decision. We append these trail entries every single time we force a variable’s assignment. The next question is how do we know which variables were decided for and which variables were inferred from propagation? Well, for decisions, there is no reason clause because there was no reason to begin with other than “because I said so.” Therefore, the variables that were decided for don’t get a reason clause, but the ones that were inferred do. This is great

¹⁷I can’t remember if I actually said that

¹⁸Honestly, “Salt” is a better name than “Saturn.” I might change it.

because it makes it easy to tell where the beginning of one decision starts and where it ends.

However, it would be nice to have some way to keep track of where the decisions are in the trail other than scanning for entries without a reason clause. Because of that, we will also keep another trail specifically for the decision markers. This is also helpful because then we can store whether or not we've "tried the other branch" when we backtrack.

While on the topic of backtracking, it would be nice to discuss it some more as well. Since we can't simply go up the call stack and undo variable assignments, we have to do it in the iterative way. When we backtrack, we first backtrack up to the decision entry. This is done by popping off the trail and undoing the propagated assignments. Once that's done, you then check if you've "tried the other branch", and if you haven't, you flip the decision literal's assignment and then retry. If you have already tried the other branch, you pop off the decision entry, and then repeat this process to the next decision entry. This process continues until you either find a decision that has not been switched yet, or you run out of decision entries.

With most of the exposition out of the way, I think it would be good to look at the changes we will make to our interface and algorithms, starting with our SAT solver.

```

1      public interface:
2          SATSolver(string input)
3          bool solveCNF()
4
5      private interface:
6          enum var
7          {
8              False
9              True
10             Unassigned
11          }
12
13          TrailEntry {
14              int lit,
15              int clauseIdx
16          }
17
18          DecisionLevel {
19              int idx,
20              bool triedOtherBranch
21          }
22
23          void parse(string equation)
24          bool propagate()
25          bool createWatched()
26
27          int numVars
28          int numClauses
29          int trailHead

```

```

30
31         int[] [] clauses
32         bool[] satisfied
33
34         TrailEntry[] trail
35         DecisionLevel[] trailStarts
36
37         optional var[] vars
38         map<int, int[]> var_to_clause
39         map<int, (int, int)> clause_to_vars

```

The first notable change is that *propagate()* now returns a boolean instead of a delta. This is because it will no longer be responsible for undoing propagation in response to a contradiction, thus achieving zero-cost backtracking. Next, we now have a *trail* and a *trailHead*, which points to the last made decision entry. *trailStarts* keeps track of all the indices of the decision entries. One last thing to note is that the *TrailEntry* data structure uses indices for clauses instead of the actual clause itself. There's a very good reason for this, and I will reveal the answer in the next chapter, but until then, why do you think it might be that way? Hint: it's not because indices are easy for accessing.

The next change we need to make is to our algorithm to create watched literals. It's the same for basically 90% of the algorithm, except for the unit clause case, which I'll highlight.

```

1         bool createWatched(int[] assignment) {
2             for (i = 0 up to numClauses) {
3                 clause = clauses[i]
4                 ...
5                 else if (clause.size == 1) {
6                     int var = clause[0]
7                     ...
8                     append(trail, {var, i})
9                     ...
10                }
11                ...
12            }
13            return True
14        }

```

Rather than enqueueing the variable, we create a new trail entry for each unit clause. These ones are all part of the **root** decision level, which we mentally reserve for the 0th decision level. If we backtrack up to this level (if it exists) or before it, we've encountered a root level contradiction, so the equation is unsatisfiable. Because not every equation has unit clauses, it's important to remember that this decision level may not always exist, and so we should be keep that in mind for our solving algorithm.

The next algorithm we're going to change is the *propagate()* algorithm because the changes are also very minimal.

```

1         bool propagate(int[] assignment) {

```



```

2         while (trailHead < trail.size) {
3             int prop = trail[trailHead].lit
4             trailHead = trailHead + 1
5             ...
6             for (i = 0 up to watched_clauses.size) {
7                 clause_idx = watched_clauses[i]
8                 int[] clause = clauses[clause_idx]
9                 (int, int) watches = clause_to_var[clause_idx]
10                ...
11                if (relocated == False) {
12                    ...
13                    int other_watch = clause[watches.second];
14                    int idx = abs(other_watch)
15                    if (assignment[idx] == Unassigned) {
16                        ...
17                        append(trail, {other_watch, clause_idx})
18                    } else {
19                        return false
20                    }
21                }
22            }
23        }
24        return true
25    }

```

The main difference is that we just *trailHead* to iterate through the trail, and instead of enqueueing, we just append to the trail. We also return True/False depending on whether or not propagation succeeded.

Lastly, we look at our *solve()* method, which has the most notable changes. There are two ways to implement this, depending on what makes sense to you. I call them “assignment before propagation” and “propagation before assignment.” I call them these because that’s what they are. Before showing the pseudocode, I want to talk about the algorithm a bit more. The beginning is mostly the same: create our watched literals, check that it succeeded, return False if it didn’t. Otherwise, we do a root level propagation and then check if it succeeds as well. However, the main confusion comes from handling the case of when there is a 0th decision level and when there isn’t. There are multiple ways to go about this. After the beginning, we enter the solving loop.

There are two main parts: assigning and backtracking, with a call to the propagate method separating them. Assigning is now scanning for the first Unassigned variable, forcing it True, creating a new trail entry, and then updating the trail head. I’m sure there’s a better way to keep track of the next unassigned variable other than scanning the entire assignment list on each loop iteration, however, for simplicity’s sake, we’re implementing it the naive way.

Since we’ve already discussed backtracking, we will now get into the algorithm itself:

```

1     bool solve() {
2         var[1..numVars] assignment

```

```

3
4     for each (var in assignment) {
5         var = Unassigned
6     }
7
8     if (createWatched(assignment) == False) {
9         return False
10    }
11
12    if (trail.size > 0) {
13        append(trailStarts, {0, true})
14
15        if (propagate(assignment, 0) == False) {
16            return False
17        }
18    }
19
20    while (True) {
21        bool succeeded = propagate(assignment)
22
23        if (!succeeded) {
24            while (trailStarts.size > 0) {
25                DecisionLevel decision = trailStarts.last
26
27                while (trail.size > decision.idx + 1) {
28                    int lit = trail.last.lit
29                    assignment[abs(lit)] = Unassigned
30                    erase(trail, trail.last)
31                }
32
33                if (decision.triedFalse == False) {
34                    TrailEntry entry = trail.last
35
36                    entry.lit = -entry.lit
37                    assignment[abs(entry.lit)] = False
38                    decision.triedFalse = True
39
40                    trailHead = decision.idx
41                    break
42                } else {
43                    int lit = trail.last.lit
44                    assignment[abs(lit)] = Unassigned
45                    erase(trail.last)
46                    erase(trailStarts.last)
47                }
48            }
49
50            if (trailStarts.size == 0) {
51                return False
52            }

```

```

53
54         continue
55     }
56 }
57
58 //Assignment code
59 int firstUnassigned = -1
60 for (i = 1 up to numVars) {
61     if (assignment[i] == Unassigned) {
62         firstUnassigned = i
63         break
64     }
65 }
66
67 if (firstUnassigned == -1) {
68     return True
69 }
70
71 assignment[firstUnassigned] = True
72 append(trail, {firstUnassigned, -1})
73 append(trailStarts, {trail.size - 1, False})
74 trailHead = trailStarts.last.idx
75 }

```

Here is the pseudocode for solving, written as “propagate then assign.” Quick question: when there’s no 0th decision level and we propagate, how do we know that none of the propagated variables will have the wrong decision level? Ie, how do we know that they’ll have a decision level of 0 instead of 1, since we haven’t made any decisions yet? This is kind of a trick question, the answer is that if there is no 0th level, *propagate()* will do nothing because the while-loop condition will immediately return False, and so the method will terminate, assigning nothing!

Instead of showing you the code for “assign then propagate” (it is NOT just taking the assigning code and moving it up before the call to *propagate()*), let’s run through how this works. When there’s a 0th level, we obviously want to mark it in our *trailStart*, but why do we want to mark that the “other branch” has already been tried? Well that’s because those assigned by the root level propagation are forced into those values, so it doesn’t make sense to swap their values to be their opposites since it will immediately be an unsatisfiable assignment anyways.

Next, in the main solving loop, the backtracking code is just “grab the last decision index, undo until the root decision, if it’s already tried False, undo the decision entirely and repeat with the previous one, and if not, try False and continue.”

Lastly, let’s talk about doing assign then propagation. Here’s why it doesn’t make total sense to start out with propagation:

- we do a root level propagation if there needs to be one.
- if there is no root level propagation to be done, it still doesn’t make sense to immediately propagate since there’s nothing to propagate.

- if there was a root level propagation, it doesn't make sense to do another one immediately since we've just completed one.

So if you wanted to do assign and then propagate, how would that be done? Earlier, I said that it's not just taking the code for assigning and then moving it before propagation, but why? Well, suppose we fail a propagation. We backtrack, and then suppose we haven't tried the other branch for that decision. We encounter a *continue* statement and we go back to the top. If we simply move the assigning code up to the top of the loop, what's the issue? Well remember, we just flipped a decision, so we should actually be *propagating* instead of assigning. If we assign immediately after flipping the decision, we now have two decisions at the same time, and that's bad because what if they are contradictory to one another, ie one would force a variable False and the other True? Unfortunately, the propagation code wouldn't catch that because it assumes everything in the trail is valid, but we just made the trail invalid. So to the propagation code, instead of it catching that a variable is forced to be False, but it's True, or vice versa, it may just ignore it because it focuses mainly on Unassigned and False literals.

So if we want to do assign then propagate, how do we make sure that we don't accidentally push something onto the trail when we shouldn't have? We need to wrap the assigning code in a conditional statement that checks that *trailHead == trail.size*, if it doesn't, then that means that there's something to be propagated on because we're not at the end of the trail, but if it does, then we've reached the end, so we should check if we've assigned everything and can be done, or if we need to assign something new.

3.6.1 Iterative DPLL - Results

Here are the results after this huge refactor:

(Assignment before propagation)

uf20.91: **328ms**, down 21ms from **349ms** by Watched Literals

UUF50.218.1000: **6369ms**, down 761ms from **7130ms** by Watched Literals

(Propagation before assignment)

uf20.91: **324ms**

UUF50.218.100: **6365ms**

```

On all 1000 equations in uf20-91:

Iterative DPLL (AbP):      328ms (~0.33 seconds).
Iterative DPLL (PbA):      324ms (~0.32 seconds).
Watched Literals:          349ms (~0.35 seconds).
Pure Literal Elimination:  1548ms (~1.55 seconds).
Unit Propagation:          868ms (~0.87 seconds).
Recursive Backtracking:    1,540,613ms (~25.68 minutes).
Brute Force:               1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):   49,874ms (~50 seconds).
Parallel:                  10,589ms (~10.5 seconds).

On all 1000 equations in UUF50.218.1000:

Iterative DPLL (AbP):      6,369ms (~6.37 seconds).
Iterative DPLL (PbA):      6,365ms (~6.37 seconds).
Watched Literals:          7,130ms (~7.13 seconds).
Pure Literal Elimination:  231,107ms (~3.85 minutes).
Unit Propagation:          140,041ms (~2.34 minutes).

```

Should you do propagation before assignment simply because it's faster? It really doesn't matter, you do you. However, for the rest of this book, we will be using assignment before propagation because it's easier for me to understand, and it follows online pseudocode examples better than propagate before assignment.

The speed gain on satisfiable problems is very miniscule, however, if you ever want proof that iterative algorithms are faster than recursive ones, the unsatisfiable problem time decrease is pretty good evidence.

Chapter 4

CDCL

Conflict Driven Clause Learning (CDCL), is *the* modern SAT solving algorithm. It builds off of the ideas of DPLL and takes it to a whole new level. The “conflict driven” part of CDCL means that we analyze conflicts (contradictions) and see what we can learn from them. The “clause learning” part is exactly what it sounds like, we add new clauses to our database. In other words, CDCL is about analyzing why conflicts occur when we solve and using those conflicts to learn new clauses that will help enforce new rules as to what values variables can take, essentially changing the CNF equation into one that’s logically equivalent, but easier to solve. I won’t lie to you, I was not able to implement this algorithm the first time I tried. I hope and pray that this time I will be successful, so that we will be successful together dear reader.

Aside: why have we been opting for indices so often? You see, although we live in the land of pseudocode or machine-agnosticism, implementations don’t (duh). So in our previous chapter in DPLL, there were times where we used clause indices instead of pointers to clauses (obviously using copies would be terrible for performance and memory), but why did we use indices? Now that we’re implementing CDCL, we’re going to be adding clauses to the database quite often, and that can trigger a resize on the database. If we used pointers to clauses, after a resize triggers, the memory address of each clause would change, thus making the pointers invalid. We would need to add new code after the insertion of each clause to make sure a resize didn’t trigger, and if one did, we’d need to update each pointer for each reason clause and watched literal, which would be very difficult because after a resize, we now have no way to tell which clauses mapped to which reasons and watched literals, because, well, the old mapping isn’t valid and contains no information about the new addresses for everything.

1

¹More information about CDCL can be found here: <https://www.cis.upenn.edu/~cis1890/files/Lecture4.pdf>

4.1 Clause Learning

The first step in creating a CDCL solver is to learn and use new clauses. There are many versions involving different ways of backtracking, variable assigning, and even how we learn new clauses, but for now, we will keep our implementation as close to our iterative DPLL version and work our way up to the more advanced and widely used versions of the algorithm.

In order to learn a new clause, we first need to understand the learning process. The big idea is that we will encounter conflicts and contradictions while assigning and propagating new literals, and we want to learn from those conflicts and contradictions in a way that would ideally prevent us from making the same mistake again. Our unit propagation method is the perfect tool for this because it already analyzes our clauses for us and can detect when a contradiction has occurred. Rather than returning True or False to indicate a successful or unsuccessful propagation respectively, when we encounter a contradiction, we want to return the clause that is unsatisfied. You may know that when we propagate, there could be more than one clause that is unsatisfied, but just not yet evaluated to be unsatisfied. Even though there could be many conflicting clauses at once, we just want to return the first one that we find. You may wonder why not learn from all of them, and the answer is that by resolving the bad decision/propagation we made, we will (hopefully) prevent ourselves from doing more propagations that lead to those extra unsatisfied clauses.

After we return the conflicting clause to the main solving method, we then want to perform “clause resolution.” As stated in the beginning, there are many versions of clause resolution, mainly differing in when they/how they stop resolving. For our basic CDCL implementation, we will resolve the conflicting clause until it contains no more propagated literals, ie, it’s a clause of only decisions. The reason why this is a good stopping point is because we made all decisions, and because of that, clause learning will essentially allow us to learn which decisions we can’t make together.

Lastly, after we resolve on the learned clause, we want to add it to the database, and we want to initialize its watched literals so that we can use it as if it were part of our initial CNF equation. Before we discuss how clause resolution works at a theory level, let’s first discuss a few implementation changes we’ll need to make.

```
1     public interface:
2         SATSolver(string input)
3         bool solveCNF()
4
5     private interface:
6         enum var
7         {
8             False
9             True
10            Unassigned
11        }
12
13        TrailEntry {
14            int lit,
```

```

15         int decisionLevel,
16         int clauseIdx
17     }
18
19     DecisionLevel {
20         int idx,
21         bool triedOtherBranch
22     }
23
24     void parse(string equation)
25     bool propagate(int[] decisionLevel, int decisionLevel)
26     bool createWatched()
27     bool createWatched(int idx, int decisionLevel)
28
29     int numVars
30     int numClauses
31     int trailHead
32
33     int[] [] clauses
34
35     TrailEntry[] trail
36     DecisionLevel[] trailStarts
37     int[] var_to_trail
38
39     optional var[] vars
40     map<int, int[]> var_to_clause
41     map<int, (int, int)> clause_to_vars

```

The *var_to_trail* field is the only change we'll need to make, and it is just a mapping from variable index to trail index. If you enforce the invariant that each variable maps to exactly one place in the trail and no two variables map to the same place in the trail at the same time, you don't need to initialize these values to anything.

The next thing we'll look at is our *propagate()* code since it needs very minimal changes. The change we're making now is that instead of returning a boolean value to indicate whether or not propagation was successful, we're now going to return a clause if we encounter a contradiction and nothing otherwise.

```

1     bool propagate(int[] assignment, int decisionLevel) {
2         while (trailHead < trail.size) {
3             int prop = trail[trailHead].lit
4             trailHead = trailHead + 1
5             ...
6             for (i = 0 up to watched_clauses.size) {
7                 clause_idx = watched_clauses[i]
8                 int[] clause = clauses[clause_idx]
9                 (int, int) watches = clause_to_var[clause_idx]
10                ...
11                if (relocated == False) {
12                    ...

```



```

13         int other_watch = clause[watches.second];
14         int idx = abs(other_watch)
15         if (assignment[idx] == Unassigned) {
16             ...
17             append(trail, {other_watch, decisionLevel,
clause_idx})
18         } else {
19             return clause
20         }
21     }
22 }
23 }
24 return null
25 }

```

And before we look at changes to our solving method, we're going to create a new overload for *createWatched()* that takes in the index of a clause. Its only purpose is to be called on newly learned clauses, which don't have the same guarantees as the other overload. You see, the learned clause should hopefully NEVER be unsatisfied when this method operates over it. It should ALWAYS return True, so if it ever doesn't, it's a good debugging tool. Here's how it works: we traverse over the clause of the passed in index, counting the number of True, False, and Unassigned literals it has. We will then search for two indices that could be watched literals. Since the clause should never be unsatisfied, we are guaranteed to always find one (if our implementation is correct). If we can't find another and our clause has more than one literal, we will initialize the second watched literal to any arbitrary place in the clause (with the caveat that it can't be the same index as our first watched literal). We will then check if the clause is unit, and if it is, we will add it to our trail, assign the Unassigned variable, etc. Otherwise, we do nothing.

```

1     bool createWatched(int index, int decisionLevel) {
2         clause = clauses[index]
3         if (clause.size == 0) {
4             return False
5         }
6
7         bool foundFirst = False
8         bool foundSecond = False
9         int firstIdx = -1
10        int secondIdx = -1
11
12        int numUnassigned = 0
13        int numFalse = 0
14        int numTrue = 0
15
16        for (i = 1 up to clause.size) {
17            int idx = abs(clause[i])
18            if (assignment[idx] == Unassigned) {

```

```

19         numUnassigned = numUnassigned + 1
20         if (foundFirst == False) {
21             foundFirst = True
22             firstIdx = i
23         } else if (foundSecond == False) {
24             foundSecond = True
25             secondIdx = i
26         }
27     } else if (evals_false(clause[i])) {
28         numFalse = numFalse + 1
29     } else {
30         numTrue = numTrue + 1
31
32         if (foundFirst == False) {
33             foundFirst = True
34             firstIdx = i
35         } else if (foundSecond == False) {
36             foundSecond = True
37             secondIdx = i
38         }
39     }
40 }
41
42 if (clause.size >= 2) {
43     if (foundSecond == False) {
44         secondIdx = (firstIdx + 1) % clause.size
45     }
46
47     append(var_to_clause[clause[firstIdx]], index)
48     append(var_to_clause[clause[secondIdx]], index)
49
50     clause_to_var[index] = {firstIdx, secondIdx}
51 } else {
52     append(var_to_clause[clause[firstIdx]])
53     clause_to_var[index] = {firstIdx, firstIdx}
54 }
55
56 if (numTrue == 0 && numUnassigned == 1) {
57     append(trail, {clause[firstIdx], decisionLevel, index})
58     var_to_trail[abs(clause[firstIdx])] = trail.size() - 1
59
60     int var = clause[firstIdx]
61     int idx = abs(var)
62
63     if (var > 0) {
64         assignment[idx] = True
65     } else {
66         assignment[idx] = False
67     }
68 }

```

```

69
70     return True
71 }

```

Because we don't know anything about this clause, we need to gather as much information as we can right away before we decide how watched literals are spread out on the clause. When the clause has more than one literal, we use " $\text{secondIdx} = (\text{firstIdx} + 1) \% \text{clause.size}$ " because that prevents us from adding a watched literal past the end of the clause, if *firstIdx* is already on the edge.

The last method we need to discuss is our *solve* loop. Clause resolution happens after backtracking because it relies on the current trail to understand ² why a conflict occurred. If we were to undo propagations, then the resolution operator would have no reference as to what "most recently" means, and therefore, clause resolution wouldn't work.

```

1  bool solve() {
2      var[1..numVars] assignment
3
4      for each (var in assignment) {
5          var = Unassigned
6      }
7
8      if (createWatched(assignment) == False) {
9          return False
10     }
11
12     int decisionLevel = 0
13
14     if (trail.size > 0) {
15         append(trailStarts, {0, true})
16
17         if (propagate(assignment, decisionLevel) == False) {
18             return False
19         }
20     }
21
22     while (True) {
23         if (trailHead == trail.size) {
24             //Assignment code
25             int firstUnassigned = -1
26             for (i = 1 up to numVars) {
27                 if (assignment[i] == Unassigned) {
28                     firstUnassigned = i
29                     break
30                 }

```

²It follows a mathematical process to transform the conflict clause into one that prevents that same conflict from ever occurring again. The algorithm does not "understand" anything. Do not anthropomorphize computer algorithms, especially artificial intelligence.

```

31         }
32
33         if (firstUnassigned == -1) {
34             return True
35         }
36
37         assignment[firstUnassigned] = True
38         append(trail, {firstUnassigned, -1})
39         append(trailStarts, {trail.size - 1, False})
40         trailHead = trailStarts.last.idx
41     }
42
43     optional int[] contradiction = propagate(assignment)
44
45     if (contradiction has a value) {
46         int[] conflict = contradiction.value
47
48         bool propagated_found = True
49         while (True) {
50             propagated_found = False
51             int idx = 0
52
53             for (i = 1 up to conflict.size) {
54                 int var = conflict[i]
55                 int v_idx = abs(var)
56
57                 int t_idx = var_to_trail[v_idx]
58                 if (trail[t_idx].reasonIdx has a value) {
59                     propagated_found = True
60                     idx = max(idx, t_idx)
61                 }
62             }
63
64             if (propagated_found == False) {
65                 break
66             }
67
68             int lit = trail[idx].lit
69             int[] reason = clauses[trail[idx].reasonIdx]
70
71             append(conflict, reason)
72
73             for (i = 0 up to conflict.size) {
74                 if (abs(conflict[i]) == abs(lit)) {
75                     swap(conflict[i], conflict.last)
76                     erase(conflict, conflict.last)
77                 } else {
78                     i = i + 1
79                 }
80             }

```

```

81
82         set<int> duplicates
83         for (i = 0 up to conflict.size) {
84             if (contains(duplicates, conflict[i])) {
85                 swap(conflict[i], conflict.last)
86                 erase(conflict, conflict.last)
87             } else {
88                 insert(duplicates, conflict[i])
89                 i = i + 1
90             }
91         }
92     }
93
94     while (trailStarts.size > 0) {
95         DecisionLevel decision = trailStarts.last
96
97         while (trail.size > decision.idx + 1) {
98             int lit = trail.last.lit
99             assignment[abs(lit)] = Unassigned
100             var_to_trail[abs(lit)] = -1
101             erase(trail, trail.last)
102         }
103
104         if (decision.triedFalse == False) {
105             TrailEntry entry = trail.last
106
107             entry.lit = -entry.lit
108             assignment[abs(entry.lit)] = False
109             var_to_trail[abs(entry.lit)] = -1
110             decision.triedFalse = True
111
112             trailHead = decision.idx
113             break
114         } else {
115             int lit = trail.last.lit
116             assignment[abs(lit)] = Unassigned
117             var_to_trail[abs(lit)] = -1
118             erase(trail.last)
119             erase(trailStarts.last)
120         }
121     }
122
123     if (trailStarts.size == 0) {
124         return False
125     }
126
127     append(clauses, conflict)
128     createWatched(clauses.size)
129
130     continue

```

```

131     }
132   }
133 }

```

I want to point out that this isn't necessarily CDCL yet. This is more like DPLL, but with adding clauses to our database. The reason as to why this is, is because whenever we recover from a contradiction, we only go back one level. You see, CDCL does away with the idea of a "tried other branch" flag and relies on the learned clauses to guide what variable assignments are possible and which ones are contradictory. However, because we only backtrack one level, there is nothing to actually force new variable assignments in cases where the learned clause isn't a unit clause. It can happen, but it doesn't always happen, so we still need to try the other assignment since for the most part, the learned clause isn't used.


4.1.1 Clause Resolution

Although we've already seen the pseudocode for clause resolution, it's not discussed very much, and I want us to understand how it works better. As we make decisions and propagations, we form this thing called an **implication graph**. In a DPLL style solver, our assignments and propagations form a directed acyclic graph (I will not be proving this), and by traversing the implication graph, we can find out what decisions led to the contradiction.

To learn more about implication graphs, let's look at one and discuss what it means. We will use the implication graphs from here (Figures 4.1 and 4.2):

<https://www3.cs.stonybrook.edu/~cram/cse505/Fall120/Resources/cdcl.pdf>

We will start with Figure 4.1:



It corresponds to the following CNF equation:

$$(\omega_1) \wedge (\omega_2) \wedge (\omega_3) = (x_1 \vee x_{31} \vee \neg x_2) \wedge (x_1 \vee \neg x_3) \wedge (x_2 \vee x_3 \vee x_4)$$

Each node of the graph follows this form: $x_i = v@d$, but what does it mean?


This is the notation for assignment of a variable at a certain decision level:

$$x_i = v@d$$

where i denotes the index of the variable being assigned, $v \in \{0, 1\}$ (alternatively, $v \in \{F, T\}$), and d is the decision level for which that assignment was executed.

So in our graph, $x_{31} = 0@3$ means variable 31 was assigned 0 (or False) at decision level 3. What about the edges? The edges are clauses, and they connect variables to each other, but what's the relationship? Let's go left to right in our implication graph. Variables 31 and 1 are assigned False at decision levels 3 and 5 respectively, and through clause 1, they both have an edge to variable 2, which is also assigned False at decision level 5. If we look at the original CNF equation, we see that in clause 1, x_1 and x_{31} are not negated, so they are False literals. That leaves x_2 Unassigned, so we propagate in order to satisfy the clause. We assign $x_2 = F$ because it's negated, so its literal will evaluate to True. Likewise, x_1 has another edge of clause 2 (ω_2) to x_3 . If we look at clause 2, we see that $x_1 = F$ forces $\neg x_3 = T \equiv x_3 = F$. Lastly, we see that x_2 and x_3 have an edge to x_4 through ω_3 , forcing its assignment to be $x_4 = T$.

That was pretty easy to understand, so now let's look at an implication graph with contradictions. From Figure 4.2:



This graph corresponds to the CNF equation:

$$(\omega_1) \wedge (\omega_2) \wedge (\omega_3) \wedge (\omega_4) \wedge (\omega_5) \wedge (\omega_6) = (x_1 \vee x_{31} \vee \neg x_2)_1 \wedge (x_1 \vee \neg x_3)_2 \wedge (x_2 \vee x_3 \vee x_4)_3$$

$$\wedge (\neg x_4 \vee \neg x_5)_4 \wedge (x_{21} \vee \neg x_4 \vee \neg x_6)_5 \wedge (x_5 \vee x_6)_6$$

Up until $x_4 = 1@5$, this graph is the same as the one from above. After that, we have some more decisions and propagations, that I bet you can read by now, up until we reach this symbol: κ , which has two edges via ω_6 . This node, κ , corresponds to a contradiction within clause ω_6 because its literals are all False.

When we resolve this conflict, we first start with the clause

$$(x_5 \vee x_6)$$

as that is the clause in which our conflict occurred. Next, we find the most recently propagated literal, which we will assume to be x_6 (the most recently propagated literal

depends on our implementation of propagation, so for simplicity's sake, we will assume numerical order to settle ties in implication orderings on the same level). x_6 was propagated because of the clause

$$(x_{21} \vee \neg x_4 \vee \neg x_6)$$

Combining this with the current conflict clause yields:

$$(x_5 \vee x_6 \vee x_{21} \vee \neg x_4 \vee \neg x_6)$$

which we process to remove all instances of the conflicting variable, leading to a resolution clause of:

$$(x_5 \vee x_{21} \vee \neg x_4)$$

In this clause, the most recently propagated literal is x_5 , so we grab its reason clause (ω_4) and combine once again to get:

$$(\neg x_4 \vee \neg x_5 \vee x_5 \vee x_{21} \vee \neg x_4)$$

Simplifying that clause gives us:

$$(\neg x_4 \vee x_{21})$$

The only propagated literal in the clause is x_4 , so we grab its reason clause (ω_3) and combine again into:

$$(x_2 \vee x_3 \vee x_4 \vee \neg x_4 \vee x_{21})$$

This simplifies down into:

$$(x_2 \vee x_3 \vee x_{21})$$

Our most recently propagated literal is now x_3 . Repeating the process again, we get an unsimplified resolution clause of:

$$(x_2 \vee x_3 \vee x_{21} \vee x_1 \vee \neg x_3)$$

Simplifying it down, we now get:

$$(x_2 \vee x_{21} \vee x_1)$$

Our last propagated literal in the clause is x_2 . We combine our incomplete resolution clause with ω_1 to get:

$$(x_2 \vee x_{21} \vee x_1 \vee x_1 \vee x_{31} \vee \neg x_2)$$

This transforms into

$$(x_1 \vee x_{21} \vee x_{31})$$

which contains only decision literals.

What does this new clause, which we'll call ω_7 , tell us? It tells us that all three decision literals **cannot** all be False (ie, at least one of them **must** be True).

4.1.2 Clause Learning - Results

Here are the results for Clause Learning:

uf20.91: **425ms**, up 98ms from Iterative DPLL

UUF50.218.1000: **9652ms**, up 3287ms from Iterative DPLL

Why did our runtime go up? Well, clause learning isn't easy, and the more contradictions we encounter, the more expensive clause learning we're performing. Another reason is that we're increasing the memory load every time we add in a new clause. Luckily, our testing sets are small, but on large problems, this would be a much larger issue. So then why do modern solvers use clause learning? Because there are ways to make it much faster. The first step would be to implement non-chronological backjumping, where instead of moving up just one decision level in our trail, we move up multiple at a time. The second step would be to implement First UIPs, which makes clause learning a less expensive operation (doesn't fix the memory load though), and it combined with non chronological backjumping allows us to instantly and fully use our learned clause to its potential. Although we have taken a small step back, we will see what leaps forward we will make in the next sections.

*Aside: Why did we get rid of pure literal elimination? The problem with pure literal elimination is that once we start learning clauses, previously pure literals may no longer be pure anymore once we add in a new clause. Because pure literals aren't born from decisions/ assignments, they aren't as "concrete" in their status. (Ie, propagation **forces** new literal assignments, pure literal elimination is more of a "rule of thumb"). If we had a pure literal, say 5, and we added in a clause containing -5 , unfortunately, 5 is no longer pure, and anything propagated because of 5 also would need to be undone because the reason why they were propagated/assigned those values in the first place doesn't hold anymore. Because we saw that pure literal elimination wasn't that helpful to begin with and with all these reasons as to why it wouldn't easily work under CDCL, it would just be much better to not have it at all.*

```

On all 1000 equations in uf20-91:

Clause Learning:           425ms (~0.43 seconds).
Iterative DPLL (AbP):      328ms (~0.33 seconds).
Iterative DPLL (PbA):      324ms (~0.32 seconds).
Watched Literals:         349ms (~0.35 seconds).
Pure Literal Elimination:  1548ms (~1.55 seconds).
Unit Propagation:          868ms (~0.87 seconds).
Recursive Backtracking:    1,540,613ms (~25.68 minutes).
Brute Force:               1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):   49,874ms (~50 seconds).
Parallel:                  10,589ms (~10.5 seconds).

On all 1000 equations in UUF50.218.1000:

Clause Learning:           9,652ms (~9.65 seconds).
Iterative DPLL (AbP):      6,369ms (~6.37 seconds).
Iterative DPLL (PbA):      6,365ms (~6.37 seconds).
Watched Literals:         7,130ms (~7.13 seconds).
Pure Literal Elimination:  231,107ms (~3.85 minutes).
Unit Propagation:          140,041ms (~2.34 minutes).

```

4.2 Non-Chronological Backjumping

An issue I mentioned earlier with our initial clause learning algorithm is that we backtrack only one level per contradiction. Not one level at a time, just one level. The reason why this is an issue is that we generally never get to use what we learned from conflict resolution to flip certain assignments, which is why we still relied on the DPLL notion of “trying the other branch.” Non-chronological backtracking, or “backjumping” is the new technique we’ll use to fully exploit the learned clause.

We only need to make a very minimal interface change. Now that we won’t use a “triedFalse” variable for our trail starts, our trail starts array can now just be an integer array where each element is just the index in the trail where a new decision level starts.

```

1      public interface:
2          SATSolver(string input)
3          bool solveCNF()
4
5      private interface:
6          ...
7          int[] trailStarts
8          ...

```

Next, since we’re only changing the backtracking logic, we can keep everything else the same in our algorithm. Before we move onto code, answer this: what decision level should be backjump to after we learn a clause? Suppose we learn this clause:

$$(1 \vee 2 \vee \neg 4)$$

And each literal's decision level is its index (1 was assigned on decision level 1, 2 was assigned on decision level 2, 4 was assigned on level 4). What decision level should be backjump to? Should we backjump to 1? 2? Stay on 4? We want to backjump to decision level 2 because that way, every literal in the clause other than 4 will be False, and 4 will be Unassigned. That way, we can propagate immediately on the learned clause! So the rule for backjumping after learning a new clause is to backjump to the second largest decision level present in the clause. Here's how we'll do it:

```

1  bool solve() {
2      var[1..numVars] assignment
3
4      for each (var in assignment) {
5          var = Unassigned
6      }
7
8      if (createWatched(assignment) == False) {
9          return False
10     }
11
12     int decisionLevel = 0
13
14     if (trail.size > 0) {
15         append(trailStarts, {0, true})
16
17         if (propagate(assignment, decisionLevel) == False) {
18             return False
19         }
20     }
21
22     while (True) {
23         if (trailHead == trail.size) {
24             //Perform an assignment
25             ...
26         }
27
28         optional int[] contradiction = propagate(assignment)
29
30         if (contradiction has a value) {
31             //resolve conflict
32             ...
33
34             if (conflict.size == 0) {
35                 return False
36             }
37
38             int maxLevel = decisionLevel
39             int secondLargest = -1

```

```

40         for (i = 0 up to conflict.size) {
41             int idx = abs(conflict[i])
42             int decisionLvl = trail[var_to_trail[idx]].
decisionLevel;
43             if (decisionLvl != maxLevel) {
44                 secondLargest = max(secondLargest, decisionLvl)
45             }
46         }
47
48         int backjumpTo = secondLargest
49
50         while (trail.size > 0 && trail.last.decisionLevel >
backjumpTo) {
51             int decision_start = trailStarts.last
52             while (trail.size > decision_start) {
53                 int lit = trail.last.lit
54                 int idx = abs(lit)
55                 assignment[idx] = Unassigned
56                 var_to_trail[idx] = -1
57                 erase(trail, trail.last)
58             }
59             erase(trailStarts, trailStarts.last)
60         }
61
62         if (trailStarts.size == 0) {
63             return False
64         }
65
66         decisionLevel = backjumpTo
67
68         trailHead = trailStarts.last
69
70         append(clauses, conflict)
71         createWatched(clauses.size)
72
73         continue
74     }
75 }
76 }

```

The code remains largely the same. However, we now want to check that the learned clause isn't empty so we can avoid doing large amounts of work. When can the learned clause become empty? To answer that question, let's answer this other question: "what does an empty clause mean?" From our section on Watched Literals, an empty clause means that the equation is unsatisfiable because no variable assignment can satisfy that clause. So if we learn an empty clause, that means that we're now learning that there is no satisfactory assignment of variables, thus our equation is unsatisfiable.

4.2.1 Non-Chronological Backjumping - Results

Okay so I know I said that we would make leaps forward in the next sections. I may have lied. Here are the results for Non-Chronological Backjumping:

uf20.91: **380ms**, down 45ms from Clause Learning.

UUF50.218.1000: **11291ms**, up 1639ms from Clause Learning.

Why did our time to solve unsatisfiable equations go up so significantly? It has to do with the fact of how we're learning. We're not trying every assignment like brute force, but we are slowly learning that no matter what the assignment is, our equation is unsatisfiable. This takes a lot of backtracking and memory to learn, possibly attributing to our time increase. It's hard to say for certain.

On all 1000 equations in uf20-91:

Backjumping:	380ms (0.38 seconds).
Clause Learning:	425ms (~0.43 seconds).
Iterative DPLL (AbP):	328ms (~0.33 seconds).
Iterative DPLL (PbA):	324ms (~0.32 seconds).
Watched Literals:	349ms (~0.35 seconds).
Pure Literal Elimination:	1548ms (~1.55 seconds).
Unit Propagation:	868ms (~0.87 seconds).
Recursive Backtracking:	1,540,613ms (~25.68 minutes).
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).


On all 1000 equations in UUF50.218.1000:

Backjumping:	11,291ms (~11.29 seconds).
Clause Learning:	9,652ms (~9.65 seconds).
Iterative DPLL (AbP):	6,369ms (~6.37 seconds).
Iterative DPLL (PbA):	6,365ms (~6.37 seconds).
Watched Literals:	7,130ms (~7.13 seconds).
Pure Literal Elimination:	231,107ms (~3.85 minutes).
Unit Propagation:	140,041ms (~2.34 minutes).

4.3 First UIPs

Our next addition to CDCL is First UIPs. UIP stands for "unit implication point," which is called a "dominator" in our implication graph. What is a dominator? A node d is said to *dominate* another node n if every path in the graph from the entry node to n passes through d . Think of a binary tree. Our entry node is the the head of the tree, and it dominates its two children. Likewise, each of those children dominates their own children nodes. So in other words, the UIP in our implication graph is kind of like the reason why we had a contradiction in the first place!

Given that summary of what a UIP is, wouldn't the first UIP be the decision we made on the current decision level? After all, we wouldn't have had a contradiction if we didn't make that decision. Well not necessarily, because in our implication graph, every decision we make only has outward edges, whereas all propagated nodes have inward edges as well. So while it is true that a first UIP could be a decision variable, it's not true that it's *always* a decision variable. Take this previous implication graph as an example:



x_4 is a UIP in our implication graph because every path to x_5 and x_6 all pass through x_4 's node.

So what does the first UIP offer us? Resolving clauses up until we encounter the first UIP maximizes the level we backjump to and it makes the clauses we learn shorter, not just in length, but also in how long it takes to resolve a clause. More details can be found here:

<https://www3.cs.stonybrook.edu/~cram/cse505/Fall20/Resources/cdcl.pdf>

How do we find the first UIP? The first UIP in a conflict is found when a the number of literals in the clause for the *current* decision level is 1. So if our current decision level is 6, we resolve only on literals from decision level 6 and stop once the clause contains only 1 literal from decision level 6.

4.3.1 Implementing First UIPs

Once again, we only need to touch the conflict resolution code:

```

1  bool solve() {
2      //setup
3      ...
4
5      while (True) {
6          if (trailHead == trail.size) {
7              //Perform an assignment
8              ...
9          }
10
11         optional int[] contradiction = propagate(assignment)
12     }
  
```

```

13         if (contradiction has a value) {
14             int curLevel = decisionLevel
15             int[] conflict = contradiction.value
16             int end = trail.size
17             int begin = 0
18             int resolveOn = 0
19
20             for (i = 0 up to conflict.size) {
21                 int idx = abs(conflict[i])
22                 int trail_idx = var_to_trail[idx]
23
24                 if (trail[trail_idx].decisionLevel == curLevel) {
25                     end = min(end, trail_idx)
26                     begin = max(begin, trail_idx)
27                     if (trail[trail_idx].reasonIdx is not optional) {
28                         resolveOn = trail[trail_idx].lit
29                     }
30                 }
31             }
32
33             while (begin != end) {
34                 int trail_idx = var_to_trail[abs(resolveOn)]
35                 int[] reason = clauses[trail[trail_idx].reasonIdx]
36                 append(conflict, reason)
37                 for (i = 0 up to conflict.size) {
38                     if (abs(conflict[i]) == abs(resolveOn)) {
39                         erase(conflict, conflict[i])
40                     } else {
41                         i = i + 1
42                     }
43                 }
44
45                 set<int> duplicates;
46                 for (i = 0 up to conflict.size) {
47                     if (contains(duplicates, conflict[i])) {
48                         erase(conflict, conflict[i])
49                     } else {
50                         duplicates.insert(conflict[i]);
51                         i = i + 1
52                     }
53                 }
54
55                 end = INTEGER_MAX;
56                 begin = 0;
57
58                 for (i = 0 up to conflict.size) {
59                     int var = conflict[i];
60                     int idx = abs(var)
61                     int trail_idx = var_to_trail[idx];
62

```

```

63         if (trail[trail_idx].decisionLevel
64             == decisionLevel)
65         {
66             end = min(end, trail_idx);
67             begin = max(begin, trail_idx);
68
69             if (trail[trail_idx].reasonIdx
70                 is not optional)
71             {
72                 resolve_var = trail[trail_idx].lit;
73             }
74         }
75     }
76 }
77
78 //backjump
79 ...
80
81 if (trailStarts.size == 0) {
82     return False
83 }
84
85 decisionLevel = backjumpTo
86
87 trailHead = trailStarts.last
88
89 append(clauses, conflict)
90 createWatched(clauses.size)
91
92 continue
93 }
94 }
95 }

```

This is not a very efficient way of computing the first UIP (which you'll see in the results). However, let's analyze it anyways. It's the normal resolution code, but we do something different, which is stop when there's only one literal from the current decision level left in the clause. In our pseudocode, we don't compute the counts of literals from the current decision level, we have the condition *begin*! = *end*. Why is that equivalent? Well *end* is the index of the least recently propagated literal in the resolution clause, and *begin* is the index of the most recently propagated literal in the resolution clause. So when the most recently and least recently propagated literals are the same, there must be only one left, therefore we can stop. Although it's clever, it's not efficient, and therefore, not very clever.

4.3.2 First UIPs - Results

Unfortunately our runtime increases again. At this point, I'm willing to accept that perhaps I was just really really good at implementing DPLL, but since I'm not good at implementing CDCL (I literally failed the first time), my implementation is not up to par. In the next section, I will identify possible improvements that can be made, and hopefully they'll make a *positive* difference.

uf20.91: **406ms**

UUF50.218.1000: **14012ms**

On all 1000 equations in uf20-91:

First UIPs:	406ms (~0.41 seconds)
Backjumping:	380ms (~0.38 seconds).
Clause Learning:	425ms (~0.43 seconds).
Iterative DPLL (AbP):	328ms (~0.33 seconds).
Iterative DPLL (PbA):	324ms (~0.32 seconds).
Watched Literals:	349ms (~0.35 seconds).
Pure Literal Elimination:	1548ms (~1.55 seconds).
Unit Propagation:	868ms (~0.87 seconds).
Recursive Backtracking:	1,540,613ms (~25.68 minutes).
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).

On all 1000 equations in UUF50.218.1000:

First UIPs:	14,012ms (~14.01 seconds).
Backjumping:	11,291ms (~11.29 seconds).
Clause Learning:	9,652ms (~9.65 seconds).
Iterative DPLL (AbP):	6,369ms (~6.37 seconds).
Iterative DPLL (PbA):	6,365ms (~6.37 seconds).
Watched Literals:	7,130ms (~7.13 seconds).
Pure Literal Elimination:	231,107ms (~3.85 minutes).
Unit Propagation:	140,041ms (~2.34 minutes).

4.4 A Second Round of Improvements

Ever since we started CDCL, I have been promising massive speedups over DPLL and yet our SAT solver has gotten slower each iteration. Before you treat me like Jesus and crucify me,³ I knew and deep down you knew that the way we were implementing was not optimal to begin with. So in a small way, this is your fault too.

³My friends told me that it was too funny not to include so it stays.

4.4.1 Removing Hashing

The first thing we’re going to do is remove the maps and instead use arrays instead so we don’t have to deal with hashing. Our interface becomes like so:

```

1      public interface:
2          SATSolver(string input)
3          bool solveCNF()
4
5      private interface:
6          ...
7          int[] [] var_to_clause
8          (int, int)[] clause_to_var

```

In case you need a reminder as it has been a while, these are the data structures used to make watched literals happen. Before they were maps where the keys were variables or clause indices, but now, each index in the arrays are the keys and the values at those indices are, well, the values. There is a small implementation detail concerning *clause_to_var*, which is that now that the number of clauses are not fixed, everytime we add in a new clause, we need to add in a new entry to *clause_to_var* before we let our *createWatched(int, int)* overload handle assigning the values to that entry. Basically, to avoid going out of bounds for the data structure, we need to change the bounds before we access the new boundary. While we could let *createWatched(int, int)* handle the allocation as well, it would make it less versatile as it would have to only operate on the last clause in the database instead of any clause in the database.

The second point to mention is that now our propagation logic doesn’t work because we need to create a mapping from literals (positive and negative) to indices (nonnegative). We use the following mapping:

$$2 * \text{abs}(\text{lit}) + (\text{lit} < 0)$$

This is how the mapping works out (left column is literal value, right column is index):

1	→	0
-1	→	1
2	→	2
-2	→	3
3	→	4
-3	→	5
...		

The last step is to remove the duplicates set we use in clause resolution. We’re going to add in an array of booleans to our SAT solver where index *i* is whether or not we’ve seen that variable during our clause resolution. However, since our (my) implementation of first UIP clause resolution is such a big fixer-upper, it deserves its own subsection.

4.4.2 Fixing Clause Resolution

Here is our current clause resolution code:

```

1  if (contradiction has a value) {
2      int curLevel = decisionLevel
3      int[] conflict = contradiction.value
4      int end = trail.size
5      int begin = 0
6      int resolveOn = 0
7
8      for (i = 0 up to conflict.size) {
9          int idx = abs(conflict[i])
10         int trail_idx = var_to_trail[idx]
11
12         if (trail[trail_idx].decisionLevel == curLevel) {
13             end = min(end, trail_idx)
14             begin = max(begin, trail_idx)
15             if (trail[trail_idx].reasonIdx is not optional) {
16                 resolveOn = trail[trail_idx].lit
17             }
18         }
19     }
20
21     while (begin != end) {
22         int trail_idx = var_to_trail[abs(resolveOn)]
23         int[] reason = clauses[trail[trail_idx].reasonIdx]
24         append(conflict, reason)
25         for (i = 0 up to conflict.size) {
26             if (abs(conflict[i]) == abs(resolveOn)) {
27                 erase(conflict, conflict[i])
28             } else {
29                 i = i + 1
30             }
31         }
32
33         set<int> duplicates;
34         for (i = 0 up to conflict.size) {
35             if (contains(duplicates, conflict[i])) {
36                 erase(conflict, conflict[i])
37             } else {
38                 duplicates.insert(conflict[i]);
39                 i = i + 1
40             }
41         }
42
43         end = INTEGER_MAX;
44         begin = 0;
45
46         for (i = 0 up to conflict.size) {

```

```

47         int var = conflict[i];
48         int idx = abs(var)
49         int trail_idx = var_to_trail[idx];
50
51         if (trail[trail_idx].decisionLevel
52             == decisionLevel)
53         {
54             end = min(end, trail_idx);
55             begin = max(begin, trail_idx);
56
57             if (trail[trail_idx].reasonIdx
58                 is not optional)
59             {
60                 resolve_var = trail[trail_idx].lit;
61             }
62         }
63     }
64 }
65 }

```

It's pretty long and inefficient. This was written before we understood implication graphs fully. Now we do understand implication graphs, so let's exploit that structure.

The idea is this: instead of rescanning over and over to make sure our resolution clause has just one current level variable, we'll keep track of which variables we've seen as well as which ones are from this decision level. Then, as we perform clause resolution, we'll update the seen and count variables to reflect the current state of the resolution clause, then once we've encountered our stopping condition, we'll use which variables we've kept marked as seen to construct our learned clause directly. This is very different from the way we've been performing clause resolution, which involved repeatedly scanning the resolution clause and recomputing how far along we are in the UIP process.

How are we using the implication graph in our new method? Well since the implication graph is an acyclic graph, that means we can use its structure for our UIP traversal. There is a relationship between the implication graph and our trail, in that our trail is just a compressed version of the implication graph, meaning it doesn't keep track of which variable assignments led to which propagations. However, they both follow the same structure in that the edge of our trail are the edges of our implication graphs and the order of variables in the trail and graphs are the same. Because the implication graph is acyclic, once we encounter a variable during clause resolution, we will never encounter it again while we're resolving. Thanks to that property, we can traverse the trail backwards without having to worry about going over the same parts of the trail/graph, which was something that we were sort of doing in the old method when we rescanned the whole conflict clause, whereas now, we are only concerned with the differences to the resolution clause made across each iteration.

Our resolution algorithm becomes as follows:

```

1  if (contradiction has a value) {

```

```

2      int[] conflict = contradiction.value
3      int levelCount = 0
4      int curLevel = decisionLevel
5
6      for each (int var in conflict) {
7          if (!seen[abs(var)]) {
8              seen[abs(var)] = True
9
10             if (trail[var_to_trail[abs(var)]].decisionLevel ==
11                 curLevel) {
12                 levelCount = levelCount + 1
13             }
14         }
15     }
16
17     for (i = trail.size down to 0) {
18         TrailEntry entry = trail[i]
19         idx = abs(entry.lit)
20
21         if (seen[idx]) {
22             int[] reason = clauses[trail[i].reasonIdx]
23
24             for each (int var in reason) {
25                 if (!seen[abs(var)]) {
26                     seen[abs(var)] = True
27
28                     int trailIdx = var_to_trail[abs(var)]
29
30                     if (trail[trailIdx].decisionLevel == curLevel) {
31                         levelCount = levelCount + 1
32                     }
33                 }
34             }
35
36             seen[idx] = False
37             levelCount = levelCount - 1
38         }
39
40         if (levelCount == 1) {
41             break
42         }
43     }
44
45     int[] learned
46     for (i = 1 up to numVars) {
47         if (seen[i] == True) {
48             int lit = trail[var_to_trail[i]].lit
49             append(learned, -lit)
50             seen[i] = False

```

```

51     }
52 }

```

Here's the rundown of what's going on. In the first block, we're marking which variables we've seen and how many of them are from the current decision level from the initial conflicting clause. Next, we start traversing our trail from the most recent trail entry, and we move towards the front of the trail. If we've seen the variable from the current trail entry, we grab its reason clause, and for each new unseen variable, we mark it as seen and increment the current decision level count. Next, we "remove" the variable we just processed by unmarking it as seen and decrementing *levelCount*.

Once we have only one literal from the current level, we construct a new learned clause by iterating over the entire *seen* array and we add the negative literal to the learned clause because the learned clause is supposed to learn the opposite of what we've done so we don't repeat it. Then we unmark everything we have seen so that way we can reuse the same array the next time there's a conflict instead of allocating a new one each time.

The results for this are very good, but I won't tell you what they are, because there's one more thing I want to do. You see, we iterate over the entire *seen* array, which is as large as the number of variables in our equation. However, it's very unlikely that we'll see many variables because first UIPs shrink how large the learned clause is, so there's less to be seen in the first place, but also because the more we learn, generally, the learned clause gets smaller because each new learned clause adds more constrictions on which assignments are valid, and as we constrict which ones are valid and which ones aren't, we inch towards learning the empty clause signaling that nothing is valid. This is just a long winded way to say that the iterating over the entire *seen* array each time, while $O(n)$, could be better.

But isn't that microoptimizing? ⁴ After all, our two test sets have 20 and 50 variables respectively, which isn't much in the grand scheme of things. And you're right, our test sets are pretty small. However, the goal is to build a SAT solver that would scale well and most SAT solvers can handle problems with millions of variables pretty well, so I'd like this one to too. Making it scale well isn't the same as microoptimizing. Microoptimizing would be like saying `++i` over `i++`. ⁵

In order to do this, we're going to keep track of the indices of which variables we've seen in an array and then iterate over that array. We will also need to keep track of which indices are in which position so we can swap and pop for efficient removal of elements. ⁶ Our object now looks like this:

```

1     public interface:
2         SATSolver(string input)
3         bool solveCNF()
4

```

⁴Microoptimizing is like cutting your hair to help you reach a weight loss goal, instead of fixing your running form so you'll be able to have better cardio.

⁵Pre vs post incrementation in loops generally does not affect runtime and whether or not it actually matters is heavily debated.

⁶I love swap and pop if only for the name. I encourage everyone to do swap and pops so we can say "swap and pop" more than we currently say "swap and pop."

```

5      private interface:
6          ...
7          int[][] var_to_clause
8          (int, int)[] clause_to_var
9
10         bool[] seen
11         int[] seenIdx
12         int[] seenPos

```

This is how these data structures change the algorithm:

```

1  if (contradiction has a value) {
2      int[] conflict = contradiction.value
3      int levelCount = 0
4      int curLevel = decisionLevel
5
6      for each (int var in conflict) {
7          if (!seen[abs(var)]) {
8              seen[abs(var)] = True
9              append(seenIdx, abs(var))
10             seenPos[abs(var)] = seenIdx.size
11
12             if (trail[var_to_trail[abs(var)]].decisionLevel ==
13                 curLevel) {
14                 levelCount = levelCount + 1
15             }
16         }
17     }
18
19     for (i = trail.size down to 0) {
20         TrailEntry entry = trail[i]
21         idx = abs(entry.lit)
22
23         if (seen[idx]) {
24             int[] reason = clauses[trail[i].reasonIdx]
25
26             for each (int var in reason) {
27                 if (!seen[abs(var)]) {
28                     seen[abs(var)] = True
29                     append(seenIdx, abs(var))
30                     seenPos[abs(var)] = seenIdx.size

```

```

31         int trailIdx = var_to_trail[abs(var)]
32
33         if (trail[trailIdx].decisionLevel == curLevel) {
34             levelCount = levelCount + 1
35         }
36     }
37 }
38
39     seen[idx] = False
40
41     //get the position of idx in seenIdx
42     int pos = seenPos[idx]
43
44     //get the last element in seenIdx
45     int last = seenIdx.last
46
47     //overwrite the element at the position pos with the
value
48     // of the last element in seenIdx
49     seenIdx[pos] = last
50
51     //mark that the last element in seenIdx is now at
position pos
52     seenPos[last] = pos
53
54     //remove the last element in seenIdx because it's also at
position pos
55     pop(seenIdx)
56
57     levelCount = levelCount - 1
58 }
59
60     if (levelCount == 1) {
61         break
62     }
63 }
64
65     int[] learned
66     while (seenIdx is not empty) {
67         int lit_idx = seenIdx.last
68
69         int lit = trail[var_to_trail[lit_idx]].lit
70         append(learned, -lit)
71
72         seen[lit_idx] = false
73         pop(seenIdx)
74         //do nothing with seenPos because we don't have to
75     }
76 }

```


4.4.3 Improvements - Results

I'll let the results speak for themselves: uf20.91: **264ms**

UUF50.218.1000: **2385ms**

On all 1000 equations in uf20-91:

Improved CDCL:	284ms (~0.28 seconds).
First UIPs:	406ms (~0.41 seconds).
Backjumping:	380ms (~0.38 seconds).
Clause Learning:	425ms (~0.43 seconds).
Iterative DPLL (AbP):	328ms (~0.33 seconds).
Iterative DPLL (PbA):	324ms (~0.32 seconds).
Watched Literals:	349ms (~0.35 seconds).
Pure Literal Elimination:	1548ms (~1.55 seconds).
Unit Propagation:	868ms (~0.87 seconds).
Recursive Backtracking:	1,540,613ms (~25.68 minutes).
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).

On all 1000 equations in UUF50.218.1000:

Improved CDCL:	2,385ms (~2.39 seconds).
First UIPs:	14,012ms (~14.01 seconds).
Backjumping:	11,291ms (~11.29 seconds).
Clause Learning:	9,652ms (~9.65 seconds).
Iterative DPLL (AbP):	6,369ms (~6.37 seconds).
Iterative DPLL (PbA):	6,365ms (~6.37 seconds).
Watched Literals:	7,130ms (~7.13 seconds).
Pure Literal Elimination:	231,107ms (~3.85 minutes).
Unit Propagation:	140,041ms (~2.34 minutes).

4.5 VSIDS

Variable State Independent Decaying Sum (VSIDS) is a type of variable assignment heuristic, basically a loosely defined rule, that is used in modern CDCL solvers. It is *the* best variable assignment/branching heuristic with many variations. In this chapter, we follow this paper

<https://arxiv.org/pdf/1506.08905>

which gives a much more in depth and theoretical explanation of how and why they work. For our purposes, we'll understand at a higher level what VSIDS are as well as how and why they work.

Currently, our variable assignment choosing heuristic has been “find first Unassigned variable, and assign it True.” With VSIDS, this becomes “find the most active

Unassigned variable, and assign it True.” The first step in VSIDS is to add in an array of **activity scores** which rank “how active” a variable is. There are two schemes for ranking activity score, which is ranking each literal (positive and negative variables), or just the variables (absolute value of the literal). For our purposes, we will rank just the variables.

The next step is after clause resolution, add to each variable’s score if it was involved in clause resolution. Again, there are multiple ways to do this. Some implementations choose only the variables in the learned clause, some implementations choose every variable involved in the resolution algorithm, even if they are resolved away, and there are more implementations than just those two. For our purposes, we will choose bumping the scores of only the variables in the learned clause.

The last step, which doesn’t happen every iteration, is to multiply each score by a decay factor ⁷ after a certain number of conflicts. The idea behind this is to even the playing field for each variable. If a variable accumulates a giant activity score, obviously it will be the first pick if it’s valid. However, activity scores don’t represent how recent the activity score bump occurred, so if that variable stopped being active, it would still be the first pick despite not actually being active relative to the learned clauses. This decay lets the variables that are starting to become more active have a better chance at competing with the potentially inactive, but score-dominating variables. As it is a decaying factor, this value is chosen to be greater than 0, but less than 1, so that the scores shrink to some percentage.

4.5.1 cVSIDS

We’re going to start out with implementing the most basic form of VSIDS and work our way up to some better versions. We’ll start with the “textbook” form, which was outlined in the section’s introduction. This form of VSIDS is called **cVSIDS** by the paper we’re using. We first add in our *activity* array into our solver object with a counter for the number of conflicts.

```

1      public interface:
2          SATSolver(string input)
3          bool solveCNF()
4
5      private interface:
6          ...
7          double[] activity
8          int numConflicts
9          ...

```

Then our solve code becomes as follows:

```

1      bool solve() {
2          //setup
3          ...

```

⁷Not to be confused with Deadpool’s “dying factor.”

```

4
5     while (True) {
6         if (trailHead == trail.size) {
7             double maxScore = -1.0
8             optional int firstUnassigned
9
10            for (i = 1 up to numVars) {
11                if (assignment[i] == Unassigned) {
12                    if (activity[i] > maxScore) {
13                        firstUnassigned = i
14                        maxScore = activity[i]
15                    }
16                }
17            }
18
19            //the rest of the assignment code is the same
20            ...
21        }
22
23        optional int[] contradiction = propagate(assignment)
24
25        if (contradiction has a value) {
26            numConflicts = numConflicts + 1
27
28            //resolve conflict using UIPs
29            ...
30            int[] learned
31            while (seenIdx is not empty) {
32                int lit_idx = seenIdx.last
33                activity[lit_idx] = activity[lit_idx] + 1
34                ...
35            }
36
37            //backjump
38            ...
39
40            //apply decay factor
41            if (numConflicts == X) {
42                for (i = 0 up to activity.size) {
43                    activity[i] *= Y
44                }
45                numConflicts = 0
46            }
47
48            continue
49        }
50    }
51 }

```

The pseudocode is very straightforward, so I won't bother to explain it much. Just so you know, if you choose to initialize *maxScore* with a value of 0.0 in the assignment code, use `>=` instead of `>` because the activity scores start out at 0.0, so not using `>=` will mess with unsatisfiable problems. In our code to apply the decay factor, I use “X” and “Y” to represent the number of conflicts and the decay factor respectively. The reason why these aren't set values, but rather placeholders is because there aren't two numbers that perform the best for *all* equations, so they should be adjusted programmatically or manually set experimentally. However, I just decided to choose two random numbers, so I used **X = 10** and **Y = 0.8**. But results for different X and Y will be shown as well.

The results were as follows: uf20.91: **279ms**

UUF50.218.1000: **1841ms**

We didn't get a big speedup for satisfiable problems, but we did shave off about half a second for the unsatisfiable problems, which I think is pretty darn good.

```
On all 1000 equations in uf20-91:

cVSIDS (X = 1, Y = 0.5):    260ms (~0.26 seconds).
cVSIDS (X = 1, Y = 0.8):    266ms (~0.27 seconds).
cVSIDS (X = 10, Y = 0.5):   270ms (~0.27 seconds).
cVSIDS (X = 10, Y = 0.8):   279ms (~0.29 seconds).
Improved CDCL:              284ms (~0.28 seconds).

On all 1000 equations in UUF50.218.1000:

cVSIDS (X = 1, Y = 0.5):    1,680ms (~1.68 seconds).
cVSIDS (X = 1, Y = 0.8):    1,705ms (~1.71 seconds).
cVSIDS (X = 10, Y = 0.5):   1,782ms (~1.78 seconds).
cVSIDS (X = 10, Y = 0.8):   1,841ms (~1.84 seconds).
Improved CDCL:              2,385ms (~2.39 seconds).
```

In the results, since we will be iterating over many versions of VSIDS, I will leave out the older results and only compare against the previous versions of our solver. In the results, it appears to be that lower X and Y values lead to faster solving times, however, I encourage you to choose your own X and Y values and see what happens for yourself.

4.5.2 mVSIDS

mVSIDS is the second form of VSIDS introduced by the paper from the section introduction. In this type of VSIDS, we bump the values of every variable we resolve on during clause resolution as well as the values of the variables in the learned clause. The multiplicative decay step is applied the same as during cVSIDS.

The algorithmic change is basically trivial:

```
1 bool solve() {
2     //setup
3     ..
```

```

4
5     while (True) {
6         //assigning
7         ...
8
9         optional int[] contradiction = propagate(assignment)
10
11        if (contradiction has a value) {
12            numConflicts = numConflicts + 1
13
14            //set up UIPs from conflict clause
15            ...
16
17            for (i = trail.size down to 0) {
18                TrailEntry entry = trail[i]
19                idx = abs(entry.lit)
20
21                if (seen[idx]) {
22                    //grab reason and update seen and levelCount
23                    ...
24
25                    activity[idx] += 1
26                    //unmark seen and update levelCount
27                    ...
28                }
29
30                if (levelCount == 1) {
31                    break
32                }
33            }
34
35            //create learned clause and update activity scores
36            ...
37
38            //backjump
39            ...
40
41            //apply decay factor
42            ...
43
44            continue
45        }
46    }
47 }

```

Here are the results using **X = 10** and **Y = 0.8**:

On all 1000 equations in uf20-91:	
mVSIDS (X = 1, Y = 0.5):	279ms (~0.28 seconds).
mVSIDS (X = 1, Y = 0.8):	277ms (~0.28 seconds).
mVSIDS (X = 10, Y = 0.5):	269ms (~0.27 seconds).
mVSIDS (X = 10, Y = 0.8):	284ms (~0.28 seconds).
cVSIDS (X = 1, Y = 0.5):	260ms (~0.26 seconds).
cVSIDS (X = 1, Y = 0.8):	266ms (~0.27 seconds).
cVSIDS (X = 10, Y = 0.5):	270ms (~0.27 seconds).
cVSIDS (X = 10, Y = 0.8):	279ms (~0.29 seconds).
Improved CDCL:	284ms (~0.28 seconds).
On all 1000 equations in UUF50.218.1000:	
mVSIDS (X = 1, Y = 0.5):	1,532ms (~1.53 seconds).
mVSIDS (X = 1, Y = 0.8):	1,526ms (~1.53 seconds).
mVSIDS (X = 10, Y = 0.5):	1,420ms (~1.42 seconds).
mVSIDS (X = 10, Y = 0.8):	1,467ms (~1.47 seconds).
cVSIDS (X = 1, Y = 0.5):	1,680ms (~1.68 seconds).
cVSIDS (X = 1, Y = 0.8):	1,705ms (~1.71 seconds).
cVSIDS (X = 10, Y = 0.5):	1,782ms (~1.78 seconds).
cVSIDS (X = 10, Y = 0.8):	1,841ms (~1.84 seconds).
Improved CDCL:	2,385ms (~2.39 seconds).

So even with the worst mVSIDS time, it beats the best cVSIDS time (per the benchmark, results may vary, these tests are very small and are not very representative of real-world SAT solver usage).

4.5.3 EMA/Normalized VSIDS

EMA stands for “Exponential Moving Average” and it’s what allows us to mathematically decide which variables’ activity scores should be decayed and by how much. If we go back to the earlier example of some variable gaining a really high activity score and then becoming inactive, but still dominating over the more recently active variables, this is what EMA helps with. EMA will decay high scoring, but actually inactive variables’ scores while mostly preserving the actually active variables’ scores.

The way it does it is with this nifty formula:

$$s_n = (1 - f)\delta_n + f * s_{n-1}$$

s_n is the **current score**, f is our **decay factor**, $\delta_n \in \{0, 1\}$ where a value of 0 corresponds to a variable not being bumped, and a value of 1 corresponds to a variable being bumped. Let’s walk through an example:

Suppose $f = 0.95$, and $s_0 = 0.0$. In the first conflict ($n = 1$), a variable was involved. Its score becomes:

$$s_1 = (1 - 0.95)(1) + 0.95(0.0) = 0.05$$

It's not much, but it's much more than 0, so this variable is considered high ranking. Now suppose in the next conflict ($n = 2$), the variable was not involved. Its score now becomes:

$$s_2 = (1 - 0.95)(0) + 0.95(0.05) = 0.0475$$

This is less than 0.05, so now it's low ranking, even if the difference is small. Now suppose it's involved in the third conflict ($n = 3$):

$$s_3 = (1 - 0.95)(1) + 0.95(0.0475) = 0.095125$$

This is much larger than what it was before.

Although you probably have guessed already, we've been implementing VSIDS in a straightforward naive way by following the described process down to the letter. For sake of simplicity, for Normalized VSIDS, we'll also implement it naively.

Implementing normalized VSIDS isn't too different to how we're already implementing it, except now, we update every score upon each conflict iteration. We'll need an additional data structure to keep track of the δ values. Since we're appending normalization on top of mVSIDS, we still update a δ value only when we resolve a variable or it appears in the learned clause. This means that we now delay the updating of our activity scores until after we construct the learned clause:

```

1  bool solve() {
2      //setup
3      ..
4
5      while (True) {
6          //assigning
7          ...
8
9          optional int[] contradiction = propagate(assignment)
10
11         if (contradiction has a value) {
12             numConflicts = numConflicts + 1
13
14             //set up UIPs from conflict clause
15             ...
16
17             for (i = trail.size down to 0) {
18                 TrailEntry entry = trail[i]
19                 idx = abs(entry.lit)
20
21                 if (seen[idx]) {
```

```

22         //grab reason and update seen and levelCount
23         ...
24
25         delta[idx] = True
26         //unmark seen and update levelCount
27         ...
28     }
29
30     if (levelCount == 1) {
31         break
32     }
33 }
34
35 //create learned clause and update activity scores
36 int[] learned
37 while (seenIdx is not empty) {
38     int lit_idx = seenIdx.last
39     delta[lit_idx] = True
40     ...
41 }
42
43 //backjump
44 ...
45
46 //apply decay factor
47 for (i = 1 up to numVars) {
48     if (delta[i] == True) {
49         activity[i] = (1 - Y) + Y * activity[i]
50     } else {
51         activity[i] = Y * activity[i]
52     }
53     delta[i] = False
54 }
55
56     continue
57 }
58 }
59 }

```

Notice how that there is no X value anymore because normalization happens upon every iteration instead of after a set number. Thankfully, our fastest times were when we updated activity scores after every iteration anyways, so we probably weren't missing some "fastest" time under a different X value.

Here are the results:

On all 1000 equations in uf20-91:	
EMA (X = 1, Y = 0.95):	266ms (~0.27 seconds).
EMA (X = 1, Y = 0.8):	287ms (~0.29 seconds).
EMA (X = 1, Y = 0.5):	280ms (~0.28 seconds).
mVSIDS (X = 1, Y = 0.5):	279ms (~0.28 seconds).
mVSIDS (X = 1, Y = 0.8):	277ms (~0.28 seconds).
mVSIDS (X = 10, Y = 0.5):	269ms (~0.27 seconds).
mVSIDS (X = 10, Y = 0.8):	284ms (~0.28 seconds).
cVSIDS (X = 1, Y = 0.5):	260ms (~0.26 seconds).
cVSIDS (X = 1, Y = 0.8):	266ms (~0.27 seconds).
cVSIDS (X = 10, Y = 0.5):	270ms (~0.27 seconds).
cVSIDS (X = 10, Y = 0.8):	279ms (~0.29 seconds).
Improved CDCL:	284ms (~0.28 seconds).
On all 1000 equations in UUF50.218.1000:	
EMA (X = 1, Y = 0.95):	1,504ms (~1.50 seconds).
EMA (X = 1, Y = 0.8):	1,497ms (~1.50 seconds).
EMA (X = 1, Y = 0.5):	1,621ms (~1.62 seconds).
mVSIDS (X = 1, Y = 0.5):	1,532ms (~1.53 seconds).
mVSIDS (X = 1, Y = 0.8):	1,526ms (~1.53 seconds).
mVSIDS (X = 10, Y = 0.5):	1,420ms (~1.42 seconds).
mVSIDS (X = 10, Y = 0.8):	1,467ms (~1.47 seconds).
cVSIDS (X = 1, Y = 0.5):	1,680ms (~1.68 seconds).
cVSIDS (X = 1, Y = 0.8):	1,705ms (~1.71 seconds).
cVSIDS (X = 10, Y = 0.5):	1,782ms (~1.78 seconds).
cVSIDS (X = 10, Y = 0.8):	1,841ms (~1.84 seconds).
Improved CDCL:	2,385ms (~2.39 seconds).

Here we see that EMA/normalizing the VSIDS doesn't make a huge contribution, but it still makes a decent one for some values of Y.

4.5.4 AMA/adaptVSIDS

Adaptive Moving Average (AMA) VSIDS is an extension to EMA VSIDS where rather than having a fixed decay factor, we choose between two different decay rates depending on the state of our learned clause. The big idea behind making the decay factor variable instead of fixed is that our solver's "productivity" isn't fixed. The learned clause affects this productivity in how high the quality of the learned clause is.

To figure out the quality of a learned clause, we have to look at its **LBD** which stands for **literal blocks distance**. LBD is defined as the number of distinct decision levels in the learned clause. A learned clause with a low LBD is considered to be higher quality than that of one with a high LBD. The paper doesn't explicitly say why this is the case,⁸ but we can make a few guesses as to why. Considering that a low LBD means that there aren't many decision levels involved in the contradiction, so that means that the contradiction is more or less "confined" to a smaller area of the search space. This indicates that the conflict has a more localized relationship with the current set of decisions. Such clauses have a higher chance to become unit or nearly unit, under a smaller amount of changes to the current assignment, better constraining future search more effectively. Contrast this to high LBD, which has multiple decision levels involved in the contradiction, meaning that the contradiction spans a broader portion of the current search space, making it less likely it'll have an immediate effect on constraining future search.

This is supported by the fact that when we learn a clause with a low LBD, we choose a higher decay factor, so that variable activity scores shrink less, meaning we remember them more in the present. Likewise, when we learn a clause with a high LBD, we choose a lower decay factor, so variable activity scores shrink faster, making older activity lose its influence faster relative to the newer active variables.

The paper introduces this variable called *lbdema*, which is defined as the exponential moving average of the LBD. Like our EMA VSIDS, this is calculated upon every iteration, and it's also a recurrence relation where the previous value affects what the next value will be. The formula is similar to the previous EMA formula:

$$lbdema_n = (1 - \alpha) * LBD_t + \alpha * lbdema_{n-1}$$

α is our **smoothing factor**, LBD_t is the LBD of the current learned clause, and $lbdema_{n-1}$ is the previous *lbdema*. In our testing of AMA VSIDS, I used $\alpha = 0.1$.

4.5.5 VSIDS - Results

The results for VSIDS vary wildly because there are so many different types beyond what we've explored, and there are so many different parameters that could affect the end results. For VSIDS, although there might look like there's a clear winner, we need to keep in mind that we're using two relatively small test sets, so drawing an actual conclusion as to which VSIDS version is "objectively" better isn't credible. Although AMA VSIDS isn't the fastest from our benchmarks, it will be the version we continue with for the rest of our paper. As for code, given how easy it is to swap out each version of VSIDS with each other, the exact code used for each version won't be provided, as the total changes were small enough that following along shouldn't be overly difficult.

Another thing to note is that at the moment, there was not much exploration for the time differences. In order to figure out why the speeds between each version was different, it would be helpful to analyze how many conflicts there were, as well as the order of deciding variables. As of now, I am just looking at the results and just accepting them as they are, not because I lack curiosity or scientific rigor, but because this is the

⁸I might be illiterate and may have just missed the reason.

second night in a row where I barely got any sleep for no reason, and I am too tired to analyze this under a microscope.

On all 1000 equations in uf20-91:	
AMA (Y = 0.75 or 0.99):	274ms (~0.27 seconds).
EMA (X = 1, Y = 0.95):	266ms (~0.27 seconds).
EMA (X = 1, Y = 0.8):	287ms (~0.29 seconds).
EMA (X = 1, Y = 0.5):	280ms (~0.28 seconds).
mVSIDS (X = 1, Y = 0.5):	279ms (~0.28 seconds).
mVSIDS (X = 1, Y = 0.8):	277ms (~0.28 seconds).
mVSIDS (X = 10, Y = 0.5):	269ms (~0.27 seconds).
mVSIDS (X = 10, Y = 0.8):	284ms (~0.28 seconds).
cVSIDS (X = 1, Y = 0.5):	260ms (~0.26 seconds).
cVSIDS (X = 1, Y = 0.8):	266ms (~0.27 seconds).
cVSIDS (X = 10, Y = 0.5):	270ms (~0.27 seconds).
cVSIDS (X = 10, Y = 0.8):	279ms (~0.29 seconds).
Improved CDCL:	284ms (~0.28 seconds).
On all 1000 equations in UUF50.218.1000:	
AMA (Y = 0.75 or 0.99):	1,629ms (~1.63 seconds).
EMA (X = 1, Y = 0.95):	1,504ms (~1.50 seconds).
EMA (X = 1, Y = 0.8):	1,497ms (~1.50 seconds).
EMA (X = 1, Y = 0.5):	1,621ms (~1.62 seconds).
mVSIDS (X = 1, Y = 0.5):	1,532ms (~1.53 seconds).
mVSIDS (X = 1, Y = 0.8):	1,526ms (~1.53 seconds).
mVSIDS (X = 10, Y = 0.5):	1,420ms (~1.42 seconds).
mVSIDS (X = 10, Y = 0.8):	1,467ms (~1.47 seconds).
cVSIDS (X = 1, Y = 0.5):	1,680ms (~1.68 seconds).
cVSIDS (X = 1, Y = 0.8):	1,705ms (~1.71 seconds).
cVSIDS (X = 10, Y = 0.5):	1,782ms (~1.78 seconds).
cVSIDS (X = 10, Y = 0.8):	1,841ms (~1.84 seconds).
Improved CDCL:	2,385ms (~2.39 seconds).

4.6 Clause Database Management

Clause database management is another essential addition to a CDCL SAT solver. For larger and larger equations, and for equations with less valid assignments, we will tend to learn more clauses. While First UIPs can efficiently compute these learned clauses, with the added benefit of them being smaller than if we didn't use the first UIP condition, the number of learned clauses still has a memory impact. Although our test sets are of 3-SAT instances, so the initial clauses are small to begin with, and the learned clauses aren't as large as they could be, for larger SAT problems, the number of learned clauses can become exponential in size.

Although some people like to say that memory management isn't important because the hardware gets better and better, the day we have to stop worrying about the amount of memory being used is the day we have unlimited memory to work with. And until that day comes (which will be never), those people can shut up and go back to coding their 2GB Chrome tab website.⁹

For this section, we will go back to using this paper:

<https://www3.cs.stonybrook.edu/~cram/cse505/Fall120/Resources/cdcl.pdf>

starting from section **4.5.4. Clause Deletion Strategies**. The three (listed) methods described are:

- *n-order learning*: Only learning clauses with n or fewer literals.
- *m-size relevance-based learning*: Learning all clauses as long as they currently imply assignments or are themselves a unit clause, only being only being discarded once they have m or more Unassigned literals.
- *k-bounded learning*: Clauses smaller than k are kept permanently, but larger clauses are discarded once they stop being unit.

The paper also mentions that *k-bounded learning* and *m-size relevance-based learning* can be combined. You may wonder why we can't combine all three, however, notice that *k-bounded learning* is an extension of *n-order learning*, as they both keep clauses smaller than some threshold, but *k-bounded learning* is more merciful to larger clauses.

The paper also describes a newer heuristic that is based on its *activity* score. The heuristic is essentially "VSIDS for clauses instead of variables." We will progress our way into implementing all of these algorithms, incrementally of course. Before we continue, it should be noted that modern SAT solvers don't have a universal clause deletion policy, and each one does its own thing, and once ours is mature enough, it will be no different.¹⁰

4.6.1 N-Order Learning

n-order learning is definitely the most aggressive clause learning/deletion policy of the bunch. I don't think I'd want to encounter it on the streets. It's very straightforward in

⁹Websites take up too much memory, much more than they need, and I will happily die on this hill.

¹⁰Does that make us all unique or all the same?

idea: if a clause has more than n literals, it won't get learned at all. This may raise some red flags, and I hope the one we're all thinking of is, how do we backjump? If we're not learning a clause, how do we know which decision level to backjump to? Well, just because we're not learning it doesn't mean we can't backjump to a higher decision level, it just means we won't add the new clause to our database. Does that put you at ease? Well it shouldn't because the reason why this was such a big red flag in the first place is that as of now, when we decide a new variable, we always start with True and rely on clause learning to force it False later. So when we backjump and start a new decision level, how do we start the new decision level as False if we didn't learn anything?

As of now, there is no way to change our course.¹¹ So what can we do? You're gonna find this funny (not really), we need to bring back our "tried second branch" flag, or at least some form of it. Without it, and I cannot emphasize this enough, we will just be stuck in an infinite loop.

So how should we pick a value of n ? There isn't a universal way to choose, so we will try the maximum clause length of the original CNF, and then try some numbers around it. Because our current testing sets are all instances of 3-SAT, it's not very fair, in my opinion, to test on different sizes of n when the data is specifically tailored to be a fixed size. Because of this, we will go to our CNF database¹² and get a new testing set, "Flat200-479" which is a SAT encoding of a 3-color problem. There are 600 variables and 2237 clauses, all satisfiable.¹³

Here are the baseline results (without n-order learning):

On all 1000 equations in uf20-91:	
AMA:	274ms (~0.27 seconds).
On all 100 equations in Flat200-479:	
AMA:	24,238ms (~24.24 seconds).
On all 1000 equations in UUF50.218.1000:	
AMA:	1,629ms (~1.63 seconds).

We will now discuss our changes to our algorithm. Firstly, we need to compute the maximum clause length. As of now, we will say that the maximum clause length of the original equation will also be the maximum clause length of any learned clause, and then see how changing this number to lower or higher changes our runtime speed. To compute the maximum clause length, we will iterate through the original clauses and take the maximum size.

Next, we want to check how large the learned clause will be. We do this by checking the size of the *seenIdx* array because it holds the variables that will become part of the

¹¹This will be addressed in the next section.

¹²<https://www.cs.ubc.ca/~hoos/SATLIB/benchm.html>

¹³As I was setting up the benchmarking test, my framework reported that the solver failed on this new test set. It turned out I gave it the wrong path, and after providing it the correct path, it passed, as we would expect. The shock of seeing the [FAIL] really threw me for a loop.

We first need to remember to change our interface to use our old `DecisionLevel` struct:

```

1      public interface:
2
3      private interface:
4          DecisionLevel {
5              int idx,
6              bool triedOtherBranch
7          }
8
9          DecisionLevel[] trailStarts
10         ...

```

The algorithmic changes are as follows:

```

1  bool solve() {
2      //setup
3      ..
4
5      while (True) {
6          //assigning
7          ...
8
9          optional int[] contradiction = propagate(assignment)
10
11         if (contradiction has a value) {
12             numConflicts = numConflicts + 1
13
14             //set up UIPs from conflict clause
15             ...
16
17             for (i = trail.size down to 0) {
18                 TrailEntry entry = trail[i]
19                 idx = abs(entry.lit)
20
21                 if (seen[idx]) {
22                     //grab reason and update seen and levelCount
23                     ...
24
25                     delta[idx] = True
26                     //unmark seen and update levelCount
27                     ...
28                 }
29

```

```

30         if (levelCount == 1) {
31             break
32         }
33     }
34
35     if (seenIdx.size > maxSize) {
36         while (trail.size > 0 && trail_starts.size > 0) {
37             DecisionLevel decision = trailStarts.last
38
39             while (trail.size > decision.idx + 1) {
40                 int lit = trail.last.lit
41                 assignment[abs(lit)] = Unassigned
42                 var_to_trail[abs(lit)] = -1
43                 erase(trail, trail.last)
44             }
45
46             if (decision.triedFalse == False) {
47                 TrailEntry entry = trail.last
48
49                 entry.lit = -entry.lit
50                 assignment[abs(entry.lit)] = False
51                 var_to_trail[abs(entry.lit)] = -1
52                 decision.triedFalse = True
53
54                 trailHead = decision.idx
55                 break
56             } else {
57                 int lit = trail.last.lit
58                 assignment[abs(lit)] = Unassigned
59                 var_to_trail[abs(lit)] = -1
60                 erase(trail.last)
61                 erase(trailStarts.last)
62             }
63
64             if (trail.size == 0) {
65                 return False
66             }
67
68             while (seenIdx.size > 0) {
69                 int idx = seenIdx.last
70                 seen[idx] = False
71                 pop(seenIdx)
72             }
73
74             for (i = 0 up to numVars) {
75                 delta[i] = False
76             }
77
78             continue
79         }

```

```

80         }
81
82         //create learned clause and update activity scores
83         int[] learned
84         while (seenIdx is not empty) {
85             int lit_idx = seenIdx.last
86             delta[lit_idx] = True
87             ...
88         }
89
90         //backjump
91         ...
92
93         //apply decay factor
94         for (i = 1 up to numVars) {
95             if (delta[i] == True) {
96                 activity[i] = (1 - Y) + Y * activity[i]
97             } else {
98                 activity[i] = Y * activity[i]
99             }
100             delta[i] = False
101         }
102
103         continue
104     }
105 }
106 }

```

So it's pretty straightforward,¹⁴ we separate our code into two branches. In the first branch where we aren't going to learn the clause, we just backtrack as we used to in our clause learning section, clear out everything related to adjusting the VSIDS and then move on, and otherwise we move on business as usual.

¹⁴I struggled with debugging for half an hour because I forgot to clear out *delta* and *seenIdx/seen*. I can see the headlines now: "Functional programmers are right about this ONE thing!!!"


```

On all 1000 equations in uf20-91:

N-order:    283ms (~0.28 seconds).
AMA:        274ms (~0.27 seconds).

On all 100 equations in Flat200-479:

N-order:    85,749ms (~1.43 minutes).
AMA:        24,238ms (~24.24 seconds).

On all 1000 equations in UUF50.218.1000:

N-order:    2,313ms (~2.31 seconds).
AMA:        1,629ms (~1.63 seconds).

```

As you can see, doing n-order learning was terrible for our runtime. Unfortunately this isn't an implementation issue in our n-order learning algorithm, unless you consider that implementing n-order learning is an implementation issue to begin with. You can probably guess why this is (we removed a ton of unit propagation!), and this is really important because it goes to show that even small initial clauses can cause much larger learned clauses. When I doubled n , the runtime was nearly the same as AMA, but a tad slower.

4.6.2 M-Size Relevance-Based Learning

The kind of clause management heuristic we're going to try is m-size relevance-based learning, where we learn every single clause, but only delete them when they stop implying assignments and have more than m Unassigned literals. To set this up, we will restart from our VSIDS code since it's not like there's anything to build off of from n-order learning. The paper we use uses the phrasing "as soon as the number of unassigned literals is greater than an integer m " which sounds like we have to keep track of the number of unassigned literals per clause. We can do this by directly tying the number of Unassigned literals to the clause in a struct, which we will do.

There is also the option of scanning the clauses in their entirety and using that information to choose which ones to keep and which ones to discard. Then, instead of doing this every conflict or whatever, we clean up the database after a certain number of conflicts.

Although adding in this struct to keep track of the number of Unassigned literals seems really simple, you have to consider that our watched literals implementation isn't a great support structure for it, since we'd need to add in a mapping from variables to clauses, then update every single clause containing a variable, and blah blah blah, it's a lot. There's also just the fact that honestly, I'm going to be debugging this for hours, the end of the quarter is draining me and I just want to graduate, and so in other words, I'm making the executive decision to not "waste" more time than I have to.

So since our clause lengths are pretty small still, we'll say that once the clause is no longer unit, we'll get rid of it, as long as it's not part of forcing an assignment. The

implementation will be extensible, ie, it won't "hardcode" that the clause must be non-unit in order to be deleted, it will still manually count the number of Unassigned literals and then compare against a threshold to decide.

Here's what we need to do to get this to work: every time there's a certain number of conflicts, we need to scan each clause if they meet our conditions. Then for each ones that do, we will mark them to be deleted. Then we just delete them! So simple! So easy! It's actually more complicated than that because of our watched literals unfortunately. We need to remove the watch lists of each clause we're deleting, update the watch lists of all variables that were watching the clause we deleted, and update the watch lists of all the clauses that get shifted in memory, which of course, is a lot easier said than done. So here's the game plan: I'm gonna *try* to get some sleep (it is 1:41AM), and when I wake up, maybe I'll have the mental fortitude to properly explain this.¹⁵

Here's what we'll need first: we're going to implement the helper functions to help update our watch lists so our watched literals will work properly. So first we'll add them to our object:

```

1      public interface:
2
3      private interface:
4          ...
5          void removeFromWatchList(int[] watchList, int idx)
6          void removeClauseFromWatchLists(int clauseIdx)
7          void replaceWatchIndex(int oldIdx, int newIdx)
8          ...

```

We will implement them in order.

```

1      void removeFromWatchList(int[] watchList, int idx) {
2          for (i = 0 up to watchList.size) {
3              if (watchList[i] == idx) {
4                  swap(watchList[i], watchList.last)
5                  pop(watchList)
6                  break
7              }
8          }
9      }

```

This algorithm is very straightforward. Given an array of clause indices that we're watching, we want to remove the passed in *clause index*, not the *element* at the passed in index. We use swap-and-pop to avoid shifting elements, which would be a bookkeeping nightmare.

```

1      void removeClauseFromWatchLists(int clauseIdx) {
2          (int firstWatch, int secondWatch) = clause_to_var[clauseIdx]
3
4          int lit1 = clauses[clauseIdx][firstWatch]

```

¹⁵My senioritis got so bad, it evolved into full blown mental illness.

```

5      int lit2 = clauses[clauseIdx][secondWatch]
6
7      int varIdx1 = var_to_widx(lit1)
8      int varIdx2 = var_to_widx(lit2)
9
10     removeFromWatchList(var_to_clause[varIdx1], clauseIdx)
11
12     if (varIdx1 != varIdx2) {
13         removeFromWatchList(var_to_clause[varIdx2], clauseIdx)
14     }
15 }

```

This helper works by first grabbing the indices of the two literals watching the clause at the passed in index. We then turn them into the actual literals by indexing the clause. Then, we get the indices of the literals to where they map in *var_to_clause*, so that way we can grab the literals' watchlists. Finally, we remove the clause index from those literals' watchlists. This is a little bit confusing because of all of the indexing and grabbing the corresponding elements we're doing.

```

1  void replaceWatchIndex(int oldIdx, int newIdx) {
2      (int firstWatch, int secondWatch) = clause_to_var(newIdx)
3
4      int lit1 = clauses[newIdx][firstWatch]
5      int lit2 = clauses[newIdx][secondWatch]
6
7      int varIdx1 = var_to_widx(lit1)
8      int varIdx2 = var_to_widx(lit2)
9
10     int[] firstWatchList = var_to_clause[var1]
11     for (i = 0 up to firstWatchList.size) {
12         if (firstWatchList[i] == oldIdx) {
13             firstWatchList[i] = newIdx
14         }
15     }
16
17     if (varIdx1 != varIdx2) {
18         int[] secondWatchList = var_to_clause[var2]
19         for (i = 0 up to secondWatchList.size) {
20             if (secondWatchList[i] == oldIdx) {
21                 secondWatchList[i] = newIdx
22             }
23         }
24     }
25 }

```

This one is pretty self explanatory. We do our indexing magic to get the watchlists for the clause's corresponding watch literals, and we replace an old index with a new index. Again, not an element *at* an old index, an element *equaling* an old index.

Here is how they are used in our m-size learning algorithm:

```

1  bool solveCNF() {
2      //setup
3      ...
4      int numConflicts = 0
5
6      while(true) {
7          if (numConflicts > threshold) {
8              bool[] remove
9
10             for (i = numClauses up to clauses.size) {
11                 int[] clause = clauses[i]
12
13                 int numUnassigned = 0
14                 for (j = 0 up to clause.size) {
15                     if (vars[abs(clause[j])] == Unassigned) {
16                         numUnassigned = numUnassigned + 1
17                     }
18                 }
19
20                 if (numUnassigned > m) {
21                     remove[i - numClauses] = True
22                 }
23             }
24
25             for (i = 0 up to trail.size) {
26                 if (trail[i].reasonIdx is not optional) {
27                     int reasonIdx = trail[i].reasonIdx.value
28
29                     if (reasonIdx >= numClauses) {
30                         remove[reasonIdx - numClauses] = False
31                     }
32                 }
33             }
34
35             for (i = clauses.size down to numClauses) {
36                 if (remove[i - numClauses] == False) {
37                     continue
38                 }
39
40                 int lastIdx = clauses.size - 1
41
42                 if (i == lastIdx) {
43                     removeClauseFromWatchLists(i)
44                     pop(clauses)
45                     pop(clause_to_var)
46                     continue
47                 }
48
49                 removeClauseFromWatchLists(i)

```

```

50         swap(clauses[i], clauses[last])
51         swap(clause_to_var[i], clause_to_var[last])
52
53         replaceWatchIndex(last, idx)
54
55         for (i = 0 up to trail.size) {
56             if (trail[i].reasonIdx is not optional) {
57                 if (trail[i].reasonIdx.value == last) {
58                     trail[i].reasonIdx.value = idx
59                 }
60             }
61         }
62
63         pop(clauses)
64         pop(clause_to_var)
65     }
66
67     numConflicts = 0
68 }
69
70 //rest of solving loop
71 ...
72 }
73 }

```

This is what's going on: first, we check that a certain number of conflicts have happened. This number is fixed. If there have been that many conflicts, we scan every clause that is **not** a part of our original CNF equation. If it has more than m Unassigned literals, we mark it for deletion. In our code, our *remove* array uses an indexing where we subtract the original number of clauses from the current clause indexing variable. This keeps the array smaller in our implementation. Next, we iterate through our trail, and we check if any of the clauses we've marked to be deleted are the reason why something was propagated. If a clause we want to delete is a reason clause, we can't delete it because it's still being used in our assignments.

Next, we iterate from the back of the clause database up to the front, excluding the original clauses. If we're not supposed to remove the current clause, we move on (duh). Otherwise, we grab the index of the last clause. If the last clause in the database is marked to be deleted, we can easily pop it off the database. Otherwise, we need to do our swap and pop nonsense. We remove the clause from its watchlists, swap the current clause with the last clause in the database, then we replace the watch indices.

After that, we traverse our trail again and check if we've swapped any literal's reason clause index. If we have, then we need to update it. After that, we can remove the clause and its watched literals from the database. After this is all over, we reset the counter for the number of conflicts and continue with the rest of the solving loop.

That's a lot of work, so how does it hold up against our other implementations?

```

On all 1000 equations in uf20-91:

M-size:      293ms (~0.29 seconds).
N-order:     283ms (~0.28 seconds).
AMA:         274ms (~0.27 seconds).

On all 100 equations in Flat200-479:

M-size:      25,681ms (~25.68 seconds).
N-order:     85,749ms (~1.43 minutes).
AMA:         24,238ms (~24.24 seconds).

On all 1000 equations in UUF50.218.1000:

M-size:      1,937ms (~1.94 seconds).
N-order:     2,313ms (~2.31 seconds).
AMA:         1,629ms (~1.63 seconds).

```

We see that m-size relevance-based learning is slightly slower than AMA VSIDS, and much faster than n-order learning (except for uf20-91). There could be some reasons for this, such as doing clause deletion too frequently. Rather than doing it a fixed number of times, we could also perform it whenever the database grows by a certain amount, say 10%. If we do this, our results speed up quite a bit on the larger equations (FLAT200.470):

```

On all 1000 equations in uf20-91:

M-size (10\%): 320ms (~0.32 seconds).
M-size:        293ms (~0.29 seconds).
N-order:       283ms (~0.28 seconds).
AMA:          274ms (~0.27 seconds).

On all 100 equations in Flat200-479:

M-size (10\%): 17,716ms (~17.72 seconds).
M-size:        25,681ms (~25.68 seconds).
N-order:       85,749ms (~1.43 minutes).
AMA:          24,238ms (~24.24 seconds).

On all 1000 equations in UUF50.218.1000:

M-size (10\%)  1,768ms (~1.77 seconds).
M-size:        1,937ms (~1.94 seconds).
N-order:       2,313ms (~2.31 seconds).
AMA:          1,629ms (~1.63 seconds).

```

Although it doesn't speed up on the smaller equations, a 7 second speedup on larger

equations shows that this scales better on larger SAT problems, which is great!

4.6.3 Combining M-Size and K-Bounded Learning

Due to the similarities of m-size learning and k-bounded learning, I felt like it didn't make much sense to implement k-bounded learning by itself, then combining it with m-size learning afterwards. Combining these two is really easy since all we're doing is checking if clauses are:

1. Greater than some size k
2. Have at least m Unassigned literals

Since we already do the second one in our m-size loop, it's just adding another condition that needs to be true in order to mark a clause to be removed.

On the other hand, what isn't easy is choosing a size k . Not all CNF equations are made equal. Some equations have small clauses (like our testing sets) and learn larger clauses, some have large clauses and learn smaller clauses, etc. Choosing a fixed k for all clauses is obviously a really bad idea. Make it too large and we just learn everything, make it too small and we forget everything. We will bring back our "maxSize" counter from n-order based learning and use that. We will use a maxSize that is twice as large as the largest clause in the original CNF¹⁶

```
On all 1000 equations in uf20-91:

MK-bound:      288ms (~0.29 seconds).
M-size (10\%): 320ms (~0.32 seconds).
M-size:        293ms (~0.29 seconds).
N-order:       283ms (~0.28 seconds).
AMA:           274ms (~0.27 seconds).

On all 100 equations in Flat200-479:

MK-bound:      19,082ms (~19.08 seconds).
M-size (10\%): 17,716ms (~17.72 seconds).
M-size:        25,681ms (~25.68 seconds).
N-order:       85,749ms (~1.43 minutes).
AMA:           24,238ms (~24.24 seconds).

On all 1000 equations in UUF50.218.1000:

MK-bound:      1,740ms (~1.74 seconds).
M-size (10\%)  1,768ms (~1.77 seconds).
M-size:        1,937ms (~1.94 seconds).
N-order:       2,313ms (~2.31 seconds).
AMA:           1,629ms (~1.63 seconds).
```

¹⁶Because adding a second condition to one conditional is easy, and iterating through the clauses to find the size of the largest clause is also easy, we will skip pseudocode.

We see that using a `maxSize` that's twice as large as the size of the largest original clause causes a pretty big slowdown on Flat200-479. This indicates that we're removing clauses too aggressively relative to the size of the clauses we're learning.

4.6.4 “Clause VSIDS”

Now we get to the more interesting part: “Clause VSIDS.”¹⁷ The idea is that just like in VSIDS for variables, we measure a clause's activity and quality (LBD), and use that to determine whether or not it should be deleted. We can still take into account things such as length, whether or not it's being used, and the other factors we've been considering in other versions as well, however, the activity and quality are the main ones.

The paper defines this heuristic as “the decision whether a clause should be deleted is based not only on the number of literals but also on its *activity* in contributing to conflict making and on the number of decisions taken since its creation.”

In order to implement this, we'll need three additional data structures:

```

1      public interface:
2
3      private interface:
4          ...
5          double[] cActivity
6          int[] creationConflict
7          int[] lbds
8          ...

```

cActivity measures a learned clause's activity. We will update the activity of every clause involved in a conflict, if it's a learned clause. Because activity is a metric we use to determine whether or not a clause should be deleted, and we never delete an original clause, we don't need to keep track of any data for any original clauses. *creationConflict* keeps track of the conflict number that a clause was created on. So if we learn a clause on conflict 50, its creation conflict is 50. We use this to measure the age of the clause, where $age = totalConflicts - creationConflict$. Lastly, *lbds* keeps track of a clause's LBD upon learn time. Even though a clause's LBD changes as our decision trail changes, we want to keep the original LBD from when the clause was learned. This is because we want to know what decision levels led to the conflict in the first place, and recomputing it describes the how “useful” the learned clause is under the current assignment, not when it was actually created. Some SAT solvers do use dynamic LBD computations, but it's a different heuristic.

In order to keep track of these metrics for each learned clause, we need to update our VSIDS code, as well as some other parts of our conflict resolution algorithm. Luckily, these changes are very minimal:

```

1      bool solveCNF() {
2          //setup
3          ...

```

¹⁷This is NOT an official or technical name as far as I am aware


```

4      int totalConflicts = 0
5      int curSize = numClauses
6
7      double activityIncrease = 1.0
8      double activityGrowth = 1.05
9      double maxIncrease = 1e100
10     double resetIncrease = 1e-100
11
12     while (true) {
13         if (activityIncrease > maxIncrease) {
14             for each (clauseActivity in cActivity) {
15                 clauseActivity *= resetIncrease
16             }
17             activityIncrease *= resetIncrease
18         }
19
20         //Current clause deletion
21         if (clauseDatabase.size grew by X%) {
22             ...
23
24             for (i = clauses.size down to numClauses) {
25                 if (remove[i - numClauses] == False) {
26                     continue
27                 }
28
29                 int lastIdx = clauses.size - 1
30
31                 if (i == lastIdx) {
32                     removeClauseFromWatchLists(i)
33                     pop(cclauses)
34                     pop(clause_to_var)
35
36                     pop(cActivity)
37                     pop(creationConflict)
38                     pop(lbds)
39
40                     continue
41                 }
42
43                 removeClauseFromWatchLists(i)
44                 swap(cclauses[i], clauses[last])
45                 swap(clause_to_var[i], clause_to_var[last])
46
47                 swap(cActivity[i], cActivity[last])
48                 swap(creationConflict[i], creationConflict[last])
49                 swap(lbds[i], lbds[last])
50
51                 ...
52
53                 pop(cclauses)

```

```

54         pop(clause_to_var)
55
56         pop(cActivity)
57         pop(creationConflict)
58         pop(lbds)
59     }
60
61     curSize = clauses.size
62 }
63
64 //assignment
65 ...
66
67 if (contradiction occurred) {
68     totalConflicts = totalConflicts + 1
69     ...
70
71     for (int trailIdx = trail.size down to 0) {
72         ...
73         int reasonIdx = trail[trailIdx].reasonIdx.value;
74         if (reasonIdx >= numClauses) {
75             cActivity[reasonIdx - numClauses] +=
activityIncrease
76         }
77     }
78     //rest of VSIDS and LBD calculation
79     ...
80
81     //backjump and learn
82     ...
83
84     append(cActivity, 0.0)
85     append(creationConflict, totalConflicts)
86     append(lbds, lbd)
87     activityIncrease *= activityGrowth
88     ...
89 }
90 }
91 }

```

In our clause deletion code, we just added code to update our new data structures to be consistent with the rest of the clause database data structures. And then in the VSIDS calculation, we added code to update a clause's activity if it's a learned clause, and at the end of the loop, we just add extra entries into each of the data structures so they grow with the database.

We use an activity increase and an activity growth variable to simulate the effects of an activity decay and bump. The idea is that instead of bumping only the clauses that were involved in a conflict and decaying the rest, we bump only the clauses involved in a conflict in a way that is as if we decayed the others. By increasing the activity

increase by some growth amount every conflict, when we add to a clause's activity, we add an ever increasing amount, so that way, more recent clauses get larger and larger activities and less recent clauses don't have such huge increases. The only caveat to this is that because conflicts happen all the time, we need to reset the activity increase once it grows to a stupidly large size. This can be thought of as our decay step.

The next step is to now use these metrics in our clause deletion code. We could do a bunch of experiments of how each metric impacts runtime and clause deletion, but frankly, I don't want to. Instead, here's what workflow we'll use:

1. Find which clauses that can be deleted by LBD, age, and activity
2. Find which clauses can be deleted via $m+k$ -bounded learning
3. Protect clauses with an $LBD \leq 2$
4. Protect clauses that are currently reason clauses
5. Delete the "bad" ones

To determine what a "bad" clause means, we need to think of what the goals of clause deletion is. For instance, do we want to delete a fixed number of clauses, or should it be more dynamic, as in, there isn't a quota to fulfill, instead we look at each clause and go "you suck, get out" or "you're good, stay?" This idea is pretty attractive, but there is something we should consider: what if there's not enough bad clauses relative to our database size, or what if there's too many bad clauses relative to our database size? Ie, what if we mark too little clauses for deletion and our database is still large, as well as marking too many clauses for deletion and then our database becomes super small again? To be honest, if there aren't enough "bad" clauses, to delete, to me it seems like most of the database is really good, so I wouldn't mark *extra* clauses for deletion, however, capping how many clauses we delete does seem like a good idea.

Because we have a clear marker of when clause deletion should occur (when the database grows by 10%), we will say that the number of clauses being deleted can't exceed this 10% threshold. So if our original database has 100 clauses and it grows to 110 clauses, we can't delete more than 9 clauses. Supposing we delete 0, and the database grows to 121, we can't delete more than 10 clauses, and so on. Another advantage of this is that the maximum number of clauses we can delete grows/shrinks with our threshold, so if in a later experiment we found that it's more beneficial to run the deletion algorithm every time the database grows by 20%, we don't have to manually update the number of clauses we're allowed to delete as well.

In order to decide which clauses to delete, we want to decide a threshold for our metrics. I used the ones in the bottom 25th percentile for their respective metric, so high LBD, high age, and low activity. We then select the clauses in those bottom metrics, and if they have at least 2 of these negative attributes, we mark them to be deleted.

After that, we do our regular $m+k$ bounded learning, then afterwards, we mark clauses that have low LBD (≤ 2), and clauses that are reason clauses as protected, so they can't be deleted even if they had at least two bad attributes.

The code is as follows:

```

1  bool solveCNF() {
2      //setup
3      ...
4
5      while (true) {
6          //Current clause deletion
7          if (clauseDatabase.size grew by X%) {
8              int numCanDelete = clauseDatabase.size - curSize
9
10             int[] candidateIdx
11             for (i = 0 up to clauseDatabase.size - numClauses) {
12                 candidateIdx[i] = i
13             }
14
15             int cutoff = candidateIdx.size * 0.25
16
17             sortIncreasing(cActivity)
18
19             double activityCutoff = cActivity[cutoff]
20
21             sortIncreasing(conflictCreation)
22
23             int ageCutoff = conflictCreation[cutoff]
24
25             sortDecreasing(lbds)
26
27             int lbdCutoff = lbds[cutoff]
28
29             for (i = numClauses up to clauses.size) {
30                 bool lowActivity = cActivity[i] <= activityCutoff
31                 bool highAge = creationConflict[i] <= ageCutoff
32                 bool highlbd = lbds[i] >= lbdCutoff
33
34                 remove[i - numClauses] = lowActivity + highAge +
highlbd >= 2
35             }
36
37             //rest of the database deletion
38             ...
39         }
40
41         //rest of solving
42         ...
43     }
44
45 }

```

```

On all 1000 equations in uf20-91:

Clause VSIDS:    347ms (~0.35 seconds).
MK-bound:        288ms (~0.29 seconds).
M-size (10\%):   320ms (~0.32 seconds).
M-size:          293ms (~0.29 seconds).
N-order:         283ms (~0.28 seconds).
AMA:             274ms (~0.27 seconds).

On all 100 equations in Flat200-479:

Clause VSIDS:    19,181ms (~19.18 seconds).
MK-bound:        19,082ms (~19.08 seconds).
M-size (10\%):   17,716ms (~17.72 seconds).
M-size:          25,681ms (~25.68 seconds).
N-order:         85,749ms (~1.43 minutes).
AMA:             24,238ms (~24.24 seconds).

On all 1000 equations in UUF50.218.1000:

Clause VSIDS:    1,804ms (~1.80 seconds).
MK-bound:        1,740ms (~1.74 seconds).
M-size (10\%)    1,768ms (~1.77 seconds).
M-size:          1,937ms (~1.94 seconds).
N-order:         2,313ms (~2.31 seconds).
AMA:             1,629ms (~1.63 seconds).

```

To be honest, I'm a little bit disappointed in these results. I was really hopeful that this heuristic would have yielded more gains in speed, not runtime. Was my implementation perfect? No. Could that have contributed to the increased runtime? Yes, but I have a feeling that it's also the testing sets being used. Again, we can't necessarily draw a conclusion on *all* SAT problems based on this data, but we can draw the conclusion that the solve is worse for *these* SAT problems.

There were some things I tried to see what could be causing this. I tried adjusting how aggressive/loose the badness score had to be in order to be marked to be deleted. A lower badness score threshold correlated with a higher runtime. Then I tried removing the cap on deletion altogether, and that didn't impact the runtime very much. Then I tried adjusting the percentile to be the 20th, 25th, and 50th, and they all had very similar runtimes as well. I even tried messing with what we consider a bad clause, and that impacted the runtime, in a bad way. Lastly, I tried removing mk-bound learning, and that also didn't make much of an impact. I then removed the k-bounded learning afterwards, and that also didn't seem to help. No matter what I tried, it appears that this heuristic was hurting the runtime on all accounts. Perhaps something to be explored at a later date.

4.6.5 Clause Database Management - Results

Here are the results for all the kinds of clause database management compared to the last implementation without any kind of management:

On all 1000 equations in uf20-91:

Clause VSIDS:	347ms (~0.35 seconds).
MK-bound:	288ms (~0.29 seconds).
M-size (10\%):	320ms (~0.32 seconds).
M-size:	293ms (~0.29 seconds).
N-order:	283ms (~0.28 seconds).
AMA:	274ms (~0.27 seconds).
EMA (X = 1, Y = 0.95):	266ms (~0.27 seconds).
Improved CDCL:	284ms (~0.28 seconds).
First UIPs:	406ms (~0.41 seconds).
Backjumping:	380ms (~0.38 seconds).
Clause Learning:	425ms (~0.43 seconds).
Iterative DPLL (AbP):	328ms (~0.33 seconds).
Iterative DPLL (PbA):	324ms (~0.32 seconds).
Watched Literals:	349ms (~0.35 seconds).
Pure Literal Elimination:	1548ms (~1.55 seconds).
Unit Propagation:	868ms (~0.87 seconds).
Recursive Backtracking:	1,540,613ms (~25.68 minutes).
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).

On all 100 equations in Flat200-479:

Clause VSIDS:	19,181ms (~19.18 seconds).
M-size (10\%):	17,716ms (~17.72 seconds).
AMA:	24,238ms (~24.24 seconds).

On all 1000 equations in UUF50.218.1000:

Clause VSIDS:	1,804ms (~1.80 seconds).
MK-bound:	1,740ms (~1.74 seconds).
M-size (10\%)	1,768ms (~1.77 seconds).
M-size:	1,937ms (~1.94 seconds).
N-order:	2,313ms (~2.31 seconds).
AMA:	1,629ms (~1.63 seconds).
EMA (X = 1, Y = 0.95):	1,504ms (~1.50 seconds).
Improved CDCL:	2,385ms (~2.39 seconds).
First UIPs:	14,012ms (~14.01 seconds).
Backjumping:	11,291ms (~11.29 seconds).
Clause Learning:	9,652ms (~9.65 seconds).
Iterative DPLL (AbP):	6,369ms (~6.37 seconds).
Iterative DPLL (PbA):	6,365ms (~6.37 seconds).
Watched Literals:	7,130ms (~7.13 seconds).
Pure Literal Elimination:	231,107ms (~3.85 minutes).
Unit Propagation:	140,041ms (~2.34 minutes).

Because of how well M-size (10%) performed, I will keep it in the overall results when compared with every other solver we've created so far:

On all 1000 equations in uf20-91:

Clause VSIDS:	347ms (~0.35 seconds).
M-size (10\%):	320ms (~0.32 seconds).
AMA (Y = 0.75 or 0.99):	274ms (~0.27 seconds).
EMA (X = 1, Y = 0.95):	266ms (~0.27 seconds).
Improved CDCL:	284ms (~0.28 seconds).
First UIPs:	406ms (~0.41 seconds).
Backjumping:	380ms (~0.38 seconds).
Clause Learning:	425ms (~0.43 seconds).
Iterative DPLL (AbP):	328ms (~0.33 seconds).
Iterative DPLL (PbA):	324ms (~0.32 seconds).
Watched Literals:	349ms (~0.35 seconds).
Pure Literal Elimination:	1548ms (~1.55 seconds).
Unit Propagation:	868ms (~0.87 seconds).
Recursive Backtracking:	1,540,613ms (~25.68 minutes).
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).

On all 100 equations in Flat200-479:

Clause VSIDS:	19,181ms (~19.18 seconds).
M-size (10\%):	17,716ms (~17.72 seconds).
AMA:	24,238ms (~24.24 seconds).

On all 1000 equations in UUF50.218.1000:

Clause VSIDS:	1,804ms (~1.80 seconds).
M-size (10\%)	1,768ms (~1.77 seconds).
AMA:	1,629ms (~1.63 seconds).
EMA (X = 1, Y = 0.95):	1,504ms (~1.50 seconds).
Improved CDCL:	2,385ms (~2.39 seconds).
First UIPs:	14,012ms (~14.01 seconds).
Backjumping:	11,291ms (~11.29 seconds).
Clause Learning:	9,652ms (~9.65 seconds).
Iterative DPLL (AbP):	6,369ms (~6.37 seconds).
Iterative DPLL (PbA):	6,365ms (~6.37 seconds).
Watched Literals:	7,130ms (~7.13 seconds).
Pure Literal Elimination:	231,107ms (~3.85 minutes).
Unit Propagation:	140,041ms (~2.34 minutes).

Sorry that this section is really bloated.

4.7 Restarts

This is the last thing we will implement in our SAT solver (for now, maybe). Restarts are exactly what it sounds like. We completely delete the trail and all notion of assignment. Why would we do that? A valid question. The idea is that perhaps while we've been solving, we got stuck in a bad part of the search space and are just getting contradiction after contradiction. Rather than staying stuck there, perhaps we could restart our solver so that it will end up in a different part of the search space that might be better.

Makes enough sense, but how do we ensure that we don't just go back to the same part of the search space again? Well luckily because we've been learning, our solver isn't exactly the same as it initially was. It "knows" more, and since knowledge is power, the new learned clauses can force certain assignments and propagations earlier and that will nudge or push us into the right direction.

So implement this, it's really really simple. All we need to do is after a certain number of conflicts, clear **ONLY** the trail. We don't clear out our VSIDS, clause measurements, reset the watched literals, nothing, nada, nahin. Why? Because the idea of restarts is that when the going gets tough, we get going. Notice how I said that we get going, not that we get forgetful? We go back to the beginning, remembering everything we've learnt, and use that to (hopefully) not repeat mistakes. If we forgot things like our VSIDS, we would make the same assignment with the exact same decision, and we'd have to rely only on our learned clauses to sway our decision making. But since VSIDS tells us which variables have been important in conflicts, it tells us which variables are best to start out with on our new try.

For our solver, I chose to restart every time we clear out the clause database. If you've been loopy-goosy about following my implementation or have been doing your own experiments (nice), you can define your own condition. Just remember these things when you're doing a restart:

- Clear the entire trail.
- The `trail_starts` array must always have a decision level of 0.
- Clearing out `var_to_trail` is not strictly necessary as long as your implementation is correct.
- If you do clear out `var_to_trail`, make sure it's resized properly.
- Reset `qHead`.
- Reset the decision level to 0.
- Unassign every variable.
- You might want to do a root level propagation.

Here are the results:

On all 1000 equations in uf20-91:

Restarts:	385ms (~0.39 seconds).
Clause VSIDS:	347ms (~0.35 seconds).
M-size (10\%):	320ms (~0.32 seconds).
AMA (Y = 0.75 or 0.99):	274ms (~0.27 seconds).
EMA (X = 1, Y = 0.95):	266ms (~0.27 seconds).
Improved CDCL:	284ms (~0.28 seconds).
First UIPs:	406ms (~0.41 seconds).
Backjumping:	380ms (~0.38 seconds).
Clause Learning:	425ms (~0.43 seconds).
Iterative DPLL (AbP):	328ms (~0.33 seconds).
Iterative DPLL (PbA):	324ms (~0.32 seconds).
Watched Literals:	349ms (~0.35 seconds).
Pure Literal Elimination:	1548ms (~1.55 seconds).
Unit Propagation:	868ms (~0.87 seconds).
Recursive Backtracking:	1,540,613ms (~25.68 minutes).
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).

On all 100 equations in Flat200-479:

Restarts:	9,428ms (~9.43 seconds).
Clause VSIDS:	19,181ms (~19.18 seconds).
M-size (10\%):	17,716ms (~17.72 seconds).
AMA:	24,238ms (~24.24 seconds).

On all 1000 equations in UUF50.218.1000:

Restarts:	1,824ms (~1.82 seconds).
Clause VSIDS:	1,804ms (~1.80 seconds).
M-size (10\%)	1,768ms (~1.77 seconds).
AMA:	1,629ms (~1.63 seconds).
EMA (X = 1, Y = 0.95):	1,504ms (~1.50 seconds).
Improved CDCL:	2,385ms (~2.39 seconds).
First UIPs:	14,012ms (~14.01 seconds).
Backjumping:	11,291ms (~11.29 seconds).
Clause Learning:	9,652ms (~9.65 seconds).
Iterative DPLL (AbP):	6,369ms (~6.37 seconds).
Iterative DPLL (PbA):	6,365ms (~6.37 seconds).
Watched Literals:	7,130ms (~7.13 seconds).
Pure Literal Elimination:	231,107ms (~3.85 minutes).
Unit Propagation:	140,041ms (~2.34 minutes).

Even though our solver is getting slower and slower on uf20-91, I think it's actually really good that for the larger problems, it's getting substantially faster. Even though we want our algorithms and programs to work really quickly on small inputs, I think it's important for us to at some point prioritize the larger inputs. And given how most SAT problems are even larger than these, I think that's a good thing.

4.7.1 Assignment Polarity

Back in the day, we had this thing called “Pure Literal Elimination,” but when we transitioned to CDCL, it went away *forever*. Well dear reader, I'm here to tell you that it's still going to be gone forever, but its spirit lives on with us. And by that, I mean, we're going to take the idea behind pure literal elimination and apply it to our restarts.

You see, as we continue learning new clauses, we also encounter new literals. Earlier, back when we switched from brute force to recursive backtracking, we noticed that solving times decreased. Other than the main structure of the solvers, the only thing that changed was in brute force, we started out with most variable assignments being False, while in recursive backtracking and beyond, they start out as True. I made a claim that this speedup was because there were less negative literals than there were positive literals. I won't say that I knew for sure, but if we stop for a second to think about it, it makes sense.

When there are more positive literals, the likelihood of a variable's final assignment being True is higher than when there are more negative literals. This was the idea behind pure literal elimination, to an extent. Pure literal elimination worked because when a literal only shows up with one polarity, if there is a solution, there is a 100% chance that there is at least one solution where that pure literal's assignment matches its polarity (T for positive, F for negative). We want to bring this idea into our assignment polarity. While we can't guarantee that literals will remain pure because of clause learning, we can at least make better decisions based on the current number of positive/negative occurrences of a variable.

To do this, we'll add in a *polarity* array to our solver:

```
1      public interface:
2
3      private interface:
4          ...
5          int[] polarity
6          ...
```

Instead of using a count for both positive and negative literals, instead, we'll use one counter for each variable. When the count is negative, we know that there are more negative occurrences of the variable, and when the count is nonnegative, we will pretend that there are more positive occurrences of the variable (even though when the count is 0, that means it's balanced).

Because we're using one counter per variable, it might be tempting to use a literal's value directly in the count. After all, if we encounter one +4 and one -4, they'd cancel

out to 0. Although this is logically correct, I'm worried about integer overflow for equations with millions of variables¹⁸

Instead we'll settle for an "if negative => count - 1, else count + 1" kind of deal. We will gather the counts in the setup phase of our solver, and we will update the counts only when we learn a new clause. However, it would be interesting to update the counts for all variables involved in a conflict, and see how that affects things. Because of how easy this code is to add as well, I'm gonna let you figure it out. Lastly, when we pick a variable via its VSIDS, we assign it based on its polarity. Also very easy to figure out, you got this, I believe in you.

Here are the results:

¹⁸ $2^{32} \approx 4.3$ billion, so if we had one million variables, we could only handle 43 occurrences of one polarity, which is pretty low, all things considered.

On all 1000 equations in uf20-91:

Polarity:	369ms (~0.37 seconds).
Restarts:	385ms (~0.39 seconds).
Clause VSIDS:	347ms (~0.35 seconds).
M-size (10\%):	320ms (~0.32 seconds).
AMA (Y = 0.75 or 0.99):	274ms (~0.27 seconds).
EMA (X = 1, Y = 0.95):	266ms (~0.27 seconds).
Improved CDCL:	284ms (~0.28 seconds).
First UIPs:	406ms (~0.41 seconds).
Backjumping:	380ms (~0.38 seconds).
Clause Learning:	425ms (~0.43 seconds).
Iterative DPLL (AbP):	328ms (~0.33 seconds).
Iterative DPLL (PbA):	324ms (~0.32 seconds).
Watched Literals:	349ms (~0.35 seconds).
Pure Literal Elimination:	1548ms (~1.55 seconds).
Unit Propagation:	868ms (~0.87 seconds).
Recursive Backtracking:	1,540,613ms (~25.68 minutes).
Brute Force:	1,770,851ms (~29.51 minutes).
Short-Circuited (Basic):	49,874ms (~50 seconds).
Parallel:	10,589ms (~10.5 seconds).

On all 100 equations in Flat200-479:

Polarity:	8,580ms (~8.58 seconds).
Restarts:	9,428ms (~9.43 seconds).
Clause VSIDS:	19,181ms (~19.18 seconds).
M-size (10\%):	17,716ms (~17.72 seconds).
AMA:	24,238ms (~24.24 seconds).

On all 1000 equations in UUF50.218.1000:

Polarity:	1,897ms (~1.90 seconds).
Restarts:	1,824ms (~1.82 seconds).
Clause VSIDS:	1,804ms (~1.80 seconds).
M-size (10\%)	1,768ms (~1.77 seconds).
AMA:	1,629ms (~1.63 seconds).
EMA (X = 1, Y = 0.95):	1,504ms (~1.50 seconds).
Improved CDCL:	2,385ms (~2.39 seconds).
First UIPs:	14,012ms (~14.01 seconds).
Backjumping:	11,291ms (~11.29 seconds).
Clause Learning:	9,652ms (~9.65 seconds).
Iterative DPLL (AbP):	6,369ms (~6.37 seconds).
Iterative DPLL (PbA):	6,365ms (~6.37 seconds).
Watched Literals:	7,130ms (~7.13 seconds).
Pure Literal Elimination:	231,107ms (~3.85 minutes).
Unit Propagation:	140,041ms (~2.34 minutes).

That's a pretty decent speedup on the larger equations. Not sure how I feel about the speedup on UUF50.218.1000, but since it's not even 100ms across all 1000 cases, I think I'll take it.

Chapter 5

Conclusion

Well that was a journey. In 9 weeks, I have learned how to build a fully functioning (but definitely not industry-level) SAT solver, and hopefully in the amount of time it took you to read through all of this, you did too. Despite how much we’ve done, there is a much larger world of SAT solving out there. CDCL and DPLL are just two algorithms out of many that can be used to solve SAT problems. Some are better than others in certain cases. That is why, like some ML models, SAT solvers also use an “ensemble-like” method of solving. Unlike ML ensembles, which takes the output of multiple models and “weighs” them to form the final output, these SAT solvers run multiple different solving algorithms and heuristics at once, and whichever gets an answer first is the winner. This is because those algorithms will all have the same output, some just get there faster.

There are also parallel algorithms for SAT solving, so instead of one thread solving one problem, you can have multiple threads solving one problem or multiple threads solving multiple problems.

Outside of solving, there is also the notion of simplification of the problem in general. These preprocessing and inprocessing algorithms are capable of removing redundant clauses and variables from the problem entirely, thus making it easier to solve. Likewise, we can also strengthen clauses. Even cooler, we can solve a CNF problem, add more clauses and variables, then use our previous solution to try to help solve those new constraints. This is called incremental SAT, which is a really important feature.

Obviously if I had more than 9 weeks ¹ to write and learn about SAT solvers, we could learn how to extend it even further, perhaps even one day being competition worthy. There is so much more out there, and I think it would be amazing if we got to explore that more in a second draft.

5.1 Improvements

Before I end this for good, I want to list some improvements needed for this paper and our solving.

¹It was a self-imposed timeline, but that’s irrelevant.

5.1.1 This Paper is Poorly Written

You know it, I know it, everyone knows it. This would never pass academic review. You're thinking "this guy has never written anything before" and you're 100% right, I haven't. I need to go back, review it, reformat, add actual citations, etc. Is it my fault it's in a poorly written state to begin with? 100%, and I while I won't be making this project my entire life until I die, I will go back every now and then to fix things, reword stuff, reformat, whatever needs done. You can help me out by first reading this, then contacting me on revisions I should make. Any advice and help is really appreciated, and I'm writing this for the people who are interested, so by showing interest in this, it tells me how much more I need to put in to make it better for y'all.

One thing I'd like to do is redo the chapter on CDCL so it's not one mega-chapter and it could be separated into its necessary parts (clause learning, backjumping) and then the things to improve it those aspects (VSIDS, first UIPs, restarts) could be in their own chapter, then lastly, a chapter on the things that we need to keep in mind when we implement on actual hardware (clause database management). This would make the flow of everything more isolated and easier to follow along rather than following in my exact footsteps.

5.1.2 The Code Needs to be Cleaned

I said to myself "I'll refactor this nicely later" and later never came, and before I knew it, I've finished writing everything I set out to write! The code is supposed to be used as a teaching aide, and when it's not very clean, it's not as good as it could be. Sure there are code comments to help explain some things where the implementation differ from the pseudocode, but that doesn't make the code more readable, just easier to understand if you *can* read it.

Speaking of pseudocode, it's not in great shape. The pseudocode is really similar to the actual code, which is a strength and a curse. I'm not great at writing pseudocode, and I think that it's a skill that I need to develop more. If I were better at writing pseudocode, the first thing I would do is reformat it, then make it less C-based so it's actual pseudocode instead of "slightly more abstract C++." This would tie in better to the theory of SAT solving, then the actual code could be used as a teaching aide on how it's implemented.

5.1.3 We Don't Explore Other Real Solvers

Since this is a paper on how to build a solver, the main focus was on how to build one, but we ignore how other solvers have been built. This could be a good chance to explore the history of SAT solving, from the people, the algorithms, and the programs that have been built throughout the years. And through exploring other solvers and algorithms, we could have new chapters and sections on how they work and how to implement them, which would be pretty cool.

5.1.4 There Are No Actual Citations

This one is pretty self explanatory. Even though I tell you where I got my facts and figures, I don't properly cite them. I'm pretty sure this paper would be enough to get me expelled for plagiarism due to how I don't use proper citations anywhere.

And with that, those are all I can think of that needs to be done for this paper. Thank you for reading until the end. Til next time or something.